

# KV Cache Compression

Gabriel Pernell, Alexander Martin



# Background & Motivation



# Prefix Caching (Prefilling)


## Request #1

A chat between a curious user and an artificial intelligence assistant. The assistant gives helpful, detailed, and polite answers to the user's questions.  
User: Hello!

## Request #2

A chat between a curious user and an artificial intelligence assistant. The assistant gives helpful, detailed, and polite answers to the user's questions.  
User: How are you?

## Request #1



## Request #2



## Request #3



■ Prefill   ■ Prefill (Cached)   ■ Decode

# Token Selection / Eviction

Tokens



*I think , therefore I am .*


Selected  
Tokens



# Reminder: Longer Context = Longer Inference Time

Long inference times motivate using a KV Cache


- Cache the keys and values to reduce redundancy and save inference time
- Memory required:  $2bndLk$
- Larger models  $\rightarrow$  many GBs of memory



# KV Cache Compression

KV cache memory consumption with larger models motivates cache compression & optimization methods.

- Existing methods use a fixed cache size per transformer layer
  - Is this efficient?
  - Authors say no, does not take into account attention varying by layer.



# Goals & Research Questions

---

- Do LLMs aggregate information in recognizable patterns across layers?
  - If so, can this inform a smarter KV cache compression method?
- Goal: develop a compression method that
  - Allocates KV cache size based on layer and attention patterns
  - Preserve long-context performance while reducing memory



# PyramidKV: Dynamic KV Cache Compression based on Pyramidal Information Funneling

Zefan Cai , Yichi Zhang , Bofei Gao , Yuliang Liu ,  
Yucheng Li , Tianyu Liu , Keming Lu , Wayne  
Xiong , Yue Dong , Junjie Hu , Wen Xiao

Presented by Gabriel Pernell




# Background & Motivation



# Reminder: Longer Context = Longer Inference Time

Long inference times motivate using a KV Cache


- Cache the keys and values to reduce redundancy and save inference time
- Memory required:  $2bndLk$
- Larger models  $\rightarrow$  many GBs of memory



# KV Cache Compression

KV cache memory consumption with larger models motivates cache compression & optimization methods.

- Existing methods use a fixed cache size per transformer layer
  - Is this efficient?
  - Authors say no, does not take into account attention varying by layer.



# Goals & Research Questions

---

- Do LLMs aggregate information in recognizable patterns across layers?
  - If so, can this inform a smarter KV cache compression method?
- Goal: develop a compression method that
  - Allocates KV cache size based on layer and attention patterns
  - Preserve long-context performance while reducing memory



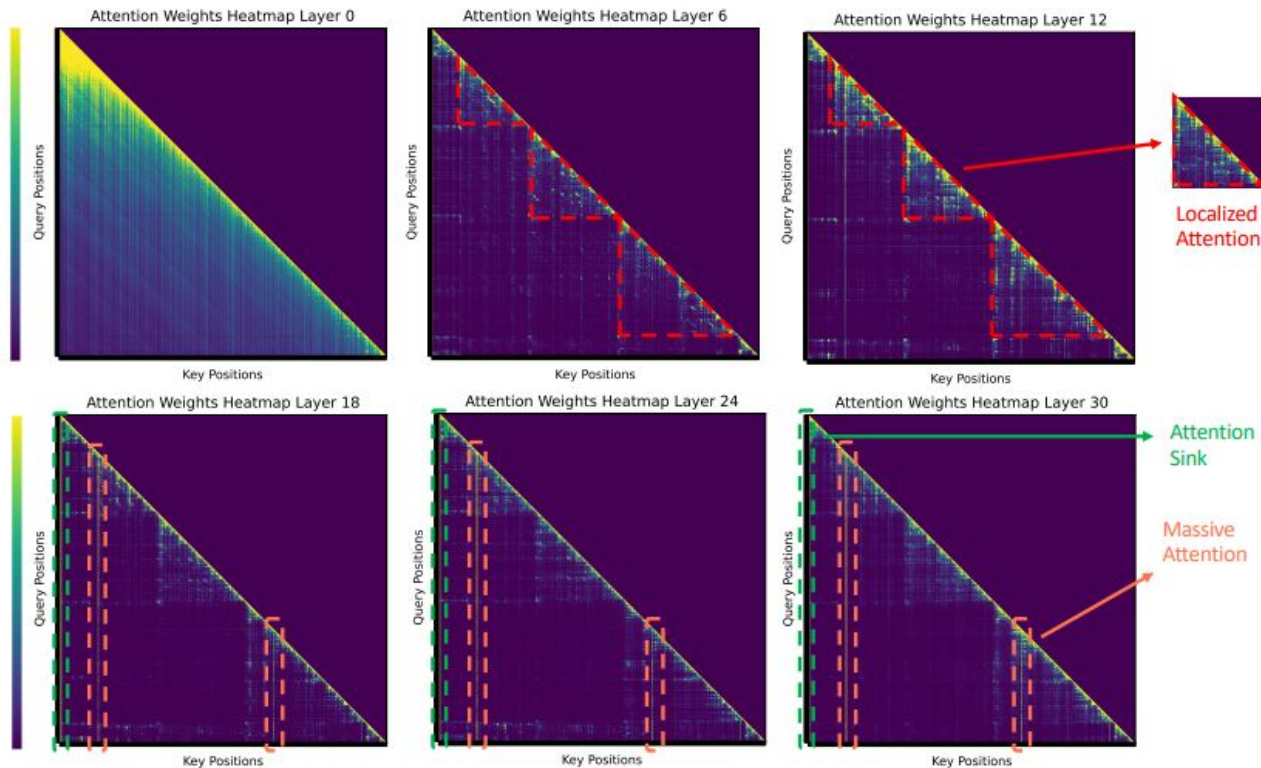
# Pyramidal Information Funneling



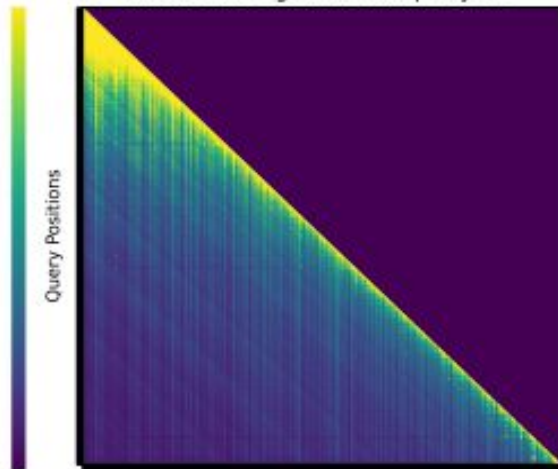
# Observational Study: Information Flow via Attention

Setup: Multi-document QA task on LLaMA: visualized attention scores across layers


- Model given several interrelated documents plus a question.



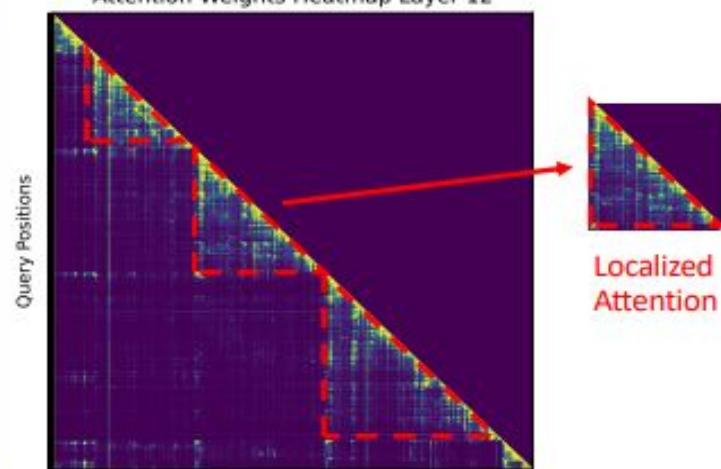
Attention Weights Heatmap Layer 0




Attention Weights Heatmap Layer 6




Attention Weights Heatmap Layer 12




Attention Weights Heatmap Layer 18



Attention Weights Heatmap Layer 24



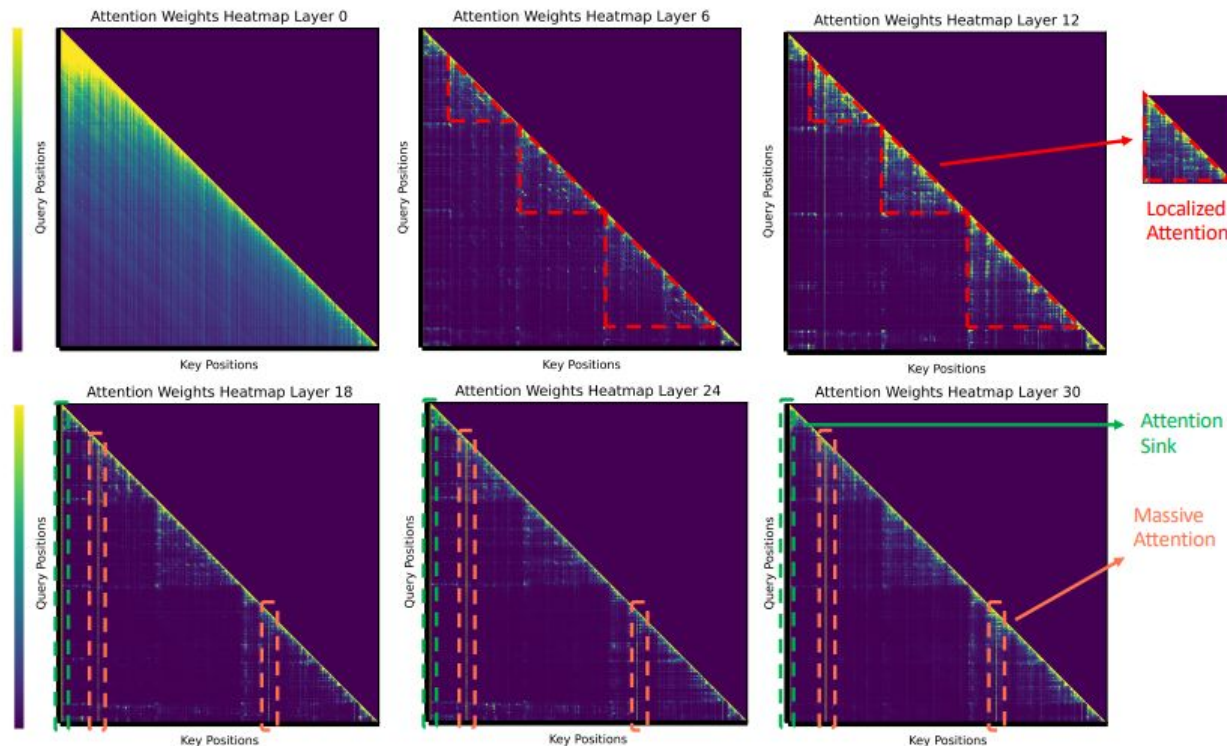
Attention Weights Heatmap Layer 30



# Findings & Interpretation

- **Findings:**

- **Lower Layers:** broad, uniform attention
- **Middle Layers:** localized attention (e.g. info within documents)
- **Upper layers:**  
“massive attention” /  
attention sink, focus on a few key tokens
- Attention narrows like a pyramid






# PyramidKV



# PyramidKV: Pyramidal KV Cache Compression Method

More cache is allocated at lower levels, less at higher levels.



# Two Key Components (1)

---

- Dynamic KV Cache Budget Allocation:

- Retain the KV cache for the last  $\alpha$  tokens (instruction tokens)

- Determine the top and bottom layer budgets:      Total cache budget:  $k^{\text{total}} = \sum_{l \in [0, m-1]} k^l$

- Use an arithmetic sequence to compute cache sizes in between, forming the pyramidal shape.

Top:  $k^{m-1} = k^{\text{total}} / (\beta \cdot m)$

Bottom:  $k^0 = (2 \cdot k^{\text{total}}) / m - k^{m-1}$

$$k^l = k^0 - \frac{k^0 - k^{m-1}}{m-1} \times l.$$



# Two Key Components (2)

---

## KV Cache Selection

- Which tokens do we keep?
  - Instruction tokens - Keep
  - Compute how much each token is attended to by the instruction tokens
  - Keep the tokens with the highest attention scores

$$s_i^h = \sum_{j \in [n-\alpha, n]} A_{ij}^h$$

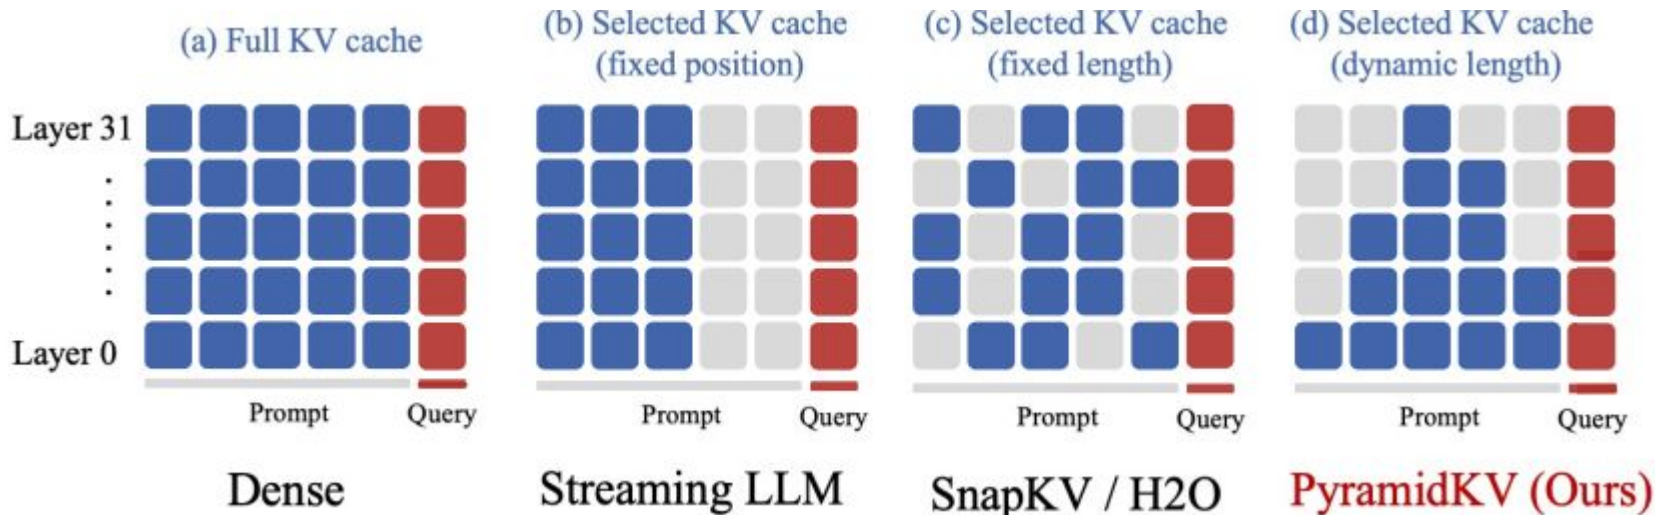


# Experiments & Results



# Setup

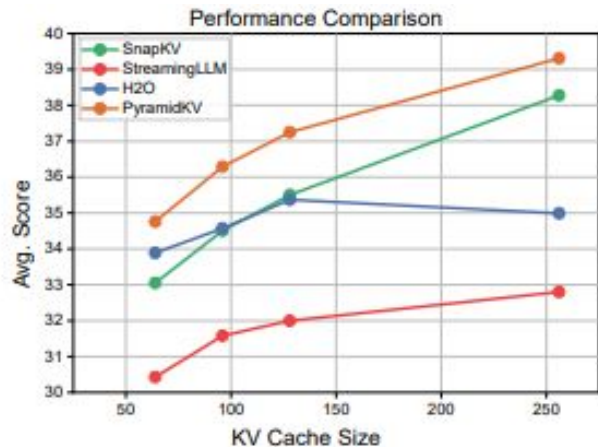
- Models: LLaMA-3-8B, LLaMA-3-70B, Mistral-7B
- Benchmark: LongBench (17 datasets across QA, summarization, code, few-shot learning).
- Baselines: FullKV, StreamingLLM, H2O, SnapKV.
- Same total KV budget on average across methods.




# Results

- PyramidKV consistently outperforms baselines, especially with small cache sizes.
- Maintains near-full performance using only 12% of full KV cache
- Even with 0.7% of KV cache, accuracy drop is minimal

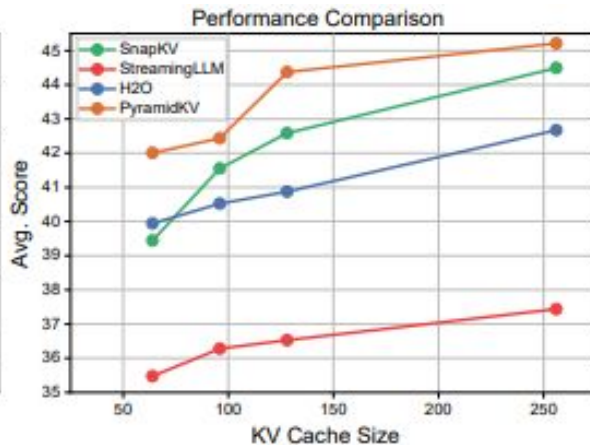
## LLaMA-3-8B



## Mistral-7B



## LLaMA-3-70B



Avg score across datasets for 64, 96, 128, and 256 cache sizes



# Results Cont'd

- PyramidKV excels with small cache sizes:

Method	Single-Document QA			Multi-Document QA			Summarization			Few-shot Learning			Synthetic		Code		Avg.
	NrtvQA	Qasper	MF-en	HotpotQA	2WikiMQA	Musique	GovReport	QMSum	MultiNews	TREC	TriviaQA	SAMSum	PCount	PRe	Lcc	RB-P	
	18409	3619	4559	9151	4887	11214	8734	10614	2113	5177	8209	6258	11141	9289	1235	4206	
LlaMa-3-70B-Instruct, KV Size = Full																	
FKV	27.75	46.48	49.45	52.04	54.90	30.42	32.37	22.27	27.58	73.50	92.46	45.73	12.50	72.50	40.96	63.91	46.55
LlaMa-3-70B-Instruct, KV Size = 64																	
SKV	23.92	31.09	36.54	46.66	50.40	25.30	18.05	21.11	19.79	41.50	<b>91.06</b>	40.26	12.00	<b>72.50</b>	43.33	57.62	39.45
SLM	22.07	23.53	27.31	43.21	<b>51.66</b>	23.85	16.62	19.74	15.20	39.50	76.89	33.06	12.00	<b>72.50</b>	40.23	50.20	35.47
H2O	25.45	34.64	33.23	<b>48.25</b>	50.30	24.88	20.03	21.50	21.39	42.00	90.36	<b>41.58</b>	12.00	71.50	43.83	<b>58.16</b>	39.94
Ours	<b>25.47</b>	<b>36.71</b>	<b>42.29</b>	47.08	46.21	<b>28.30</b>	<b>20.60</b>	<b>21.62</b>	<b>21.62</b>	<b>64.50</b>	89.61	41.28	<b>12.50</b>	<b>72.50</b>	<b>45.34</b>	56.50	<b>42.01</b>
LlaMa-3-70B-Instruct, KV Size = 2048																	
SKV	26.73	45.18	47.91	<b>52.00</b>	<b>55.24</b>	30.48	28.76	22.35	27.31	72.50	92.38	45.58	12.00	<b>72.50</b>	<b>41.52</b>	<b>69.27</b>	46.36
SLM	26.69	41.01	35.97	46.55	52.98	25.71	27.81	20.81	27.16	69.00	91.55	44.02	12.00	72.00	41.44	68.73	43.96
H2O	<b>27.67</b>	<b>46.51</b>	<b>49.54</b>	51.49	53.85	29.97	28.57	<b>22.79</b>	<b>27.53</b>	59.00	<b>92.63</b>	<b>45.94</b>	12.00	72.50	41.39	63.90	45.33
Ours	27.22	46.19	48.72	51.62	54.56	<b>31.11</b>	<b>29.76</b>	22.50	27.27	<b>73.50</b>	91.88	45.47	12.00	72.50	41.36	69.12	<b>46.55</b>







# Needle In A Haystack Experiment

- Purpose: test long-context factual retrieval
- Result: LLaMa-3-70B achieves 100% accuracy with 128 KV entries using PyramidKV, matching full cache performance
- Significance: PyramidKV preserves long-range memory and retrieval ability.


LLaMA-3-70B - 8K Context Size



(a) FullKV, KV Size = Full, acc 100.0



(c) SnapKV, KV Size=128, acc 98.6



(b) PyramidKV, KV Size=128, acc 100.0




(d) H2O, KV Size=128, acc 82.3

# KVzip: Query-Agnostic KV Cache Compression with Context Reconstruction

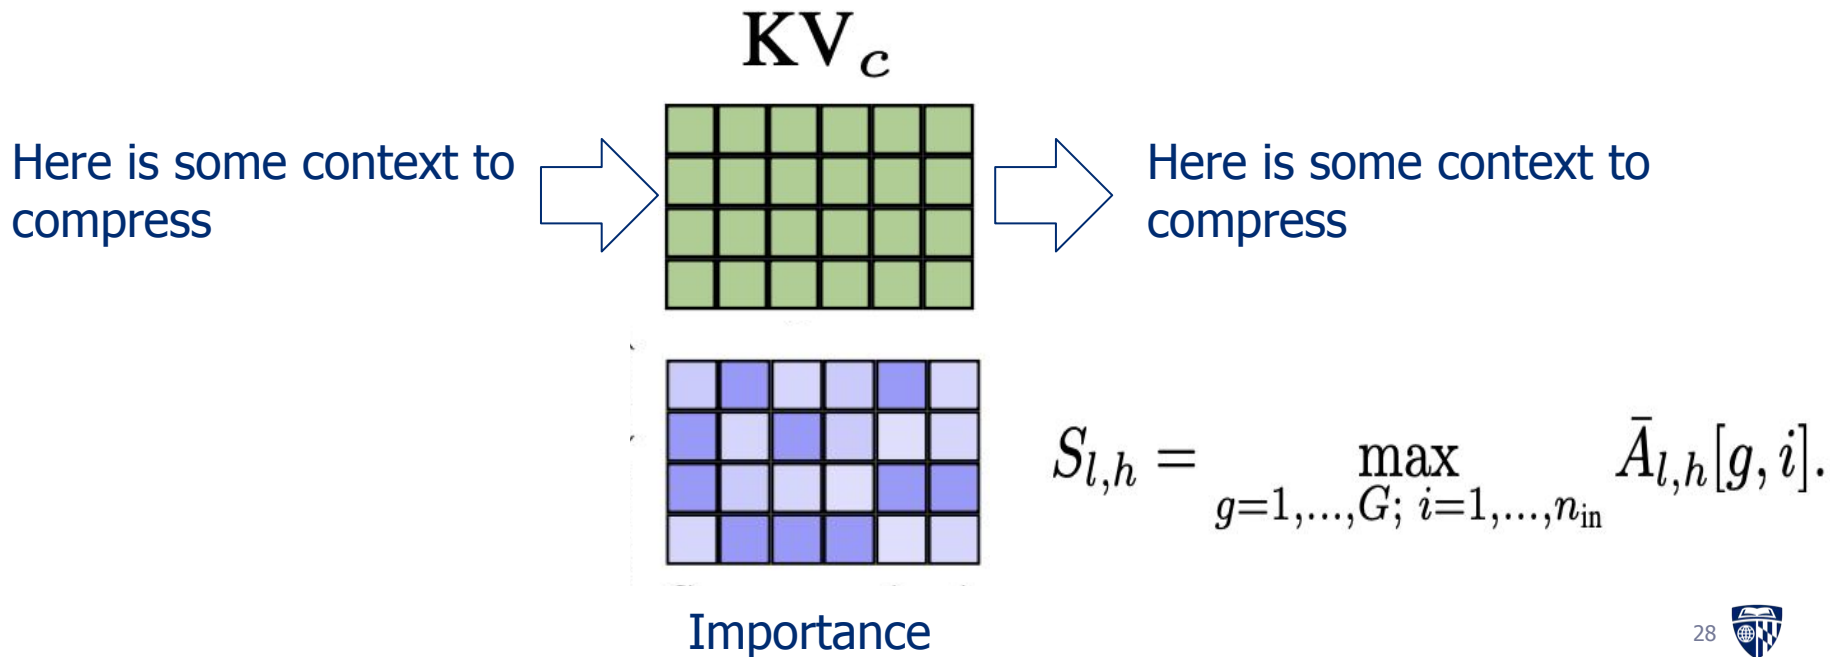
Jang-Hyun Kim, Jinuk Kim, Sangwoo Kwon,  
Jae W. Lee, Sangdoo Yun, Hyun Oh Song

# KVzip




# KVzip: How to evict?

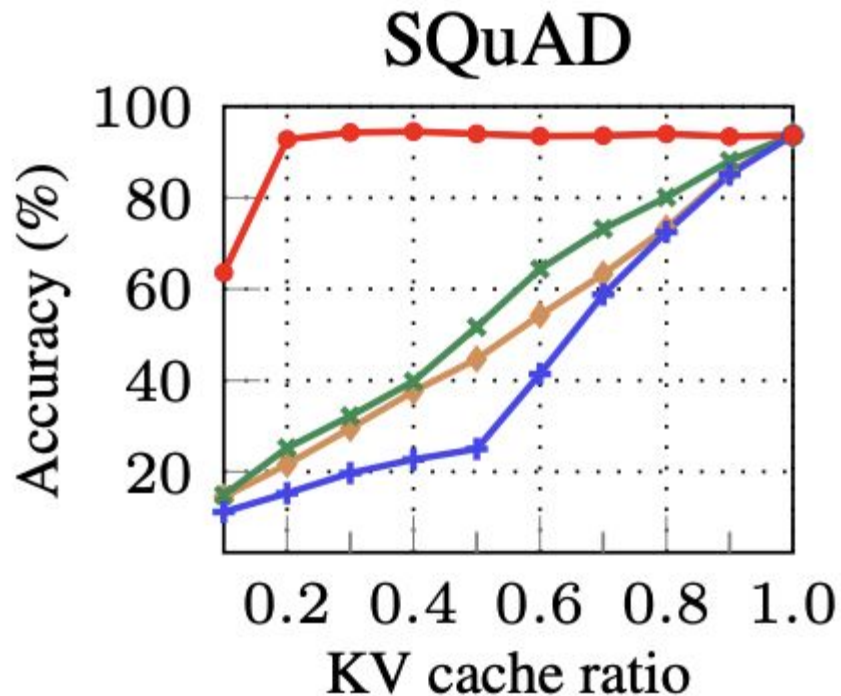
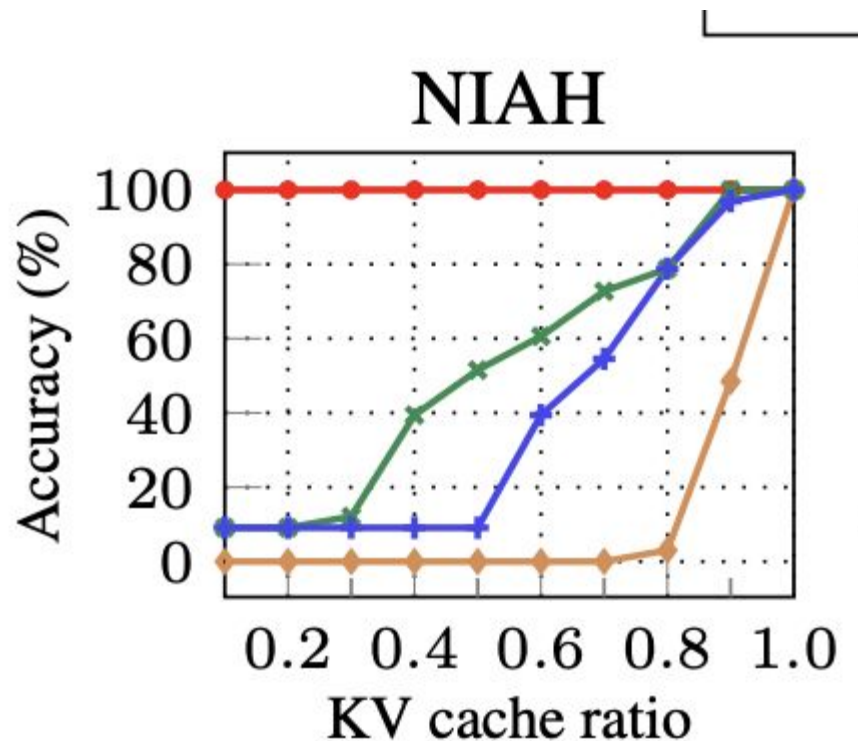
Evict tokens that don't contribute to reconstructing the context



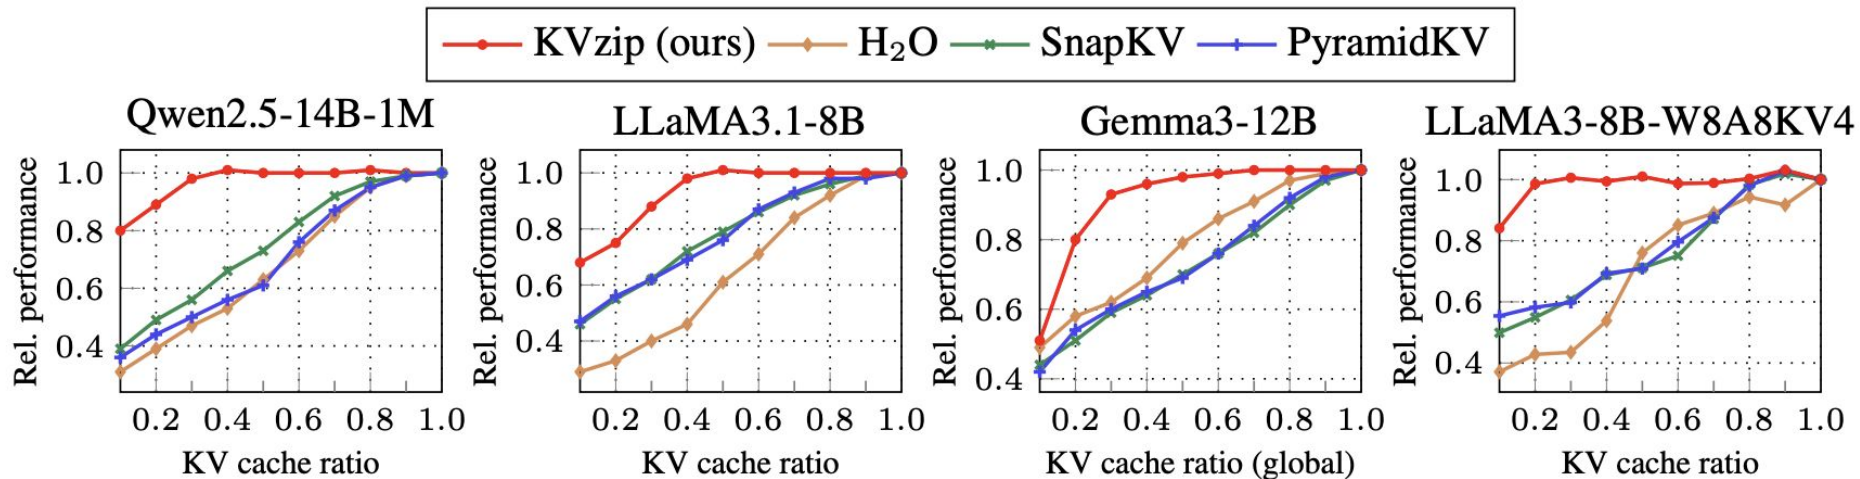
# KVzip: Attention Sparsity



# KVzip: Performance



# KVzip: Model Comparison



# Conclusions & Critiques





# Conclusions

- Main takeaway: PyramidKV mirrors attention naturally funneling through layers
- Performance:
  - Preserves accuracy while using only 12% of KV cache
  - Preserves long-context understanding ability
- Efficiency: Up to 90% GPU-memory reduction, minimal runtime overhead:

cache size	Memory	Compression Ratio	QMSum	TREC	TriviaQA	PCount	PRe	Lcc
512	428M	6.3%	22.80	71.50	90.61	5.91	69.50	58.16
1024	856M	12.5%	22.55	71.50	90.61	5.91	69.50	58.16
2048	1712M	25.0%	22.55	72.00	90.56	5.58	69.25	56.79
Full	6848M	100.0%	23.30	73.00	90.56	5.22	69.25	58.76

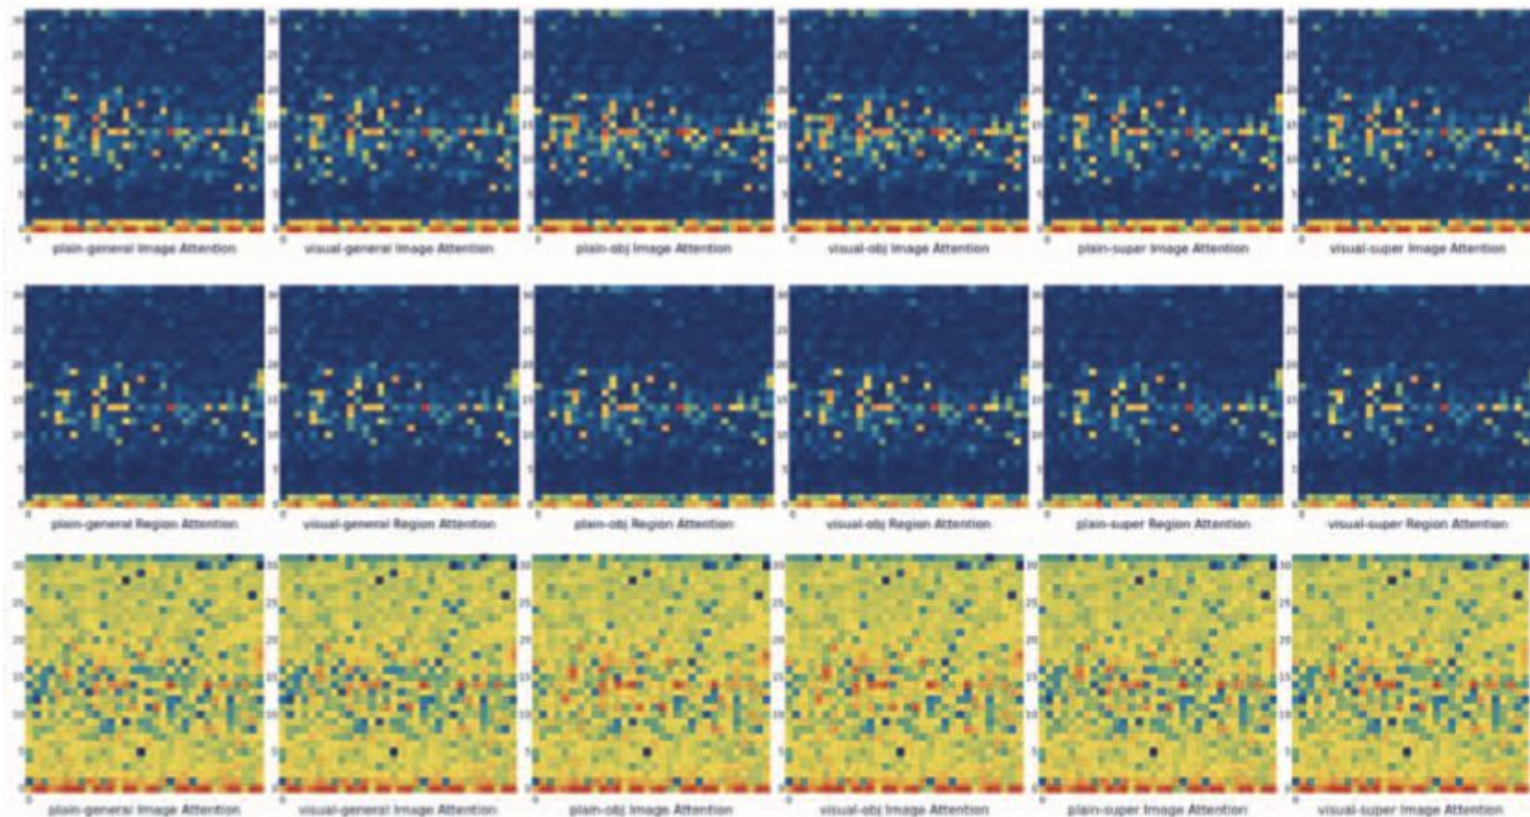
# Critiques

---

- My thoughts: Very impressive results, lots of memory reduction with a simple implementation.
- Limits:
  - Evaluated only on 3 English models.
    - No multilingual testing.
  - The pyramid could fail on tasks with different attention shapes.
    - The observed attention phenomenon was only for multi-doc QA tasks.



# Critiques: Attention Sparsity, Do you really want that?



# Thank You!

**Any Questions?**

Gabriel Pernell, Alexander Martin

