

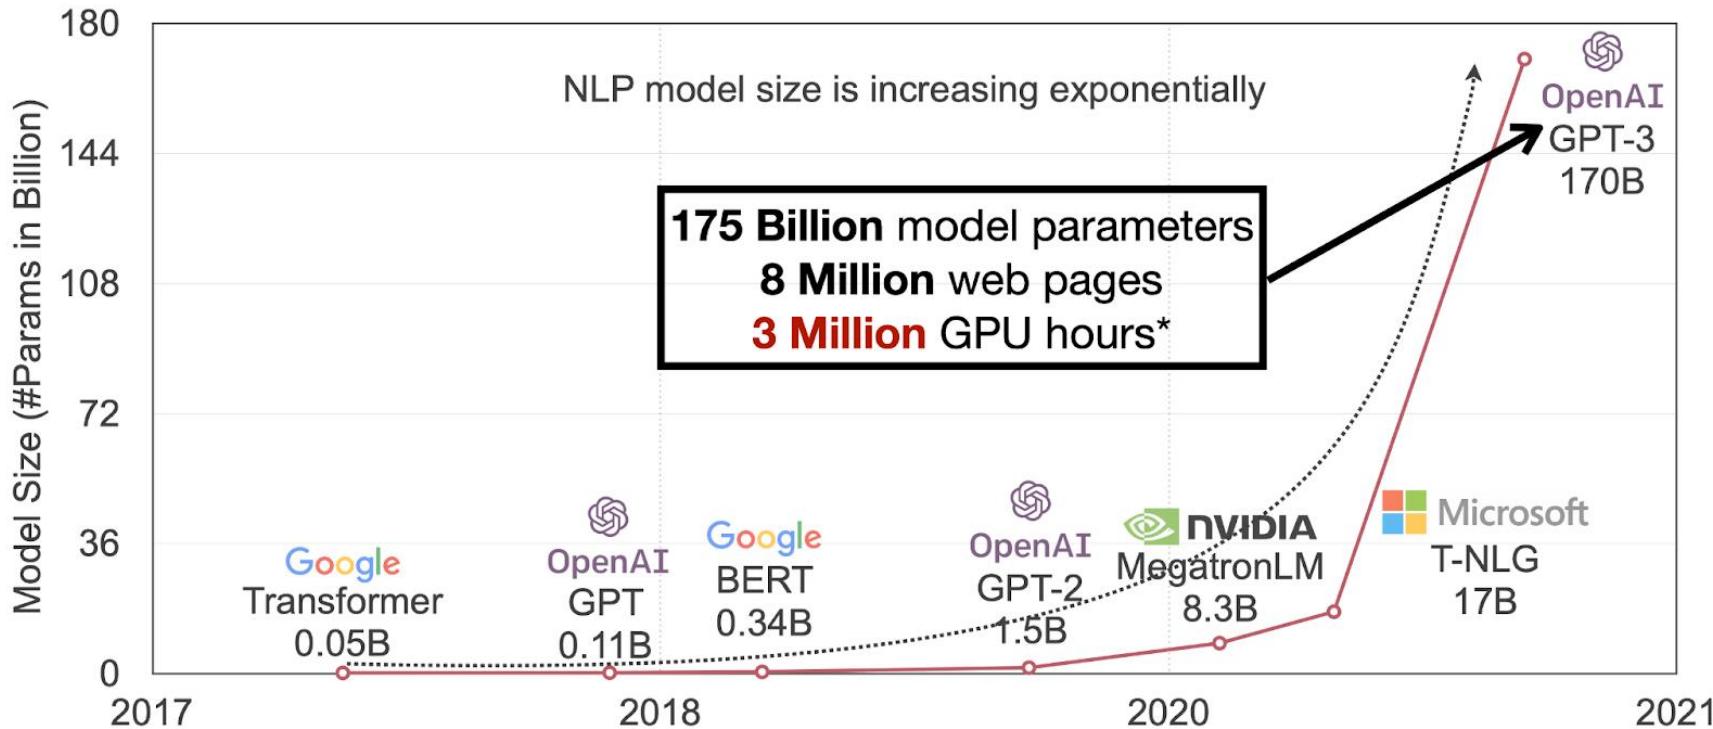


# Model Efficiency

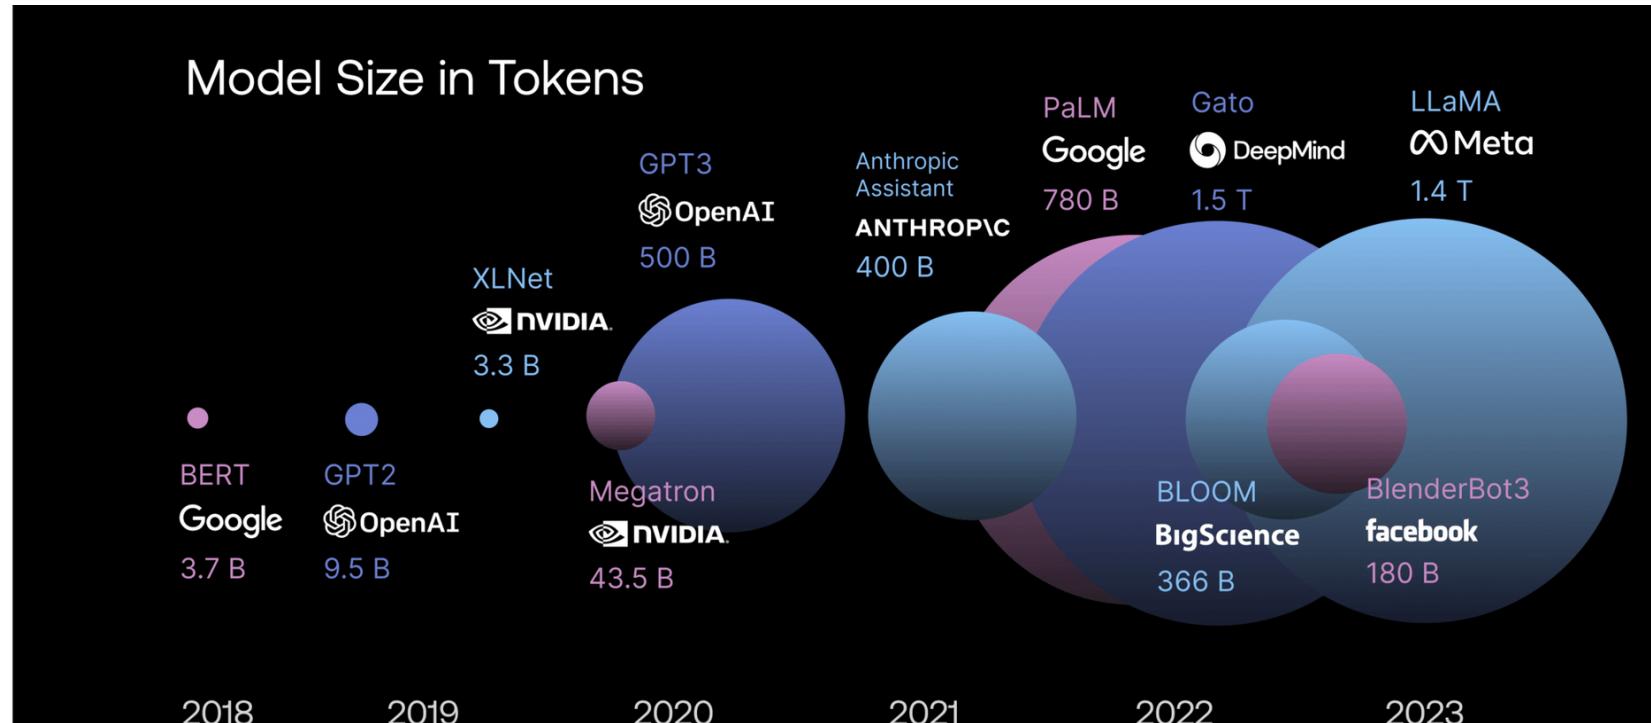
CSCI 601-471/671 (NLP: Self-Supervised Models)

<https://self-supervised.cs.jhu.edu/sp2025/>

# Our models are getting larger!



# And consumes a lot of data!



# Motivation

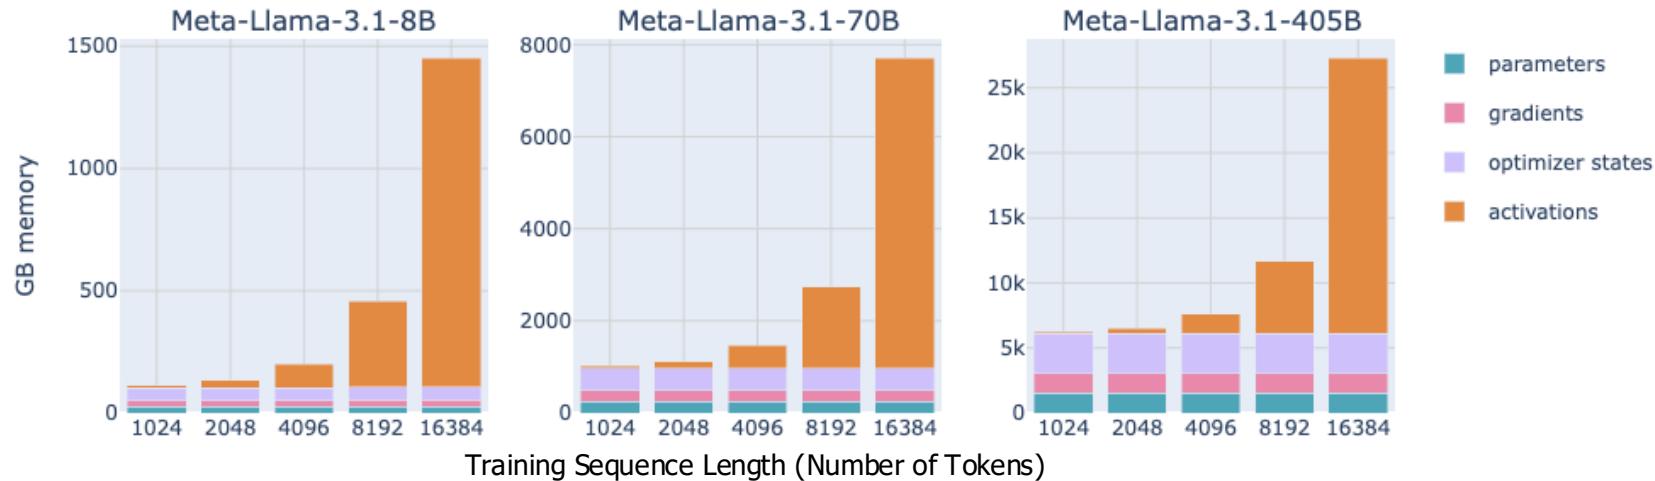
---

How much GPU memory (**at least**) do we need to perform inference/training?  
(batch size=1, ignoring the KV cache)

Model Size (Llama 3 Arch)	Inference Memory (~2x model size)	Training Memory (~7x model size)
8B	16GB	60GB
70B	140GB	500GB
405B	810GB	3.25TB

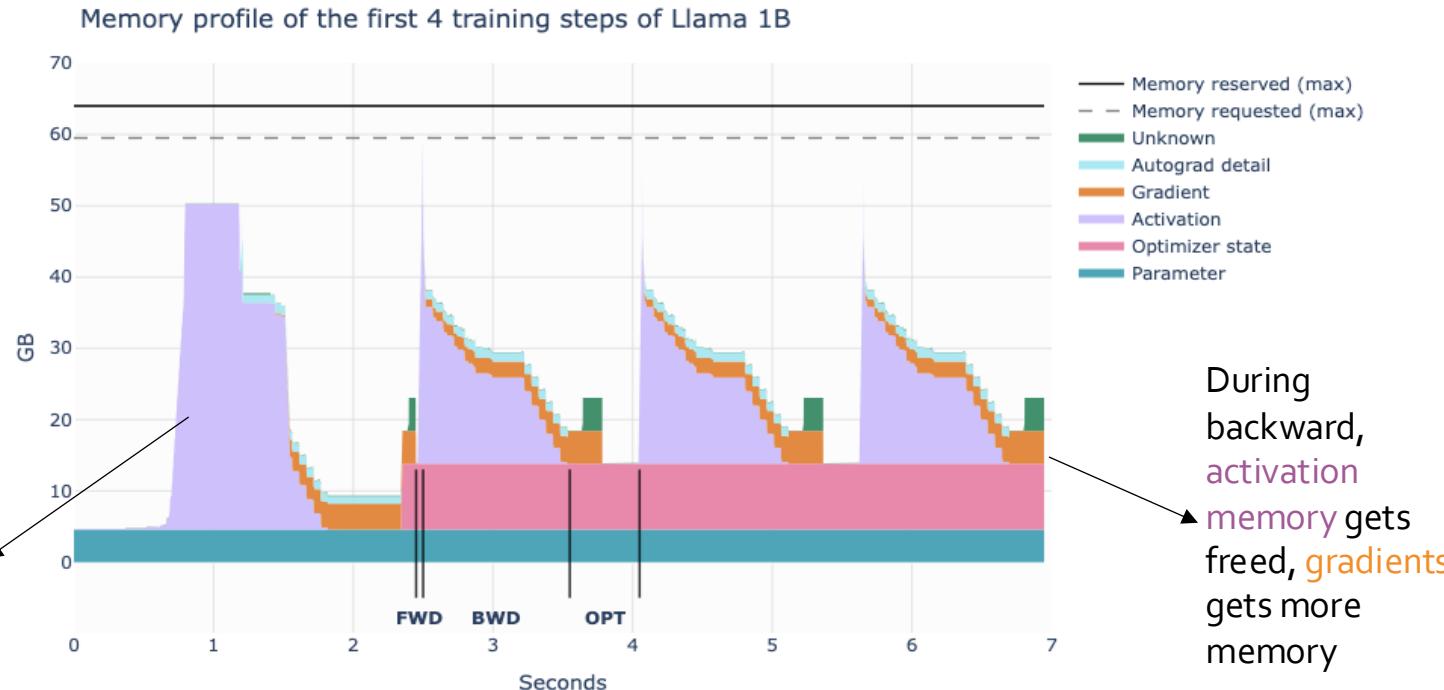
# Where did all the memory go?

Longer sequences require much more memory in training!



Source: <https://nanotron-ultrascale-playbook.static.hf.space/dist/index.html>

# Memory consumption is not static



Source: <https://nanotron-ultrascale-playbook.static.hf.space/dist/index.html>

# Model Efficiency: Topics

---

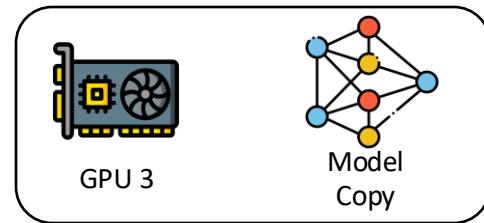
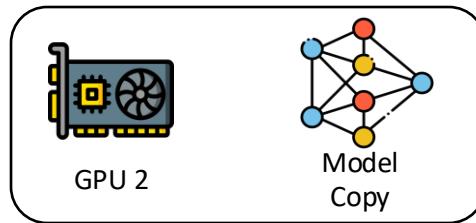
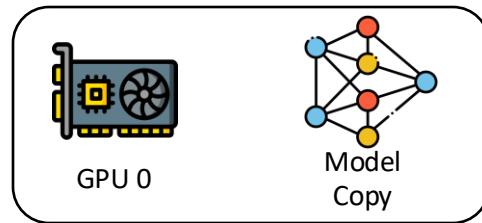
1. Distributed Training
2. Quantization (Post Training Quantization)
3. Distillation

**Chapter goal:** Getting comfortable with various mathematical and systems foundations for efficient deployment of LLMs.

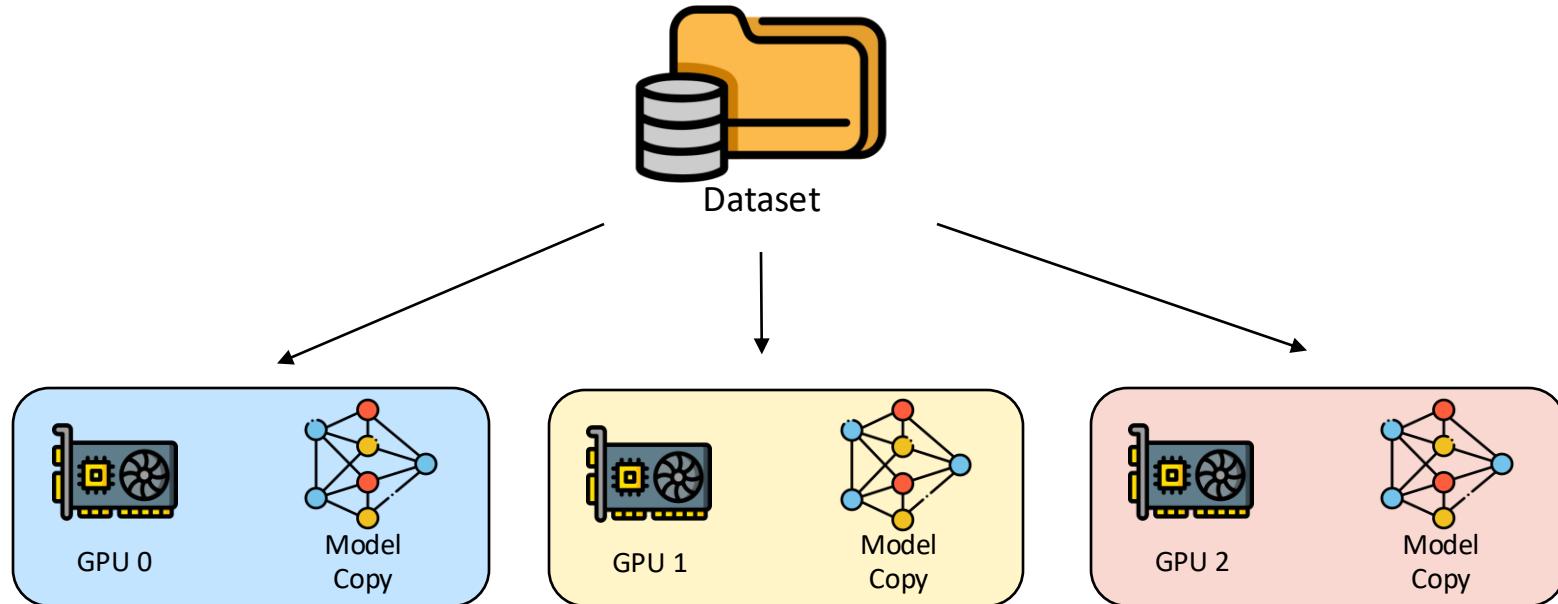
# Distributed Training

# Distributed Training

1. Naïve Data Parallelism
2. Sharding Optimizer States (ZeRO, FSDP)
3. Model Parallelism (Tensor Parallelism, Pipeline Parallelism)



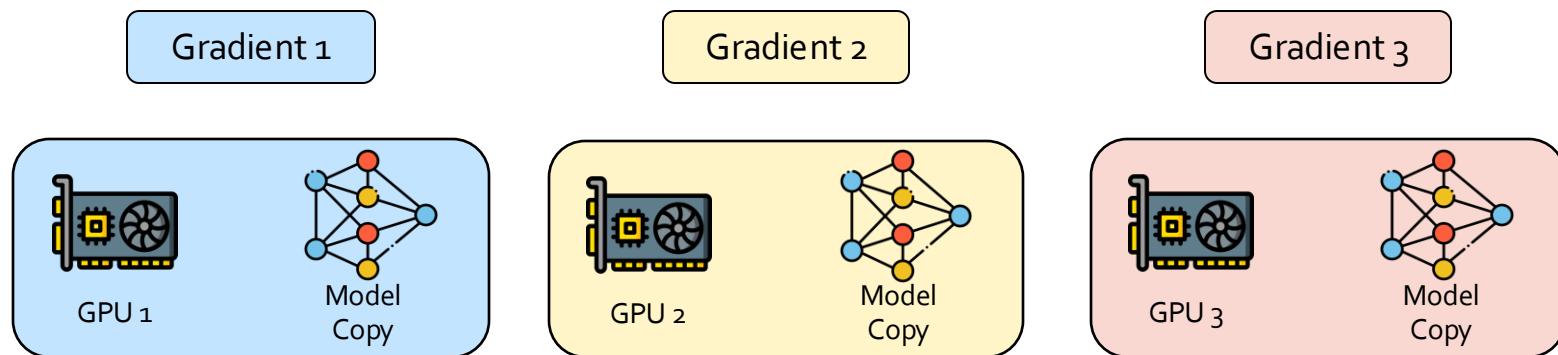
# Naïve Data Parallelism



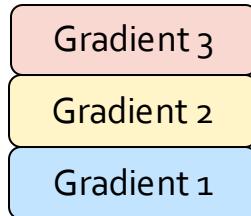
First, we want to shard the dataset and feed them into different GPUs  
How do we update the parameters?

# Naïve Data Parallelism

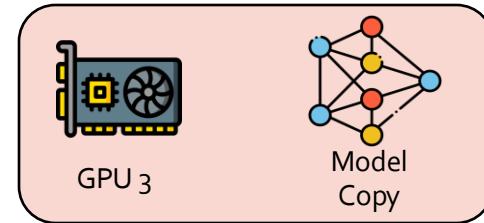
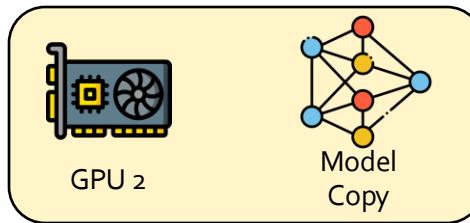
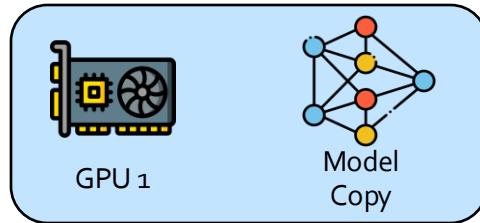
Each GPU compute gradient with a single shard of data



# Naïve Data Parallelism

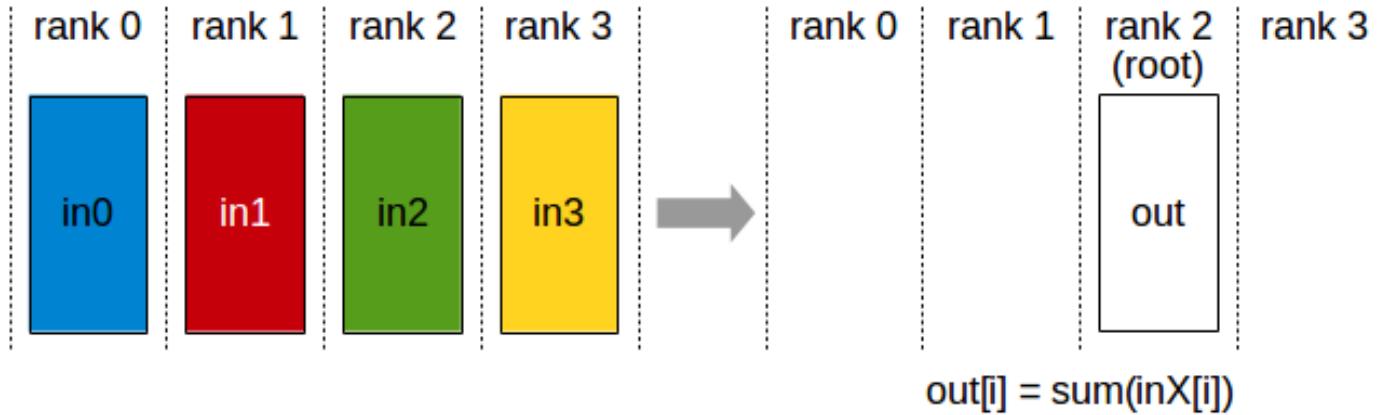


One GPU accumulates the gradients  
**(*reduce* in `torch.distributed`)**

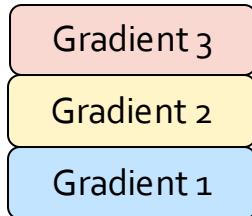


# NCCL Operations: Reduce

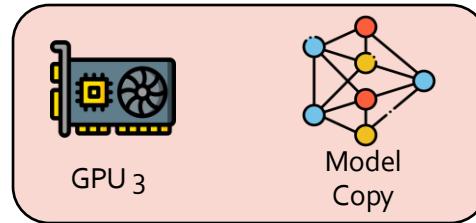
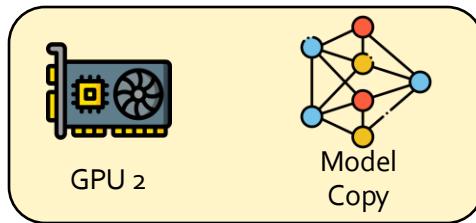
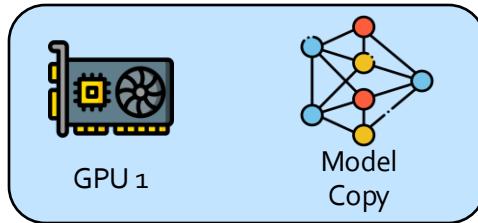
- Nvidia Collective Communications Library (NCCL) - A library developed to provide inter-GPU communications primitives (operations)
- Reduce: \*Sums\* over all \*tensors\* and stores it in a root GPU



# Naïve Data Parallelism

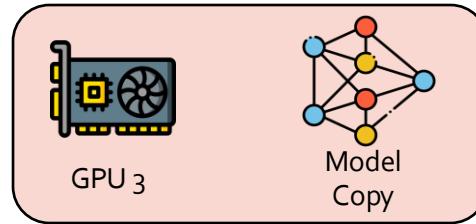
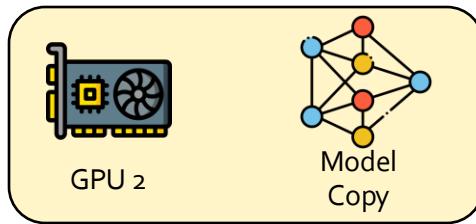
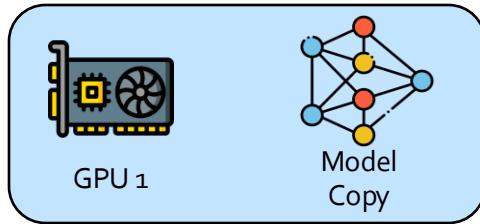
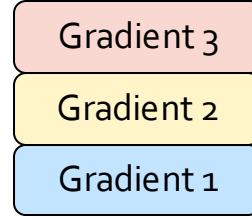
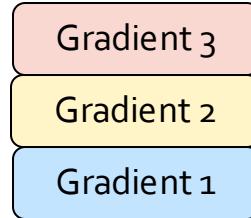
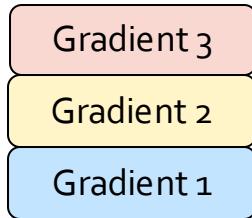


One GPU accumulates the gradients  
**(*reduce* in `torch.distributed`)**



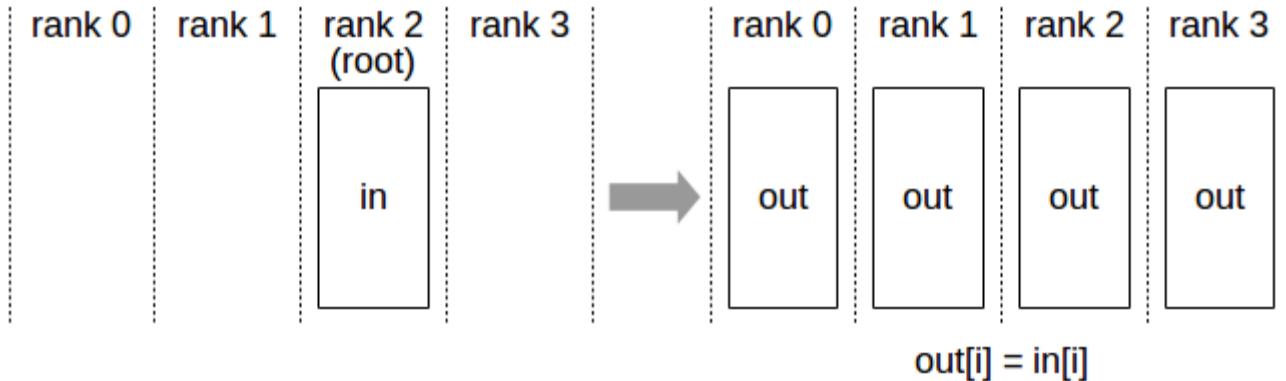
# Naïve Data Parallelism

And send the accumulated gradient to all other GPUs (**broadcast** in `torch.distributed`)



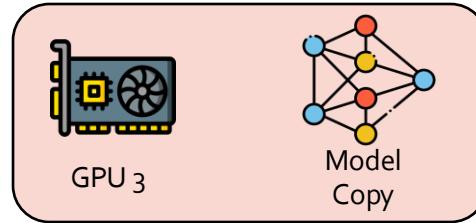
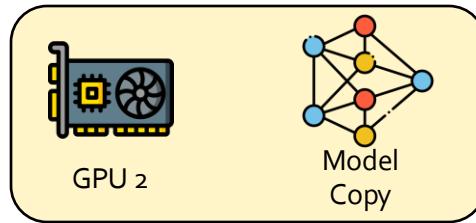
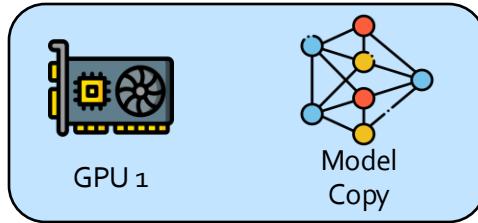
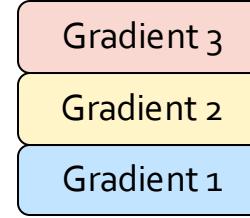
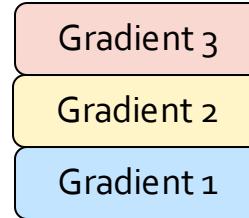
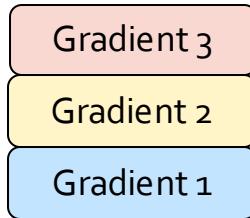
# NCCL Operations: Broadcast

- Broadcast: Duplicates one tensor to all GPUs



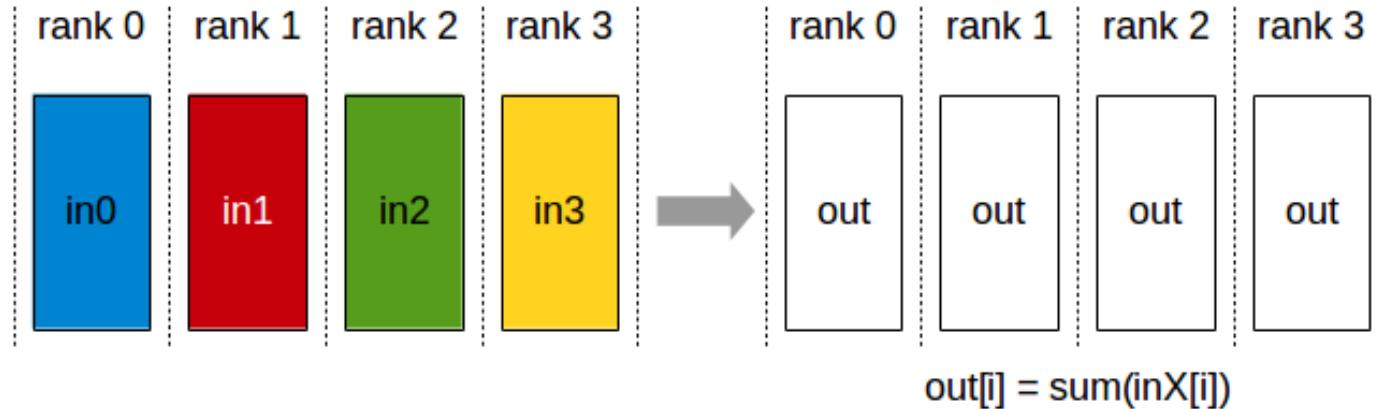
# Naïve Data Parallelism

Accumulate gradients across all GPUs and perform gradient updates  
(`all_reduce` in `torch.distributed`)



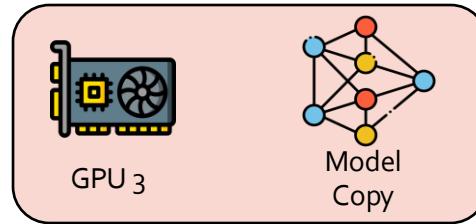
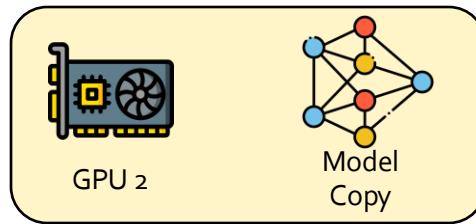
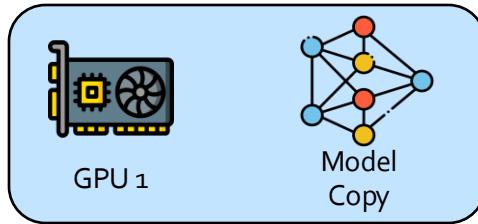
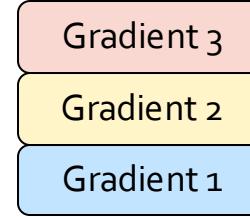
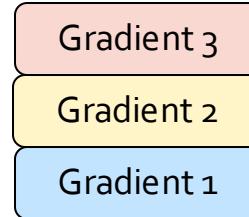
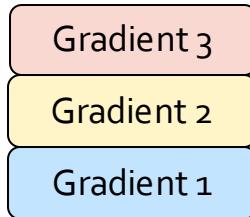
# NCCL Operations: All Reduce

- All Reduce = Reduce + Broadcast  
= Sum over input tensors, then duplicate it to all GPUs



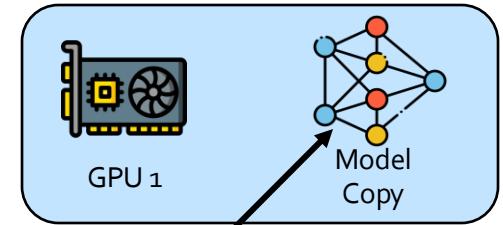
# Naïve Data Parallelism

Accumulate gradients across all GPUs and perform gradient updates  
(`all_reduce` in `torch.distributed`)



# What is wrong with Naïve DP

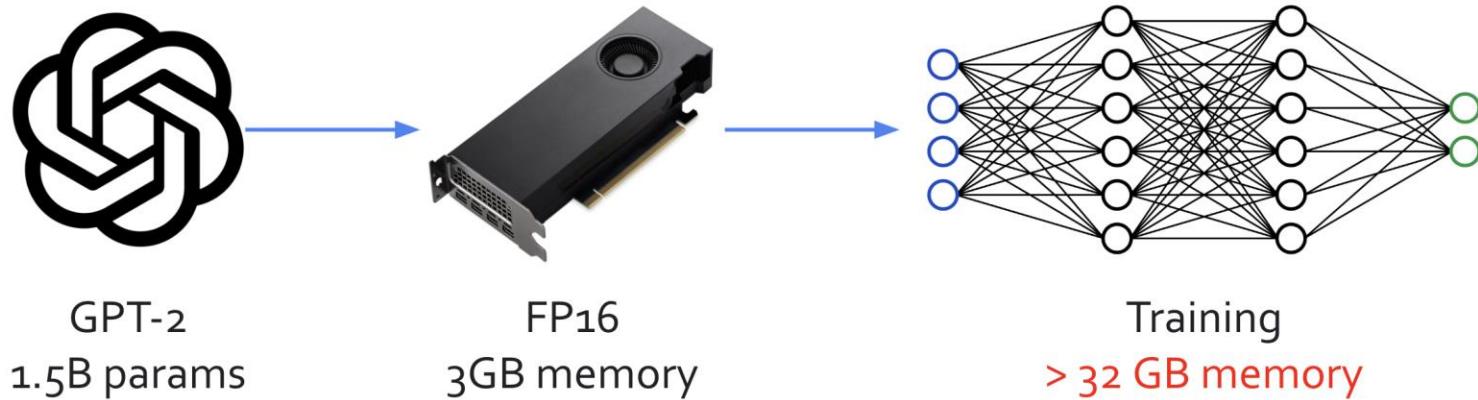
- Consumes too much memory in each GPU!



We need to store **5** copies of weights,  
which occupies **16** bytes per param

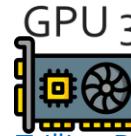
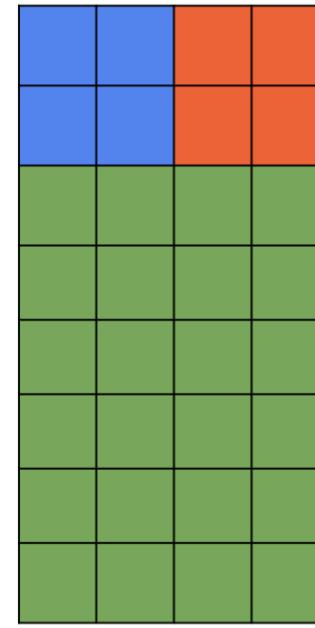
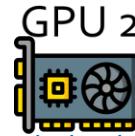
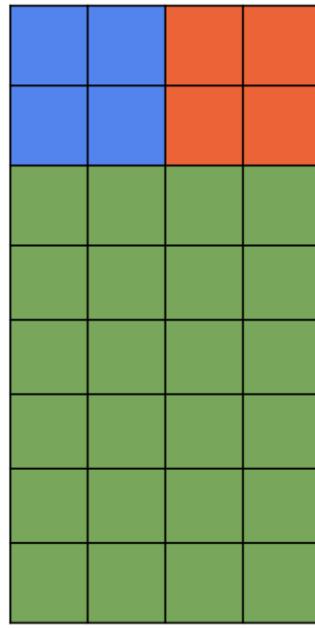
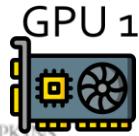
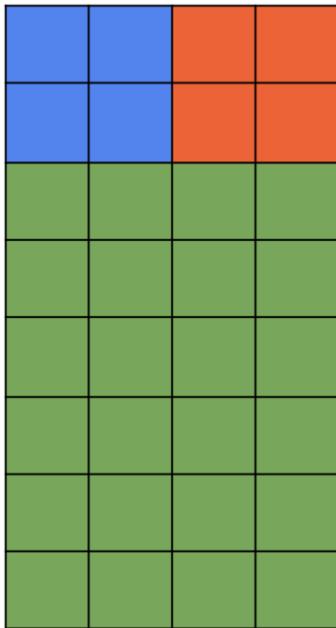
- 2 bytes for FP/BF16 model params
- 2 bytes for FP/BF16 gradients
- 4 bytes for FP32 master weights  
(the thing you accumulate into in SGD, used in mixed precision training)
- 4 bytes for FP32 Adam first order estimates
- 4 bytes for FP32 Adam second order estimates

# What is wrong with Naïve DP



Most of the memory are occupied by optimizer states.  
Some are also occupied by *residual states*: activations, buffers and fragmented memory

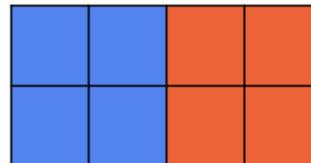
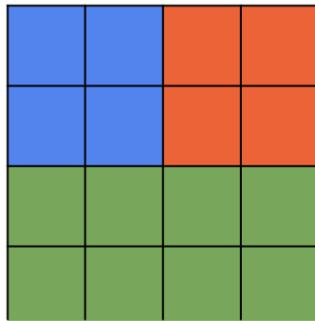
# Naïve DP – Requires too much memory!



Parameters  
Gradients  
Optimizer States

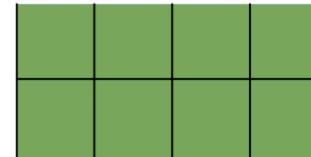
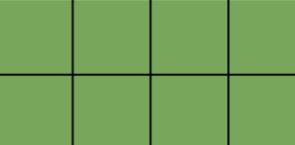
Memory/GPU for a 7.5B model:  
 $7.5\text{B} * 16 \text{ bytes} = 120 \text{ GB!}$

# ZeRO Stage 1: Sharding Optimizer States

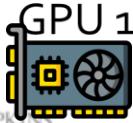


Legend:

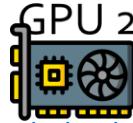
- Blue square: Parameters
- Orange square: Gradients
- Green square: Optimizer States



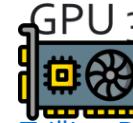
Memory/GPU for a 7.5B model:  
 $7.5\text{B} * (2+2+4) \text{ bytes} = 60 \text{ GB!}$



GPU 1



GPU 2

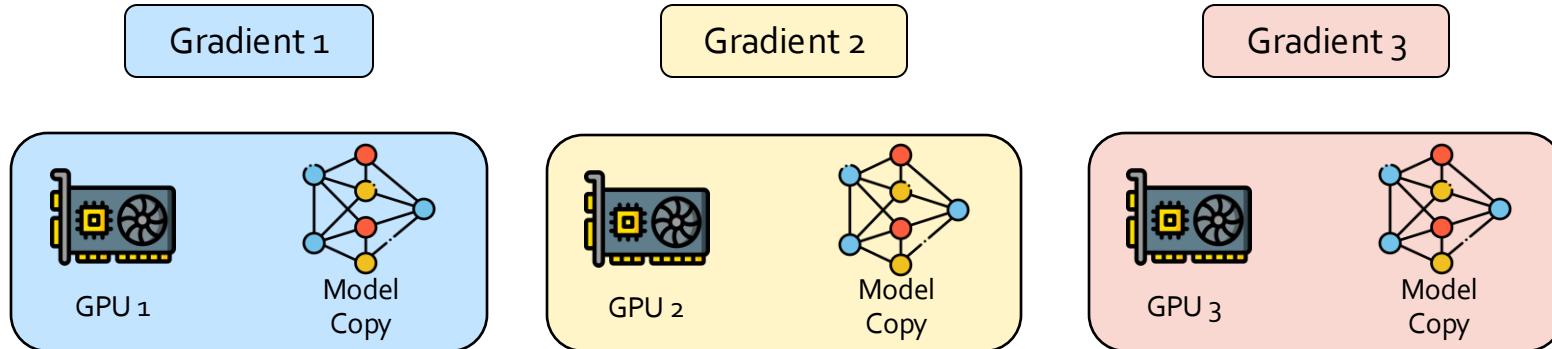


GPU 3

# ZeRO Stage 1: How it works

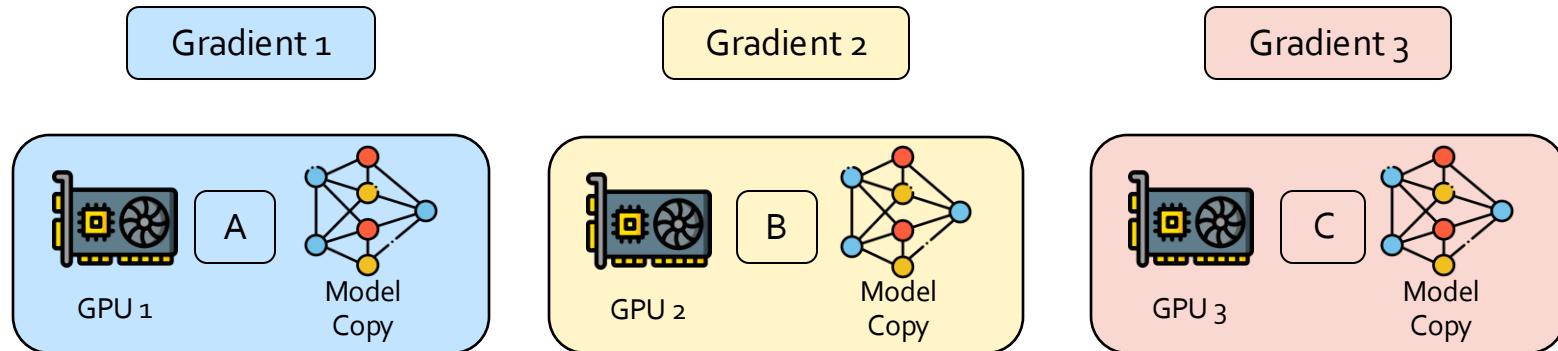
Update Parameters

Each GPU compute gradient with a single shard of data  
(The same as naïve DP)



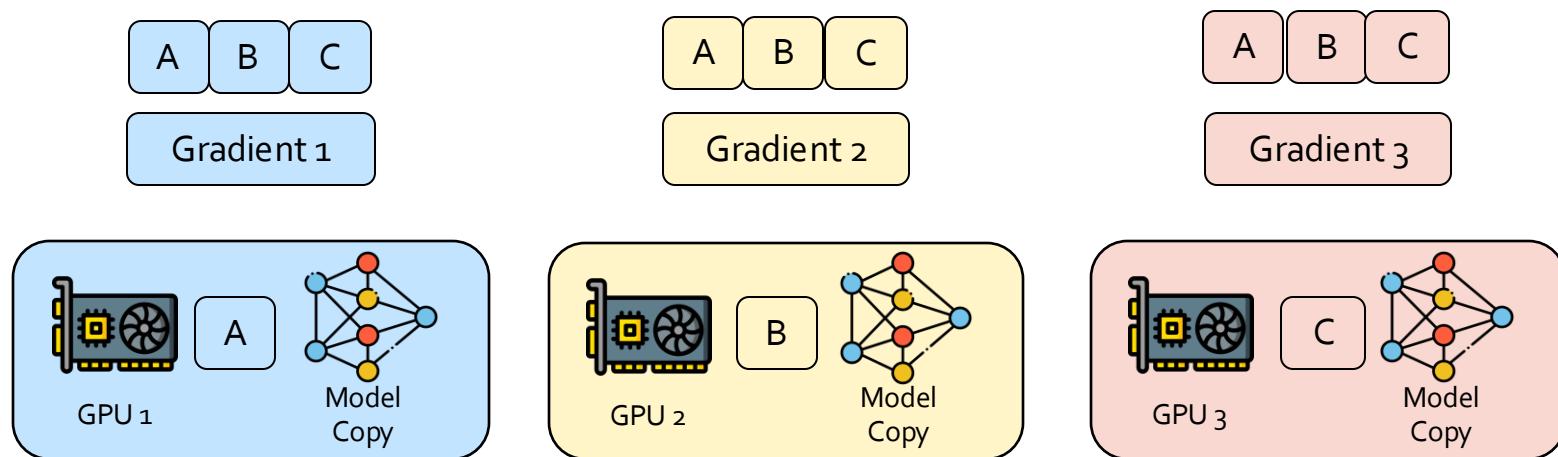
# ZeRO Stage 1: How it works

Assuming that  
GPU1 stores parameter states for parameters A,  
GPU2 stores states for params B,  
GPU3 stores states for params C



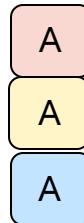
# ZeRO Stage 1: How it works

Split / shard the gradients into 3 parts!



# ZeRO Stage 1: How it works

Each GPU accumulates gradients of the params whose optimizer states the GPU is storing (`reduce_scatter` in `torch.distributed`)



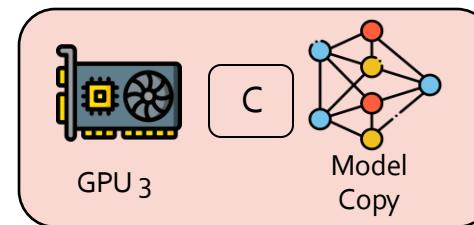
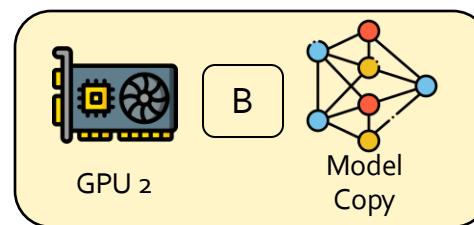
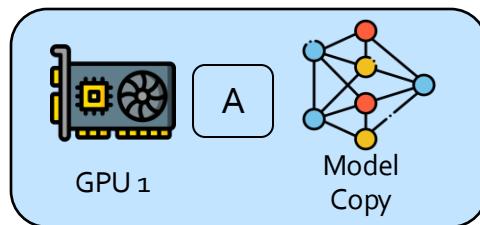
Gradient 1



Gradient 2

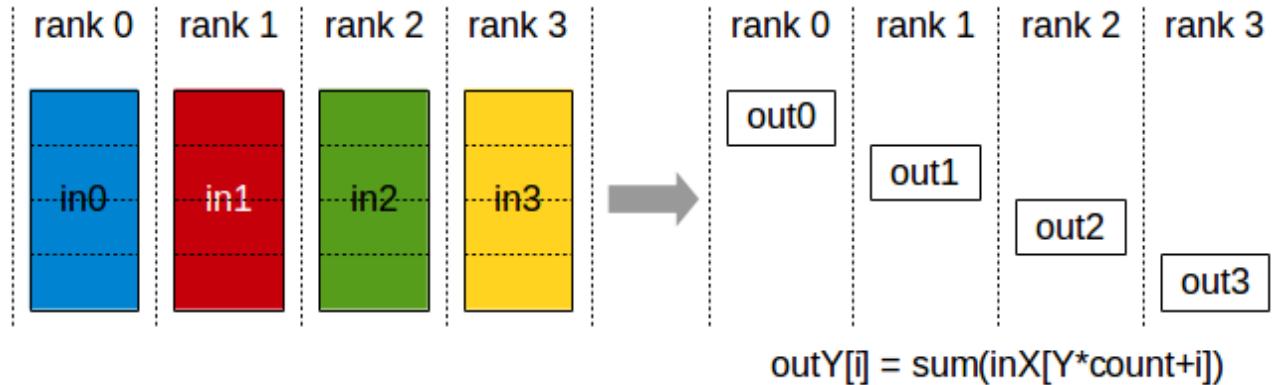


Gradient 3



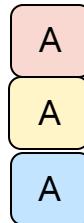
# NCCL Operations: Reduce Scatter

- reduce\_scatter: each GPU stores the sum of a shard of the input.
- all\_reduce: one GPU stores the sum over all the input.



# ZeRO Stage 1: How it works

Each GPU accumulates gradients of the params whose optimizer states the GPU is storing (`reduce_scatter` in `torch.distributed`)



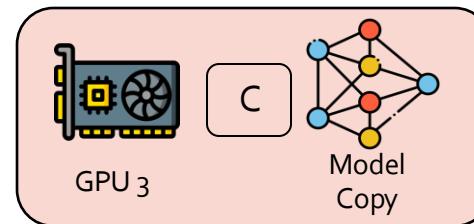
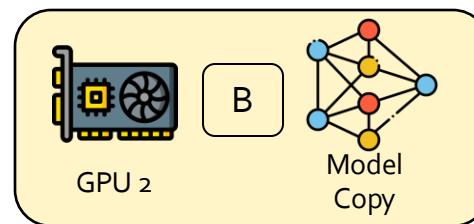
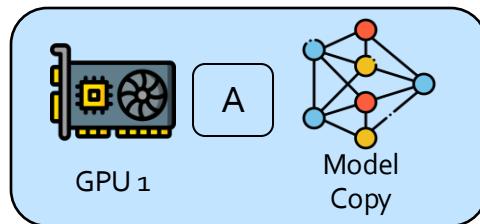
Gradient 1



Gradient 2

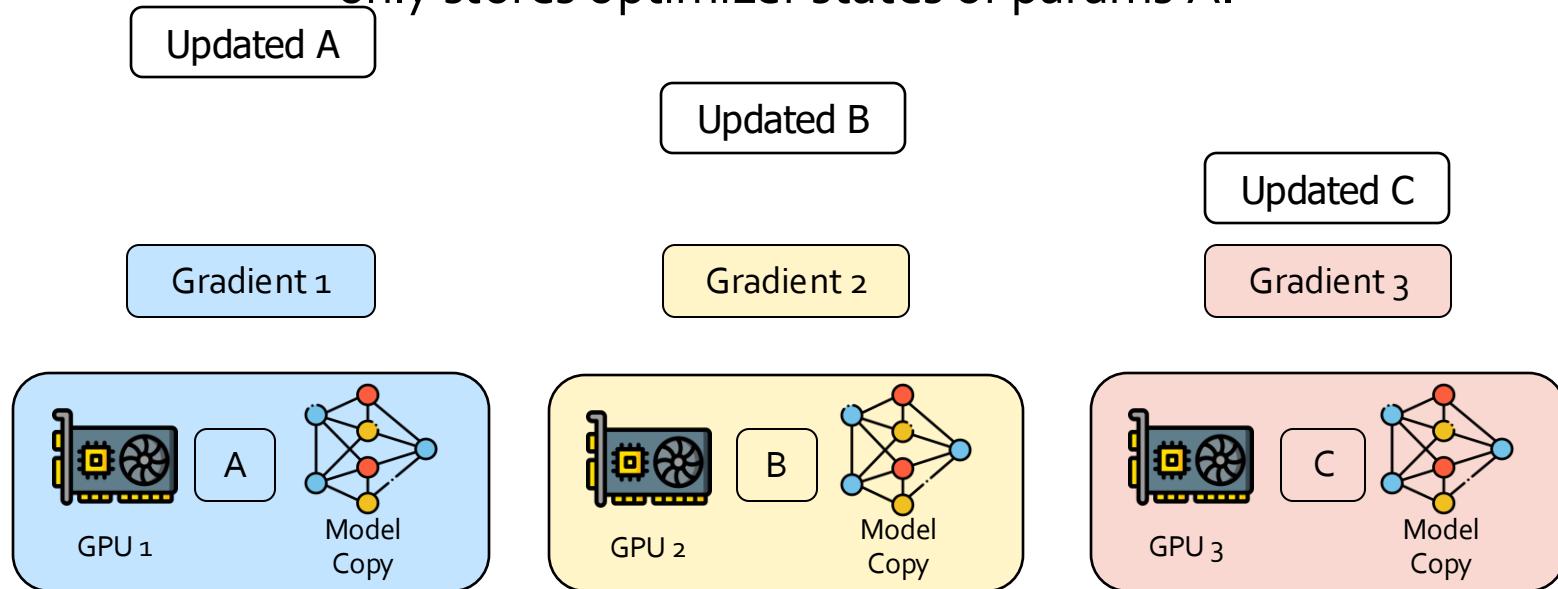


Gradient 3



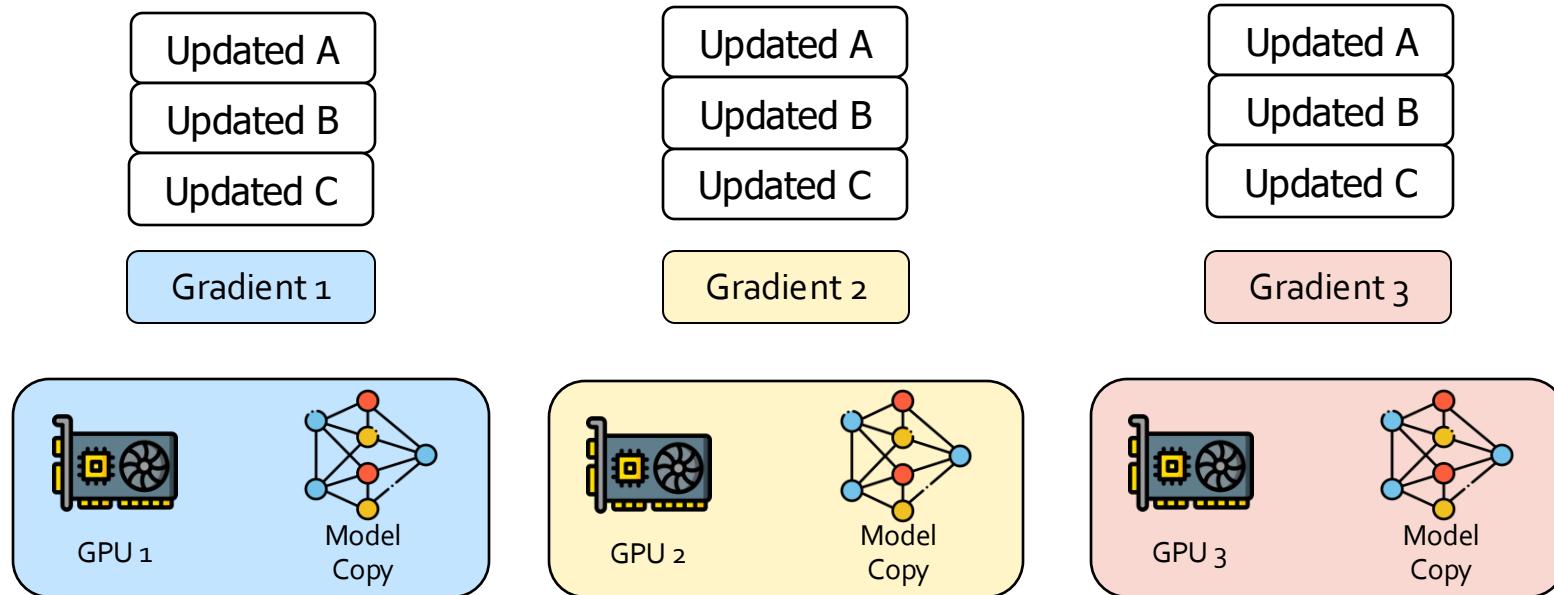
# ZeRO Stage 1: How it works

GPU1 : update params A; GPU2: Updates Params B; GPU3: updates params C. GPU1 can only update params A since it only stores optimizer states of params A.



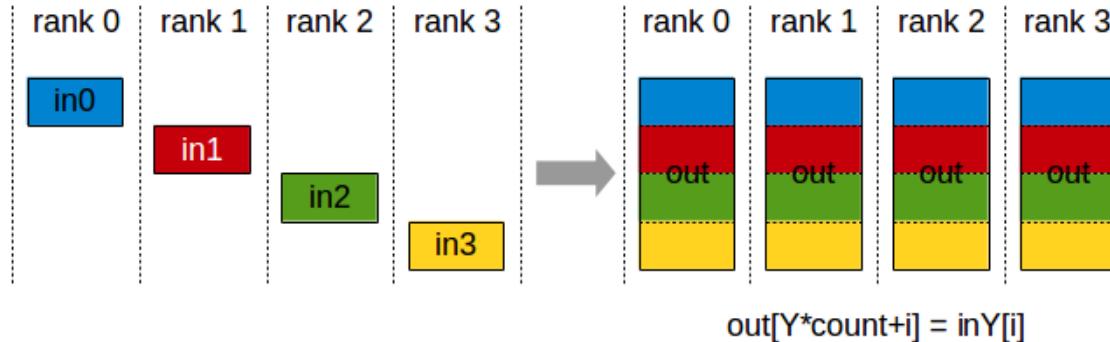
# ZeRO Stage 1: How it works

Each GPU sends updated params to every other GPU.  
Finishing `optimizer.step()`. (**all\_gather** in `torch.distributed`)



# Quiz: NCCL Operations: All Gather

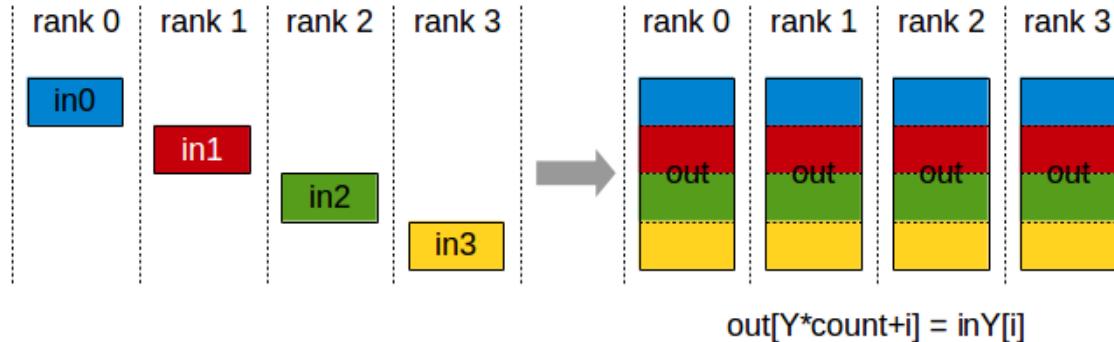
- all\_gather: every GPU performs a \_\_\_?\_\_\_ operation in parallel.



- A.Reduce      B. Broadcast      C. Reduce\_scatter

# NCCL Operations: All Gather

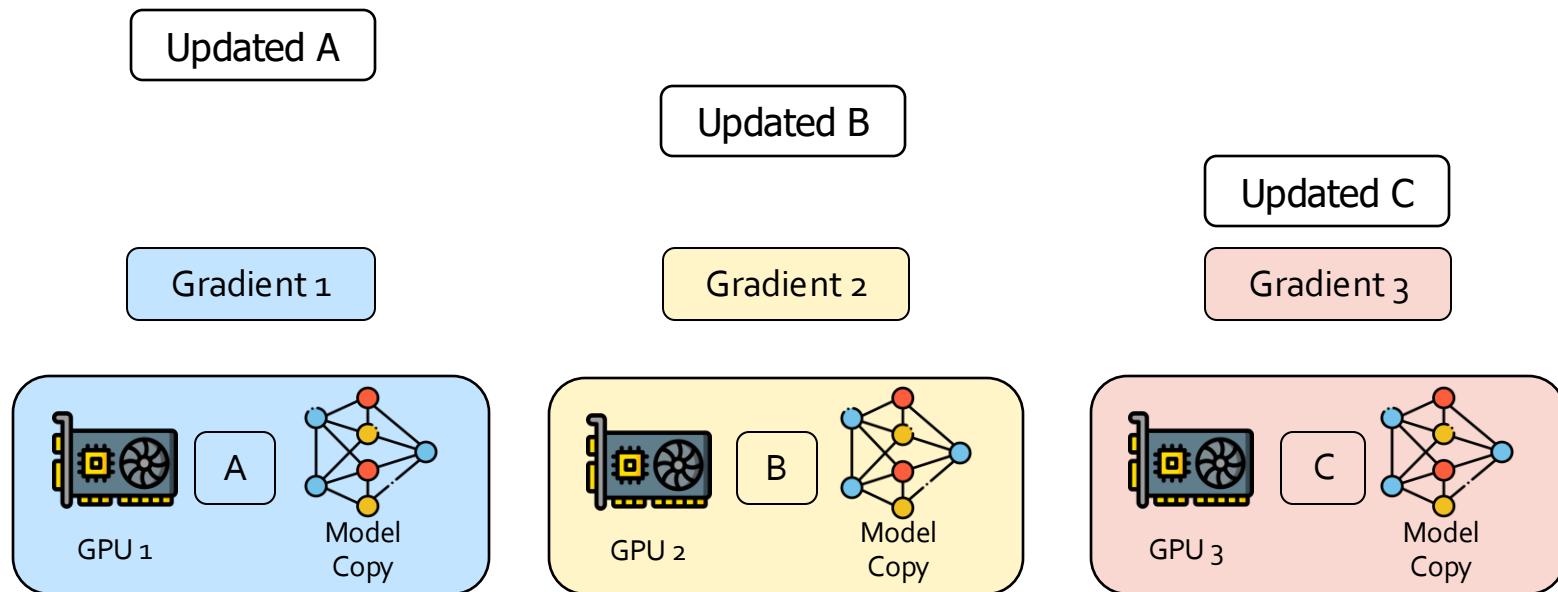
- all\_gather: every GPU performs a broadcast operation in parallel.



- A.Reduce      B. Broadcast      C. Reduce\_scatter

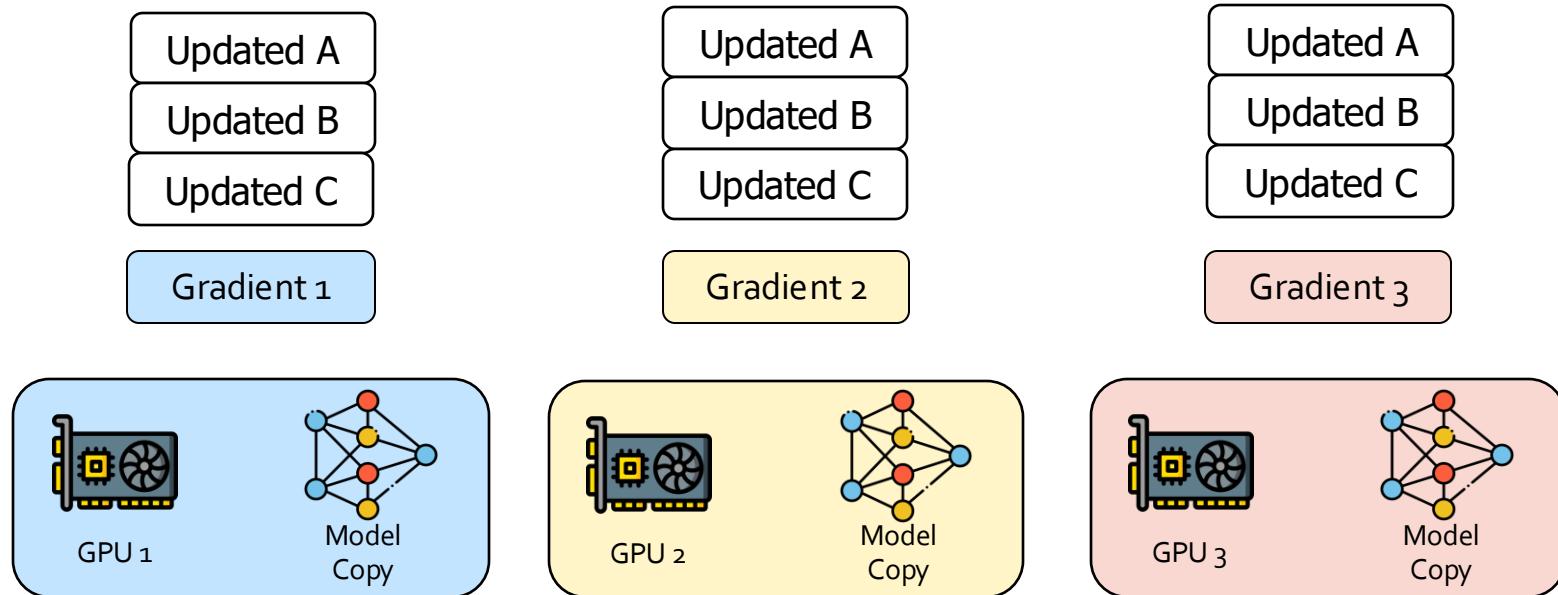
# ZeRO Stage 1: How it works

Before all\_gather



# ZeRO Stage 1: How it works

After all\_gather, every GPU has a updated copy of the model



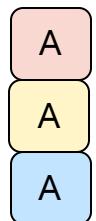
# Summary: ZeRO 1

---

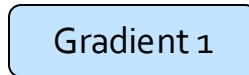
- `reduce_scatter` on the gradients: splitting the gradients into different GPUs
- Each GPU individually perform gradient updates
- `all_gather` on updated parameters
- Basically free! (Compared to Naïve Data Parallelism)

# ZeRO Stage 1: How it works

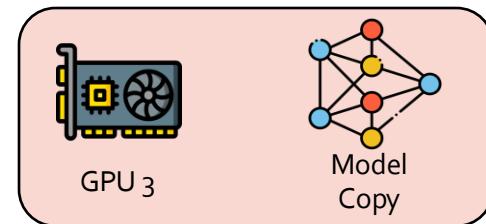
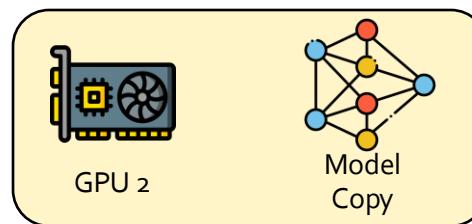
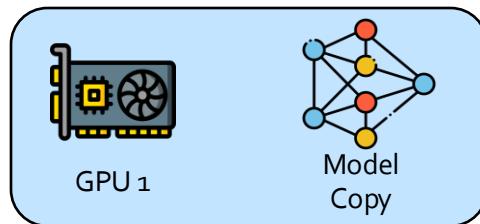
Notice: Aside from the forward pass, GPU 1 only needs gradients A, but in fact it stores A and B and C



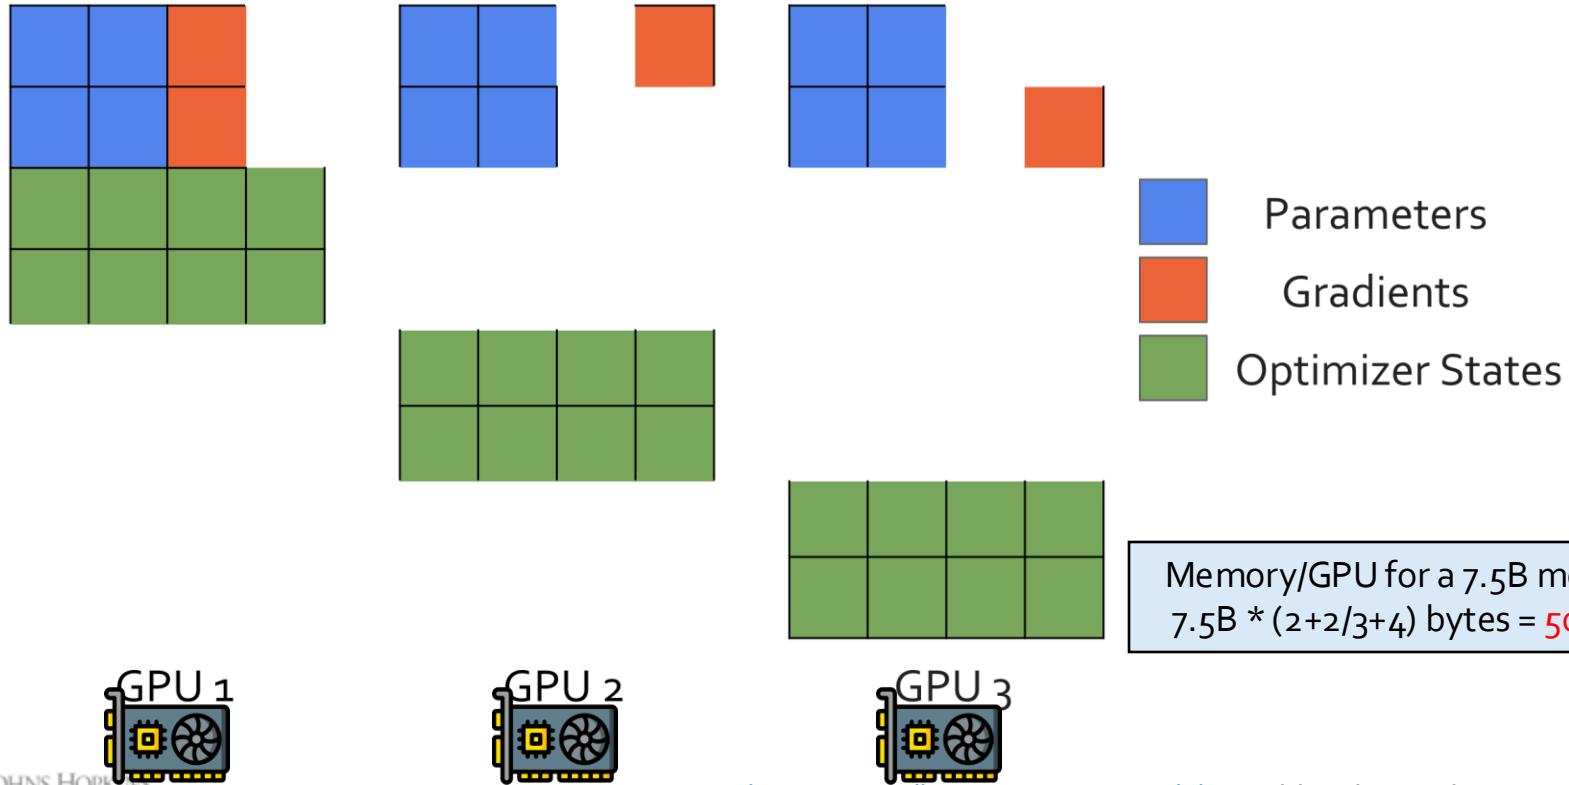
You only need these



Hey GPU1, you don't need this (can be large)

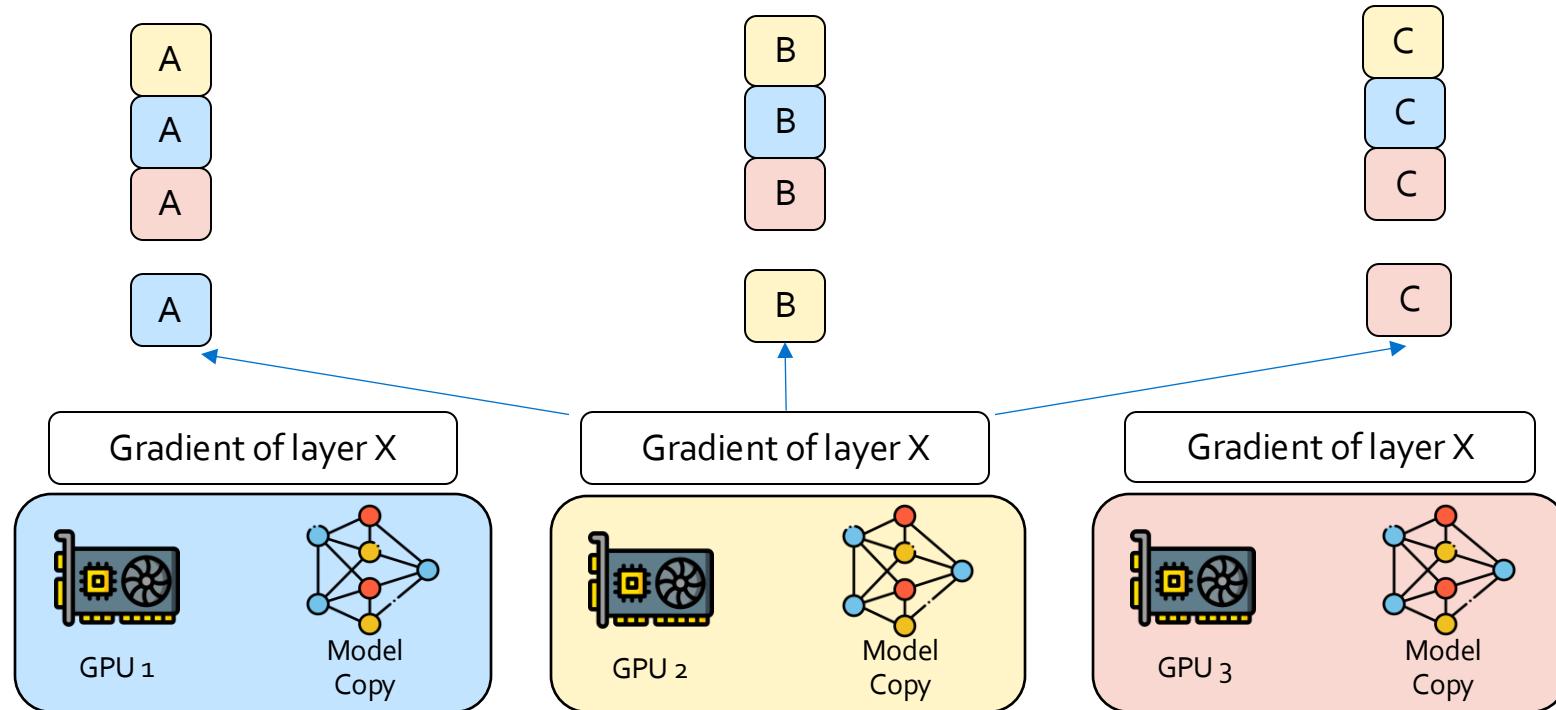


# ZeRO Stage 2: Sharding Gradients

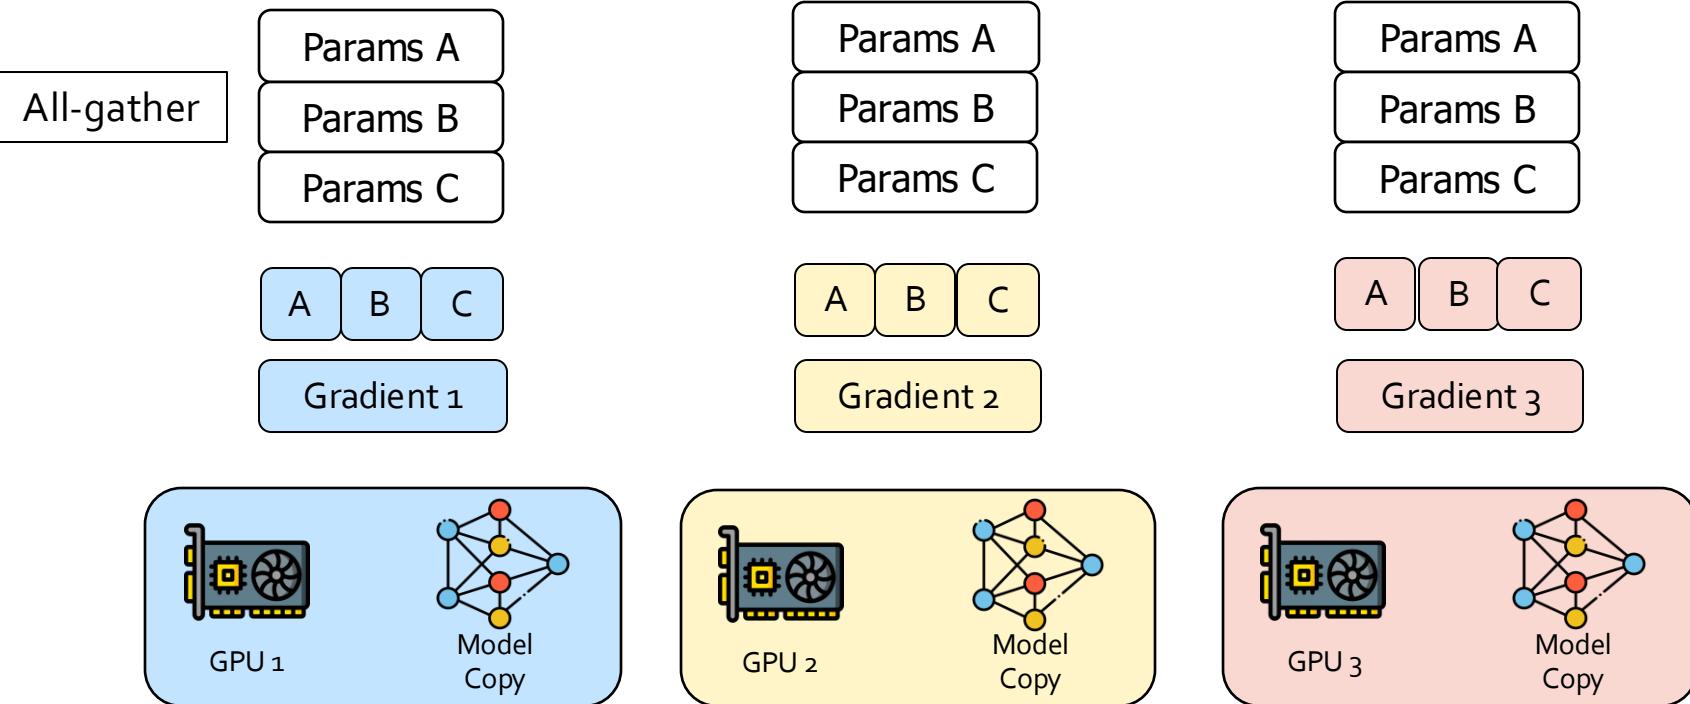


# ZeRO Stage 2: How it works

Splitting the gradient of a **single layer** during backprop, then immediately shard it!



# ZeRO Stage 2: How it works

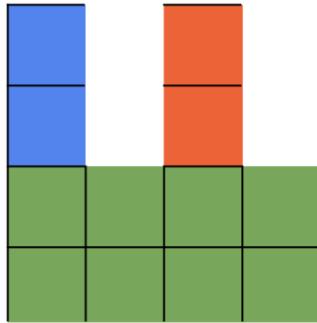


# Summary: ZeRO 12

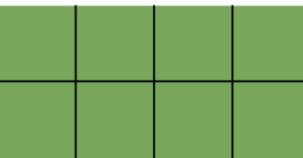
---

- ~~reduce\_scatter~~ on the gradients: splitting the gradients into different GPUs
  - Calculate gradients layer by layer and perform ~~reduce\_scatter~~, once layer is done, free the gradient
- 
- Each GPU individually perform gradient updates
  - ~~all\_gather~~ on updated parameters
  - Almost free!

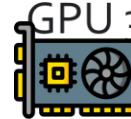
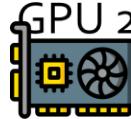
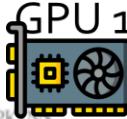
# ZeRO-3 (aka FSDP): Shard Everything!



Parameters  
 Gradients  
 Optimizer States

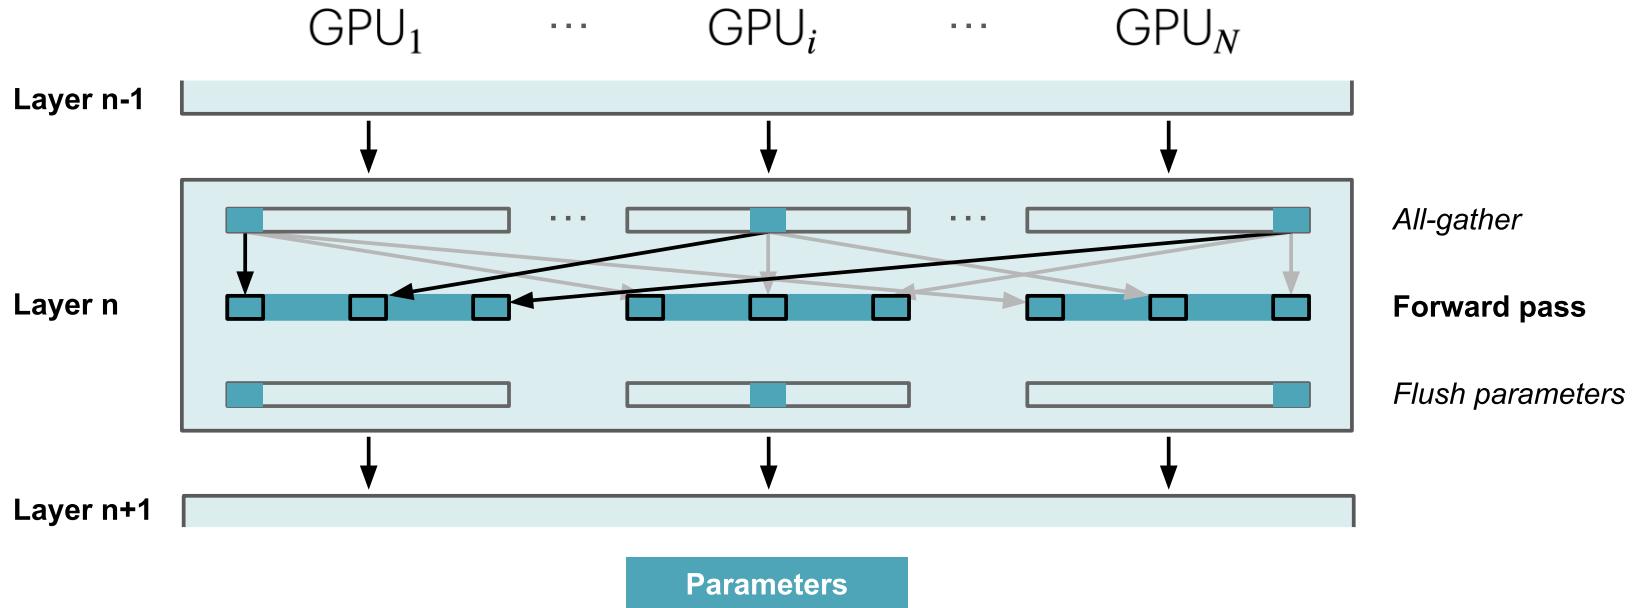


Memory/GPU for a 7.5B model:  
 $7.5\text{B} * (2/3 + 2/3 + 4)$  bytes = **40 GB!**



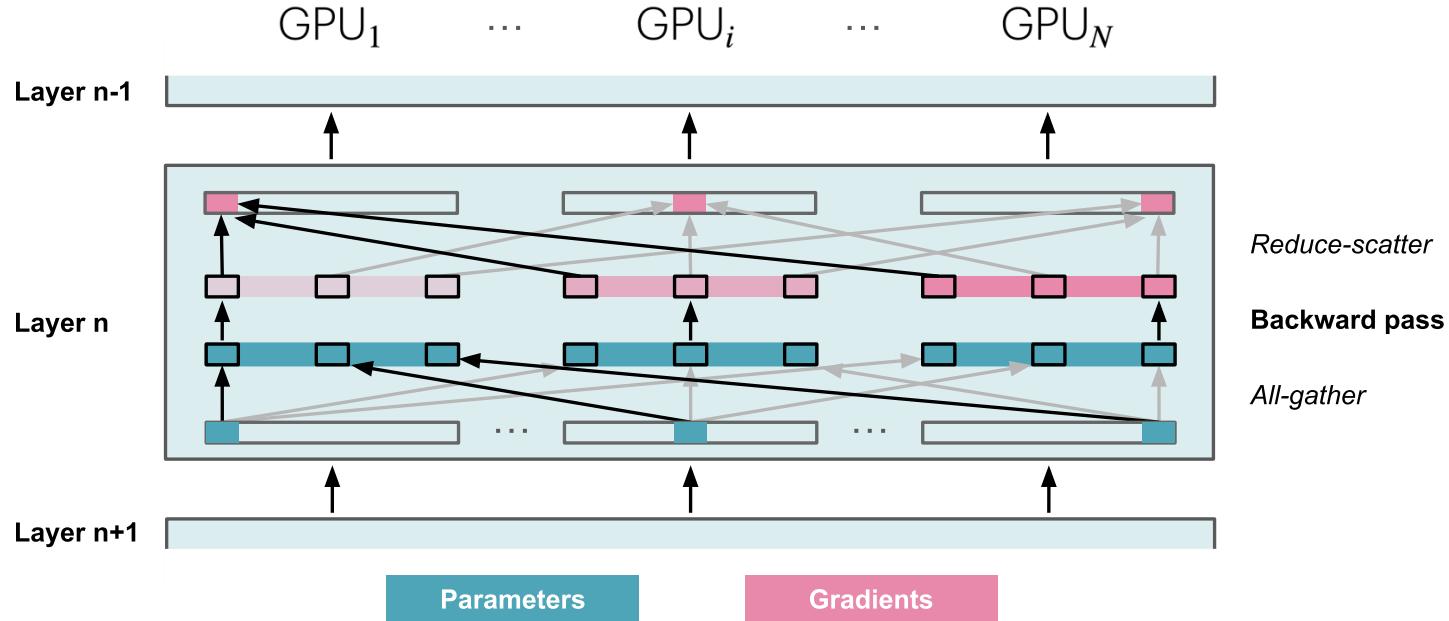
# ZeRO Stage 3: How it works (simplified)

During forward pass, the parameters are gathered on-demand



# ZeRO Stage 3: How it works (simplified)

During backward pass, the gradients are scattered (Reduce\_Scatter)



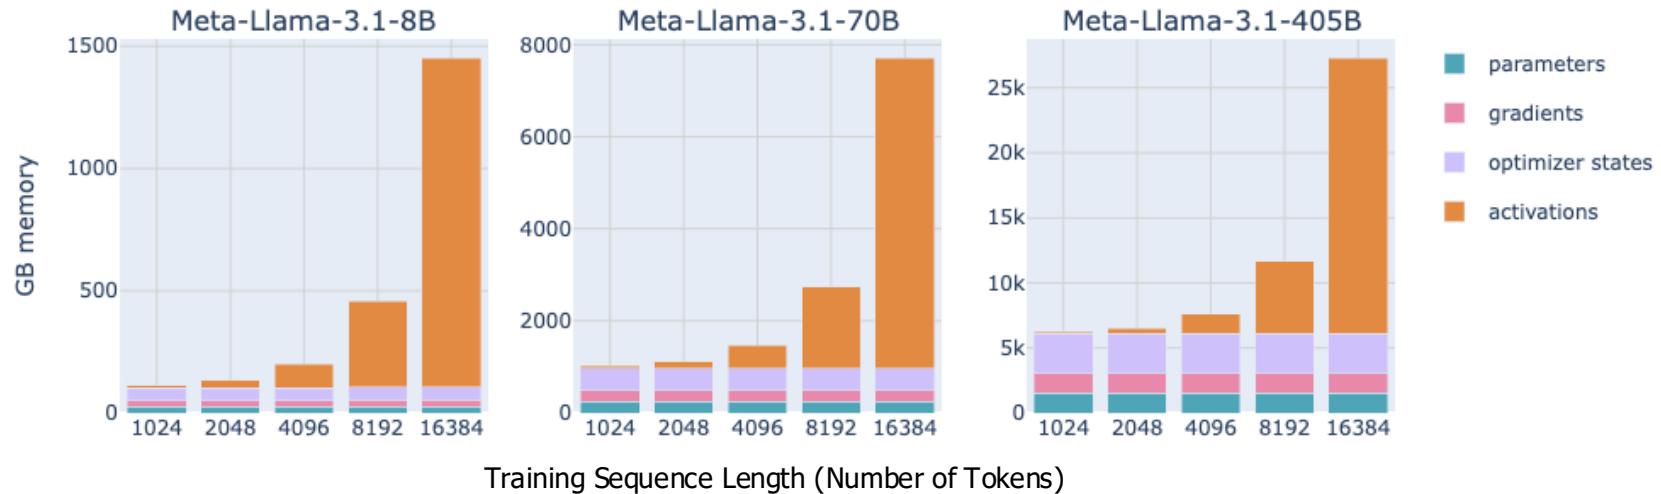
# Communication Costs

---

- Naïve Data Parallel: 2x parameter (all\_reduce)
- ZeRO-1: 2x parameter (reduce\_scatter + all\_gather) - this is free! Might as well always use it.
- ZeRO-2: 2x parameter (reduce\_scatter + all\_gather + overhead) - this is (almost) free!
- ZeRO-3: 3x parameter – which can be quite slow.

# Where did all the memory go?

So far, we dealt with the **optimizer states**  
but what about the **activations**?



Source: <https://nanotron-ultrascale-playbook.static.hf.space/dist/index.html>

# Prefix Caching

---

but what about the activations?

<System> You are a helpful assistant ... <System>

<User> I want to know how can I use the coffee machine <User>

<System> You are a helpful assistant ... <System>

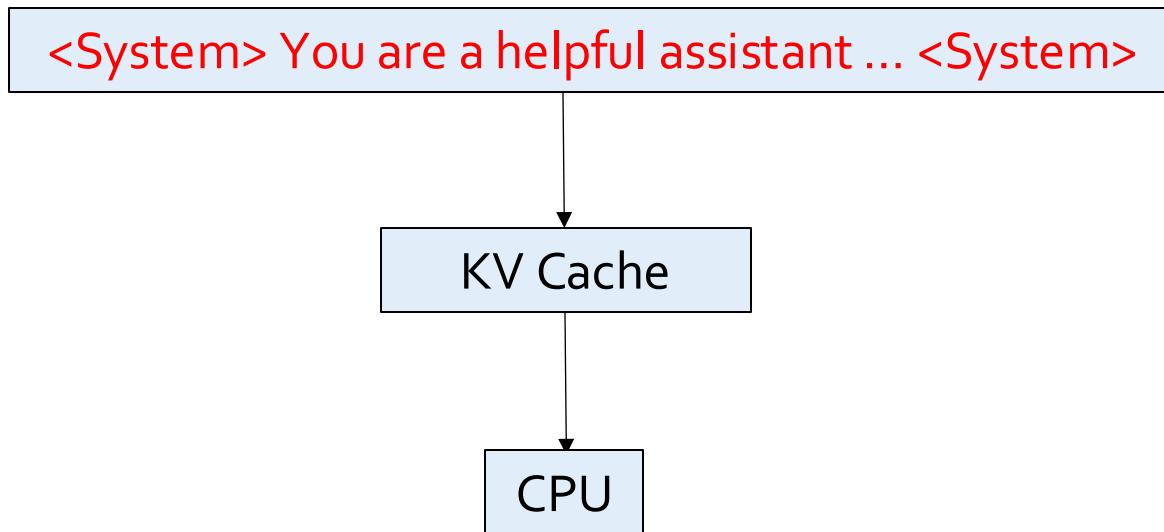
<User> Write the code for training my language model. <User>

<System> You are a helpful assistant ... <System>

<User> Help me revise my email ... <User>

# Prefix Caching

Storing the **activations** in CPU and retrieve it when needed.

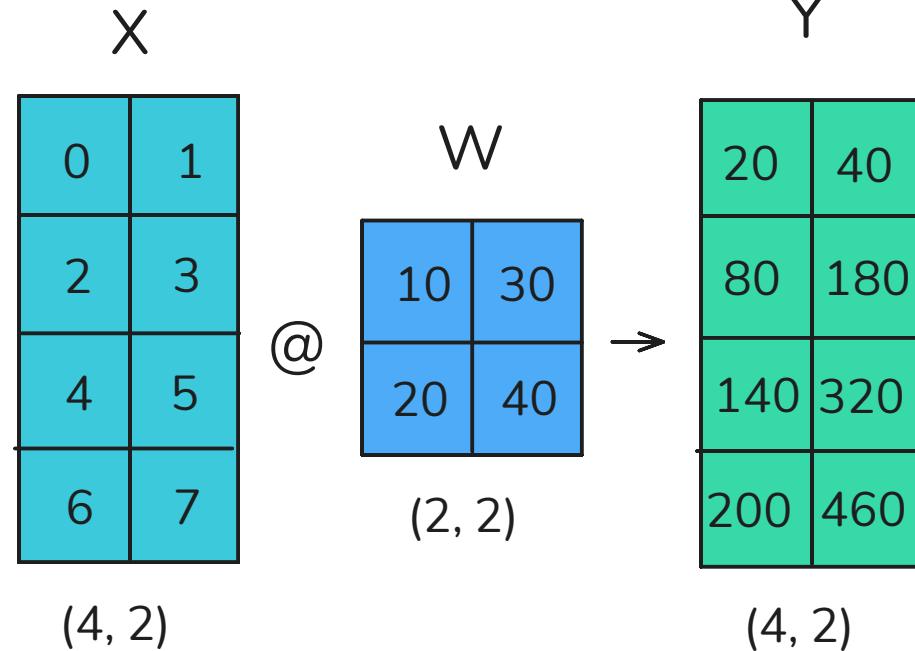


But, can we slice the activations to fit them in different GPUs?  
- Yes, by Tensor Parallelism

# Tensor Parallelism

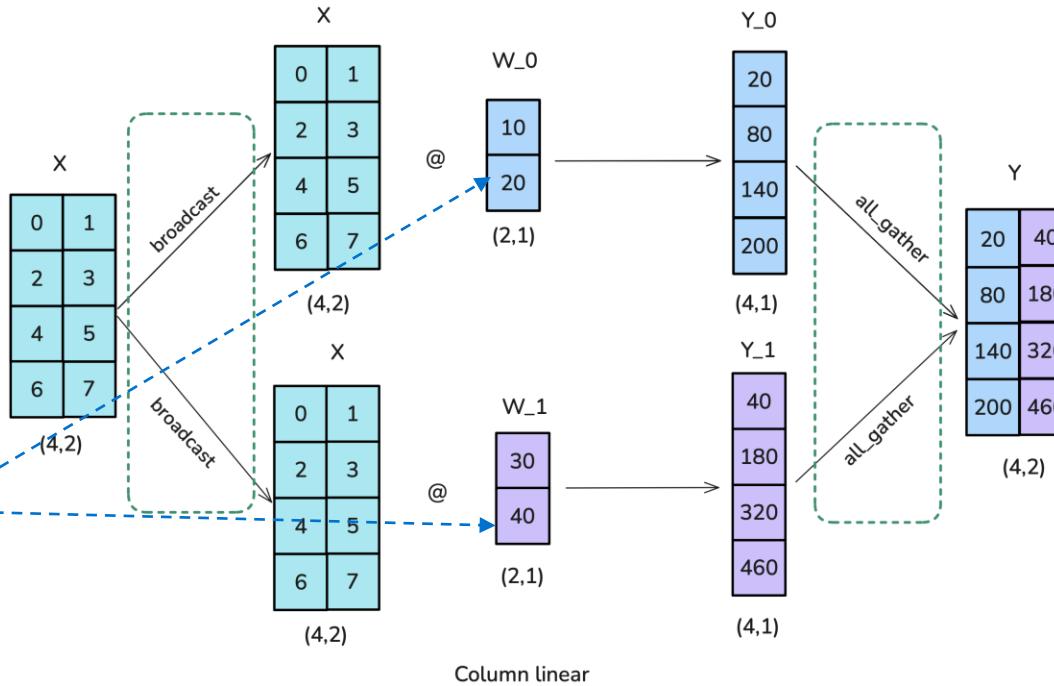
We can either cut  
the weights W into  
two columns  
(Column Parallelism)

or into two rows  
(Row Parallelism)



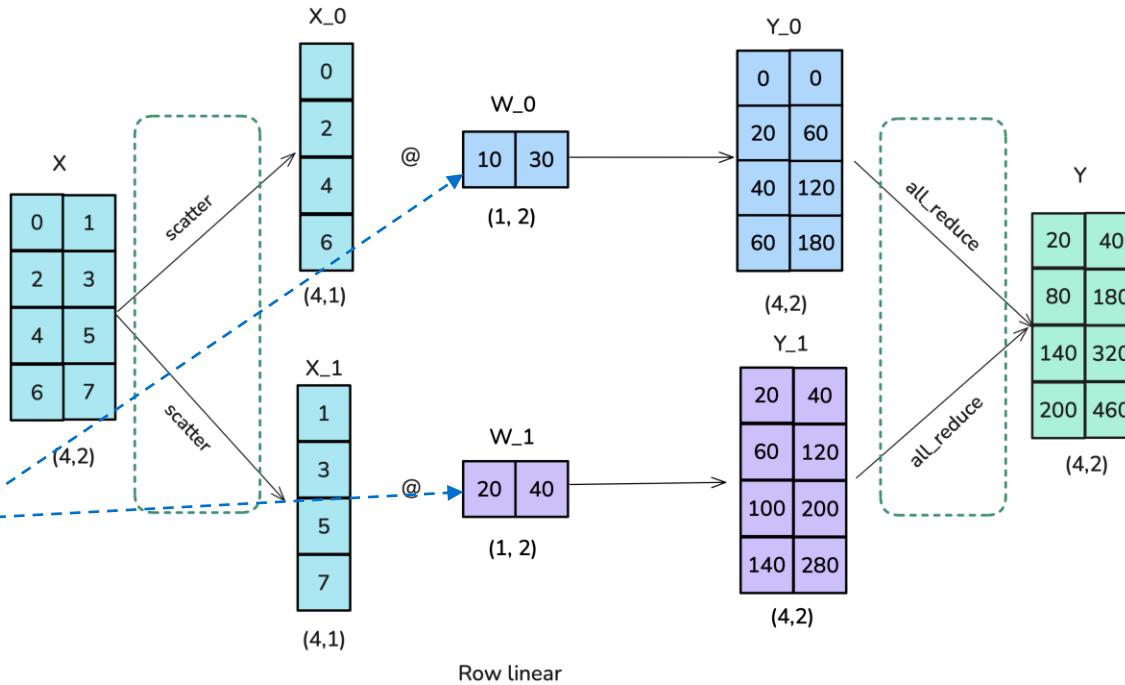
# Column-wise Tensor Parallelism

Cuts the weight matrix W into 2 columns



Source: <https://nanotron-ultrascaling-playbook.static.hf.space/dist/index.html>

# Row-wise Tensor Parallelism



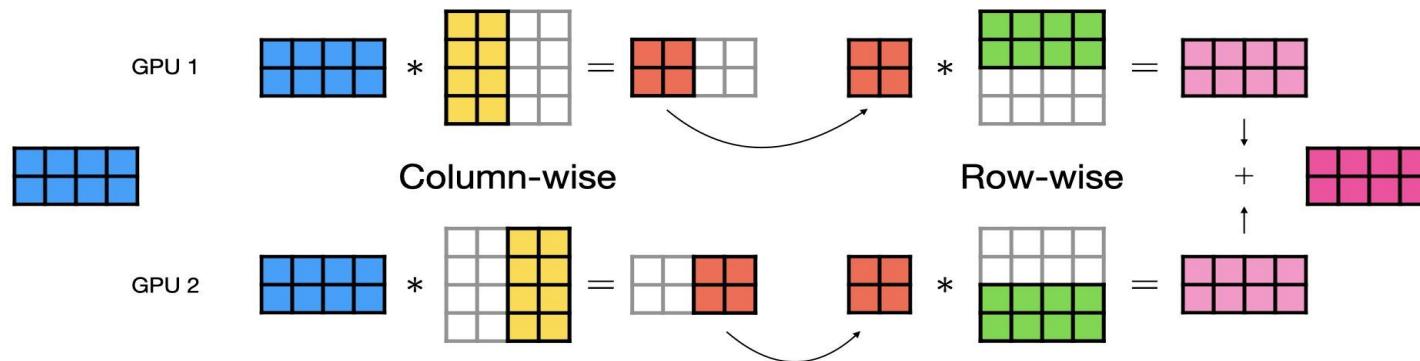
Source: <https://nanotron-ultrascale-playbook.static.hf.space/dist/index.html>

# Tensor Parallelism

Computing matrix multiplications without storing internal activations (e.g.  $xW_1$ )

$$X \quad \quad W_1 \quad \quad W_2 \quad \quad Y$$
$$\begin{matrix} \text{blue grid} \\ * \end{matrix} \quad \quad \begin{matrix} \text{yellow grid} \\ * \end{matrix} \quad \quad \begin{matrix} \text{green grid} \\ = \end{matrix} \quad \quad \begin{matrix} \text{pink grid} \end{matrix}$$

In Feed-Forward Networks,  
The dimension of  
 $W_1$  is usually 4x the hidden  
dimension.



# Tensor Parallelism: Llama Feed-Forward

```
self.w1 = ColumnParallelLinear(  
    dim, hidden_dim, bias=False, gather_output=False, init_method=lambda x: x  
)  
self.w2 = RowParallelLinear(  
    hidden_dim, dim, bias=False, input_is_parallel=True, init_method=lambda x: x  
)  
self.w3 = ColumnParallelLinear(  
    dim, hidden_dim, bias=False, gather_output=False, init_method=lambda x: x  
)  
  
def forward(self, x):  
    return self.w2(F.silu(self.w1(x)) * self.w3(x))
```

activations are element-wise operations, can be parallelized

Source: <https://github.com/meta-llama/llama/blob/main/llama/model.py>

# Tensor Parallelism: Llama Attention

Column Parallel for Query, Key and Vector and Row Parallel for attention output

```
self.wq = ColumnParallelLinear(  
    args.dim,  
    args.n_heads * self.head_dim,  
    bias=False,  
    gather_output=False,  
    init_method=lambda x: x,  
)  
  
self.wk = ColumnParallelLinear(  
    args.dim,  
    self.n_kv_heads * self.head_dim,  
    bias=False,  
    gather_output=False,  
    init_method=lambda x: x,  
)  
  
self.wv = ColumnParallelLinear(  
    args.dim,  
    self.n_kv_heads * self.head_dim,  
    bias=False,  
    gather_output=False,  
    init_method=lambda x: x,  
)
```

```
self.wo = RowParallelLinear(  
    args.n_heads * self.head_dim,  
    args.dim,  
    bias=False,  
    input_is_parallel=True,  
    init_method=lambda x: x,  
)
```

Source: <https://github.com/meta-llama/llama/blob/main/llama/model.py>

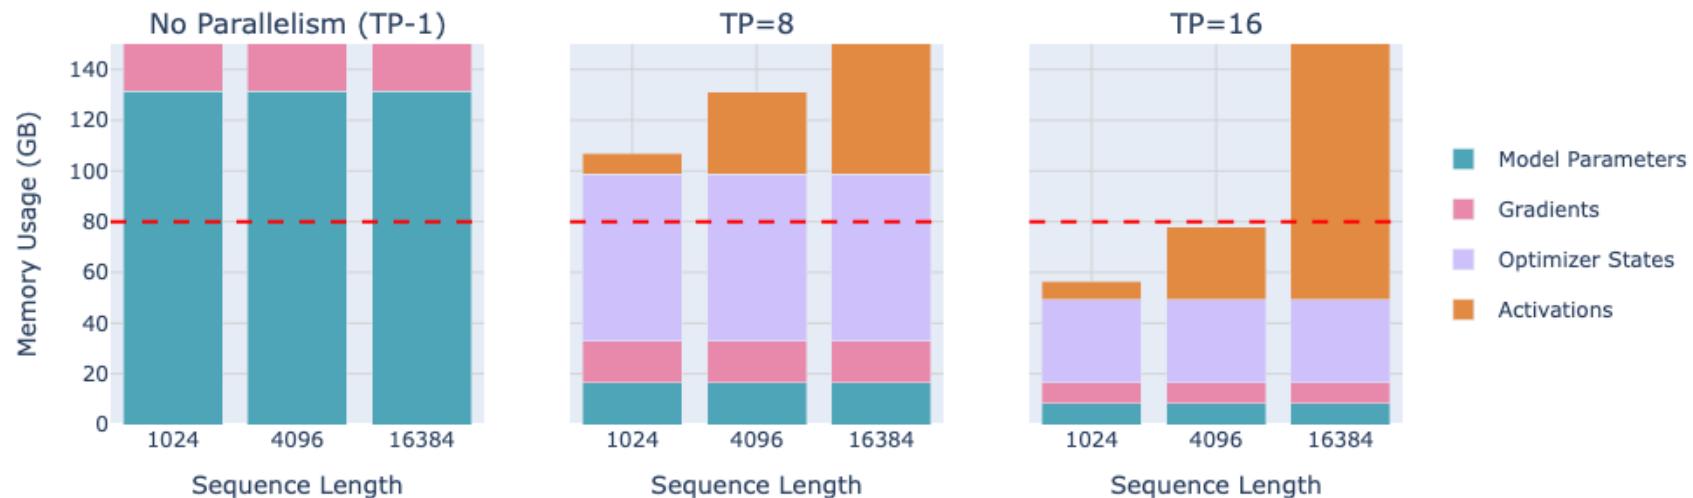
# Summary so far

---

- Data Parallelism
  - Naïve Data Parallelism
  - NCCL Operations  
(reduce, all\_reduce, reduce\_scatter, broadcast, all\_gather)
  - ZeRO-1, ZeRO-2, ZeRO-3
- Prefix Caching
- Tensor Parallelism
  - Row-wise Tensor Parallelism
  - Column-wise Tensor Parallelism

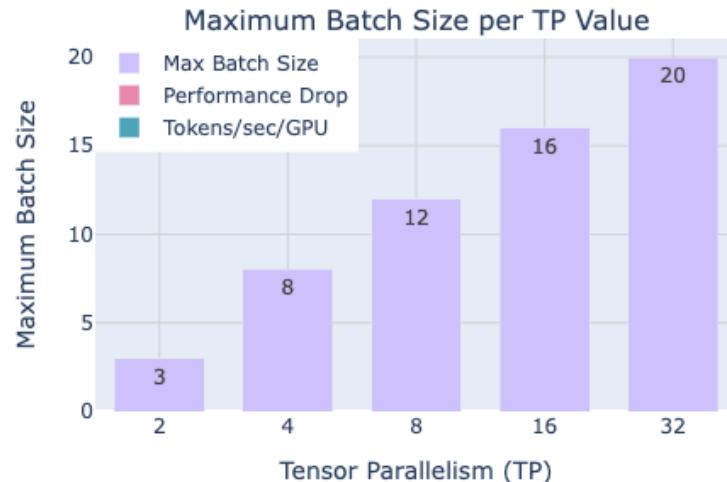
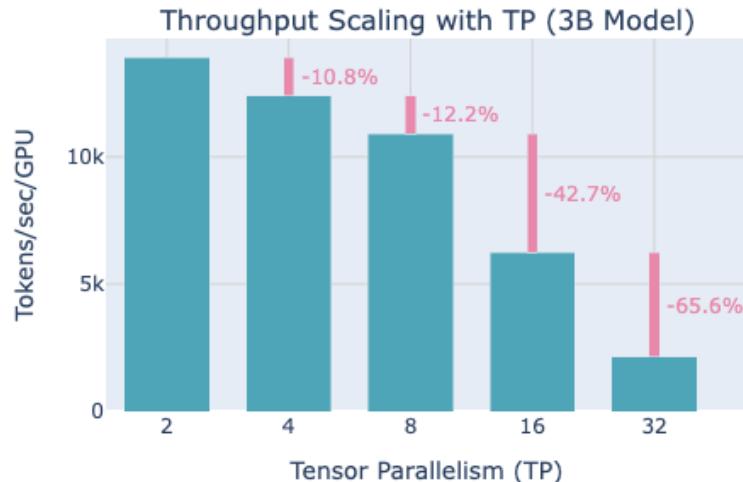
# Tensor Parallelism

Memory Usage for 70B Model



Source: <https://github.com/meta-llama/llama/blob/main/llama/model.py>

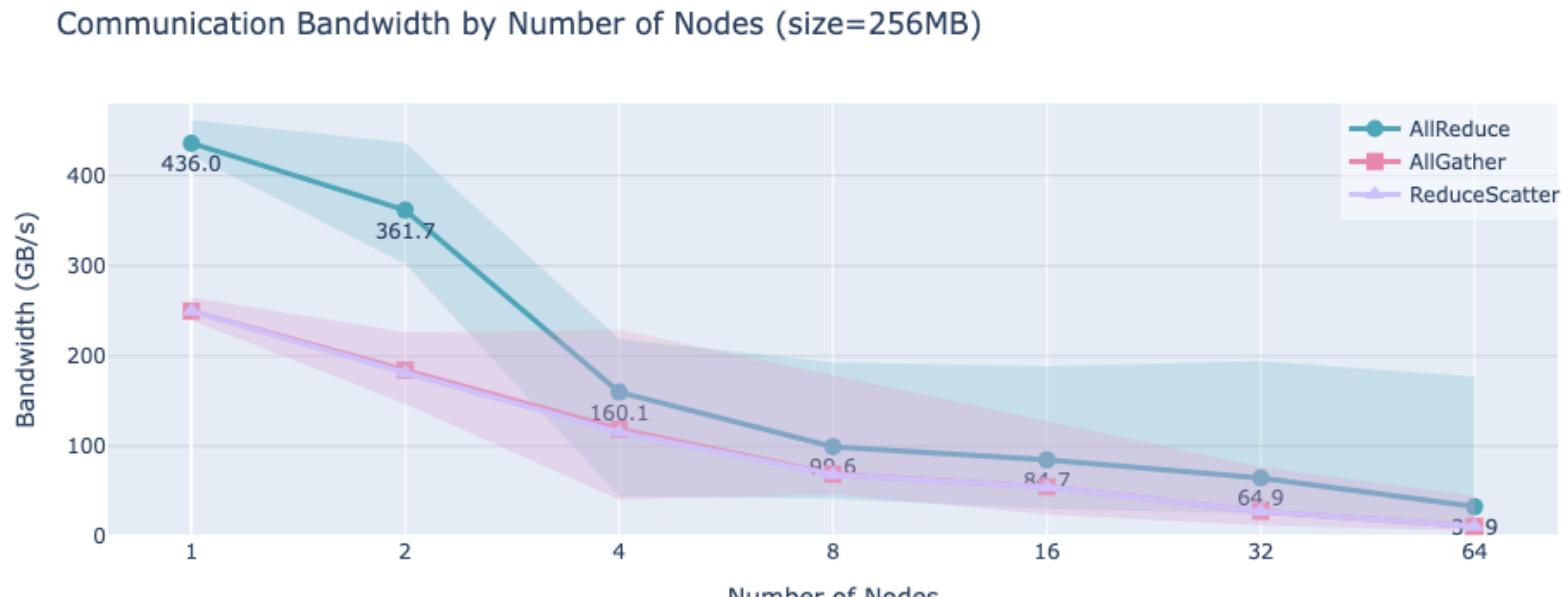
# Throughput Scaling of Tensor Parallelism



A large drop in throughput when scaling beyond 8 GPUs (one node)

Source: <https://nanotron-ultrascale-playbook.static.hf.space/dist/index.html>

# Throughput Scaling of Tensor Parallelism



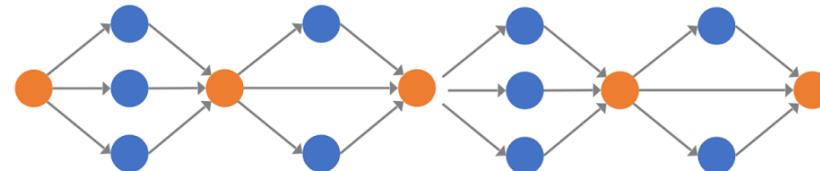
Throughput drops significantly once we go beyond one node!

Source: <https://nanotron-ultrascaling-playbook.static.hf.space/dist/index.html>

# Pipeline Parallelism

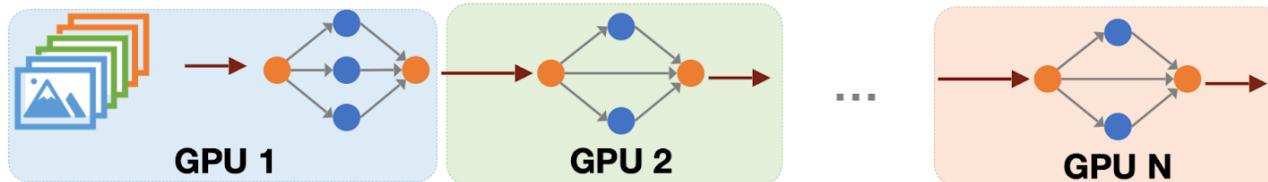


Training Dataset



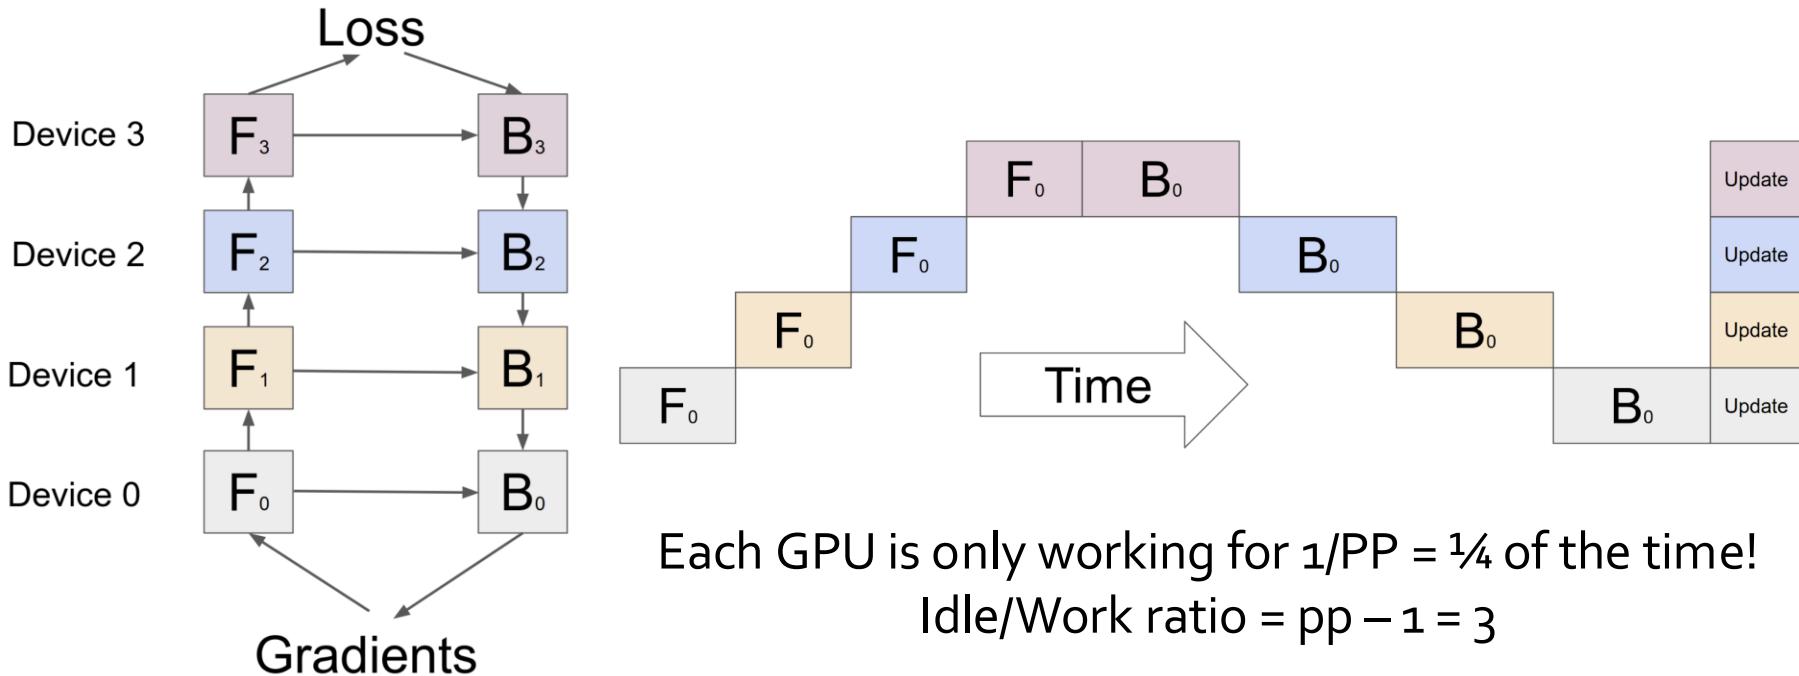
ML Model

Shard each layer of the model into individual GPUs:  
Prevents the cost of syncing params



Credit: Song Han (MIT)

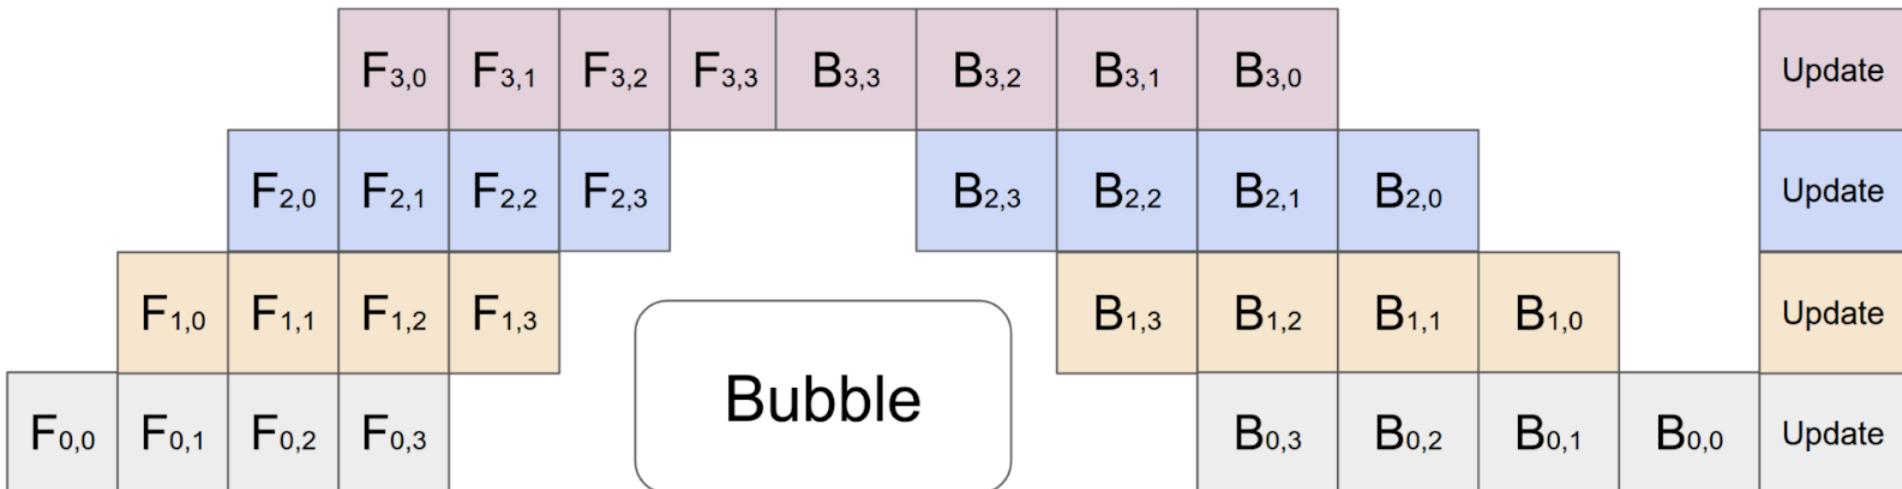
# Pipeline Parallelism



[GPipe: Easy Scaling with Micro-Batch Pipeline Parallelism](#) (Huang et al., NeurIPS 2019)

# Pipeline Parallelism: Improvement

Solution: Splitting the data into mini-batches! (AFAB)



$$\text{Idle / Work Ratio} = \text{PP-1} / M = 3 / 4$$

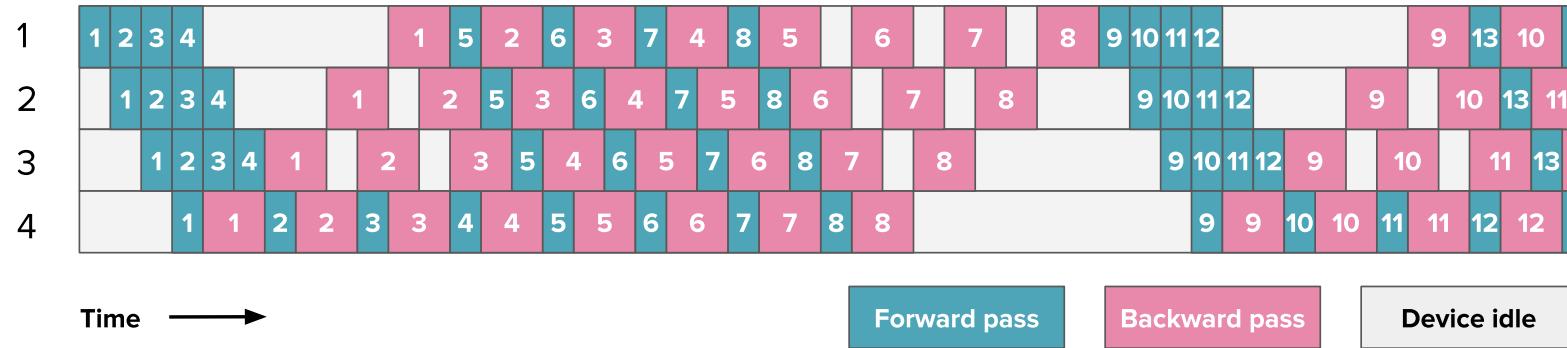
[GPipe: Easy Scaling with Micro-Batch Pipeline Parallelism](#) (Huang et al., NeurIPS 2019)

# Pipeline Parallelism

A cleverer version of AFAB: 1 Forward 1 Backward (1F1B)

Idea: Do backward as early as possible, releasing activations on the fly

GPU

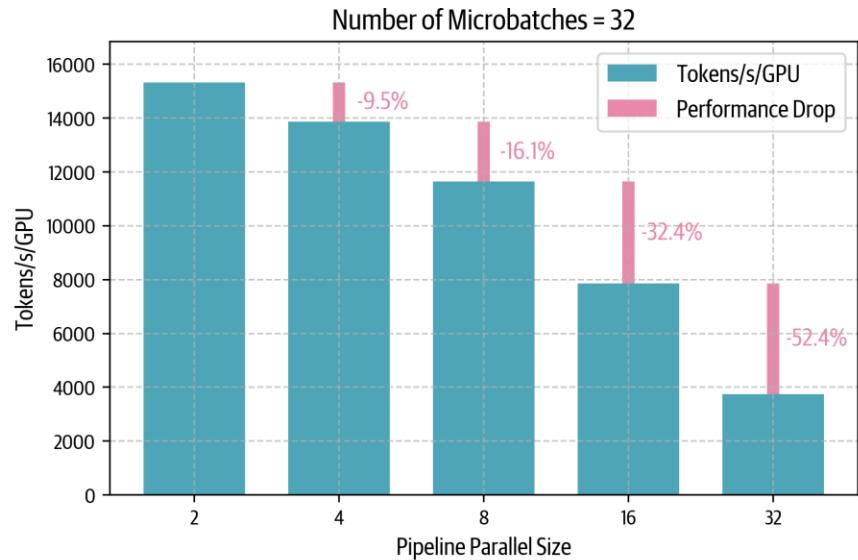
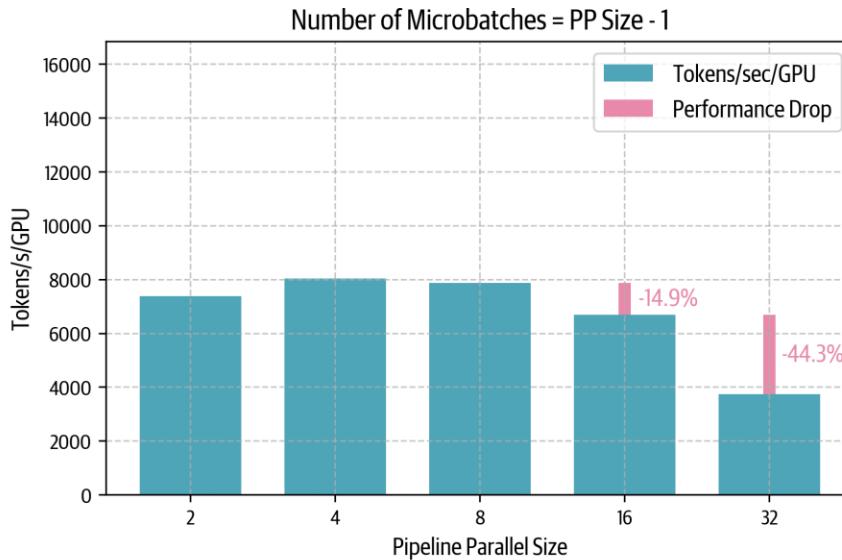


Roughly the same Idle/Work Ratio but less memory  
(as you only need to store  $p=4$  activations rather than  $m=8$ )

[GPipe: Easy Scaling with Micro-Batch Pipeline Parallelism](#) (Huang et al., NeurIPS 2019)

# Pipeline Parallelism Throughput

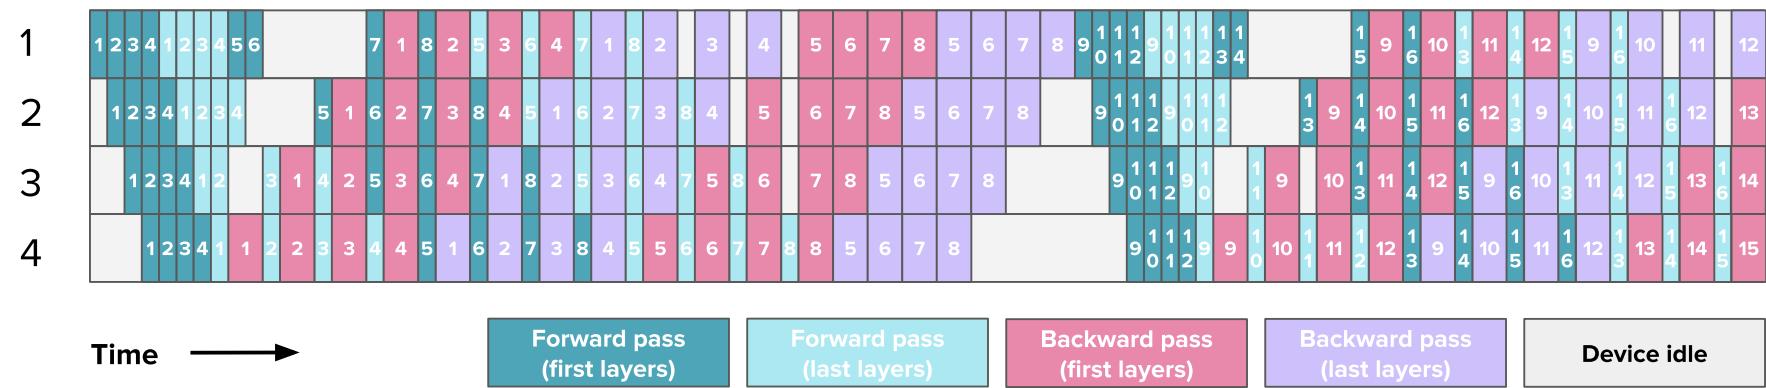
Throughput Scaling with Pipeline Parallelism (1F1B schedule)



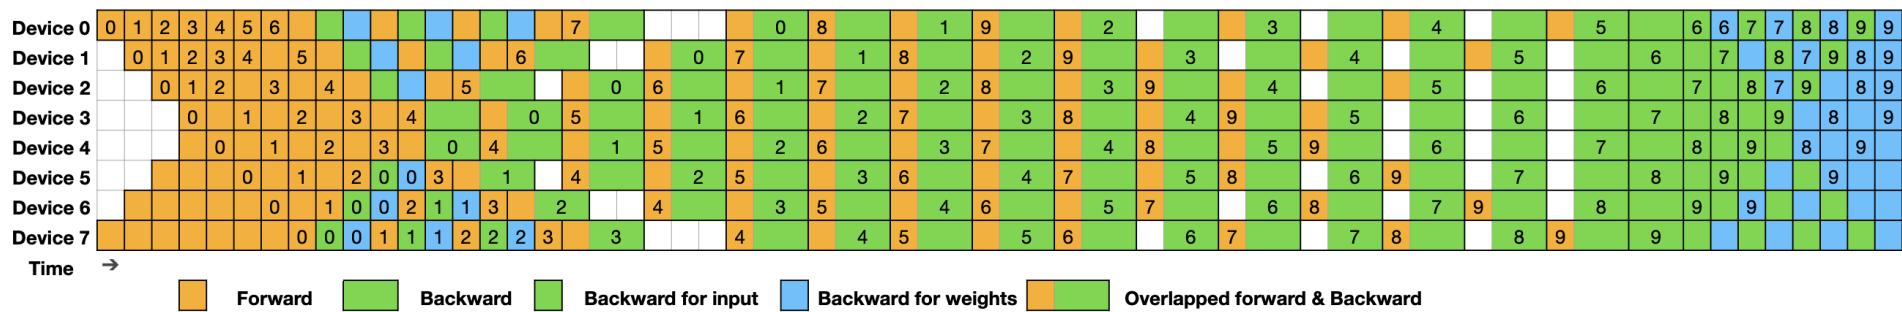
A **small** drop in throughput when scaling beyond 8 GPUs (one node)  
but a large drop as we increase the microbatch number

# Interleaving Pipeline Parallelism (LLama3)

GPU



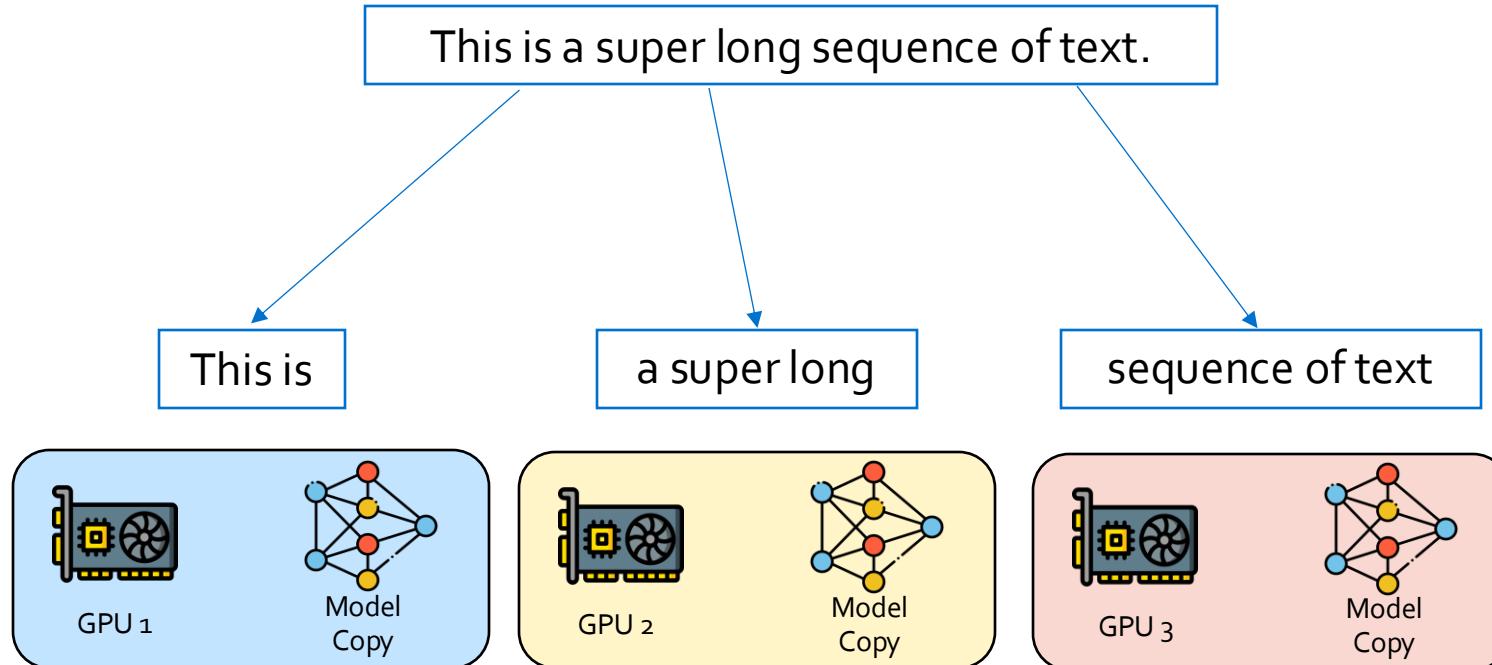
# Interleaved Pipeline Parallelism (DeepSeek)



backprop for weights (blue) can be computed at any time!  
We fill in the bubble with weight back propagation.

# What about (super) long sequences?

Suppose we want to split the sequence into different GPUs

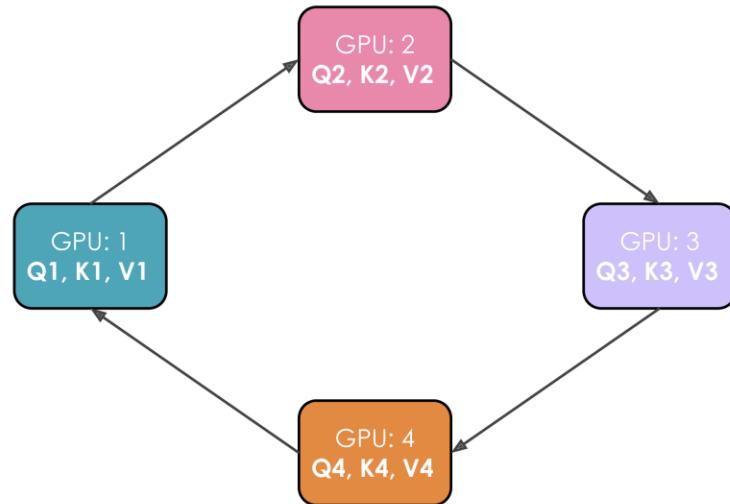


# What about (super) long sequences?

---

- Feed Forward Network / LayerNorm is not affected by splitting the sequence, each token is processed individually
- But what about attention? Each token needs to compute dot product with every other token.

# Context Parallelism (Ring Attention)

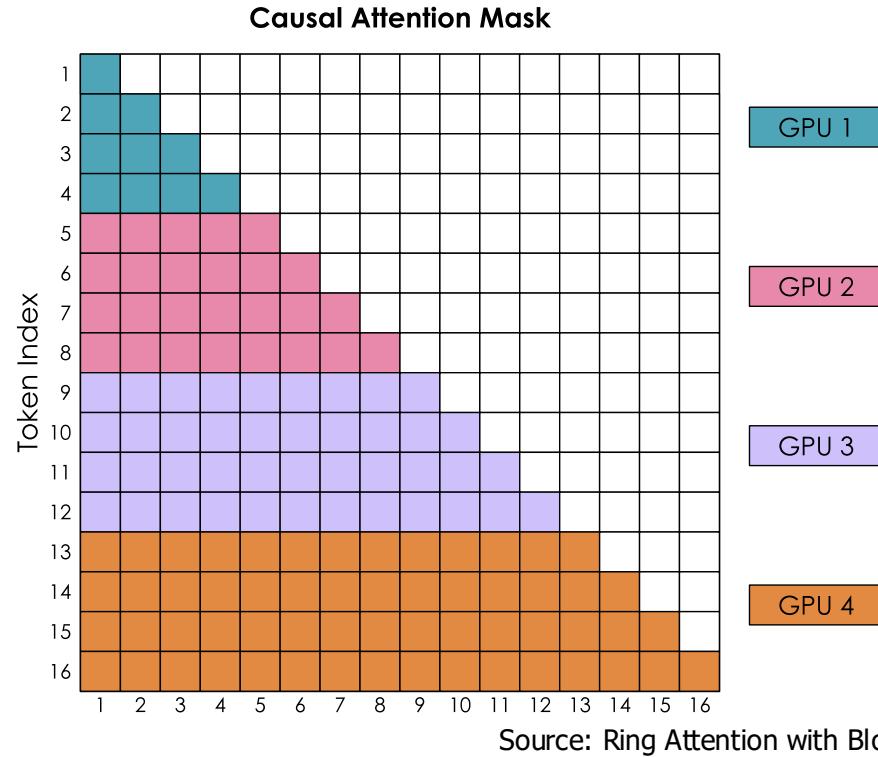


Just pass the Key, Value pairs around!

However, attention mask is usually causal – Q<sub>1</sub> does not need K<sub>2</sub>, V<sub>2</sub>, ...

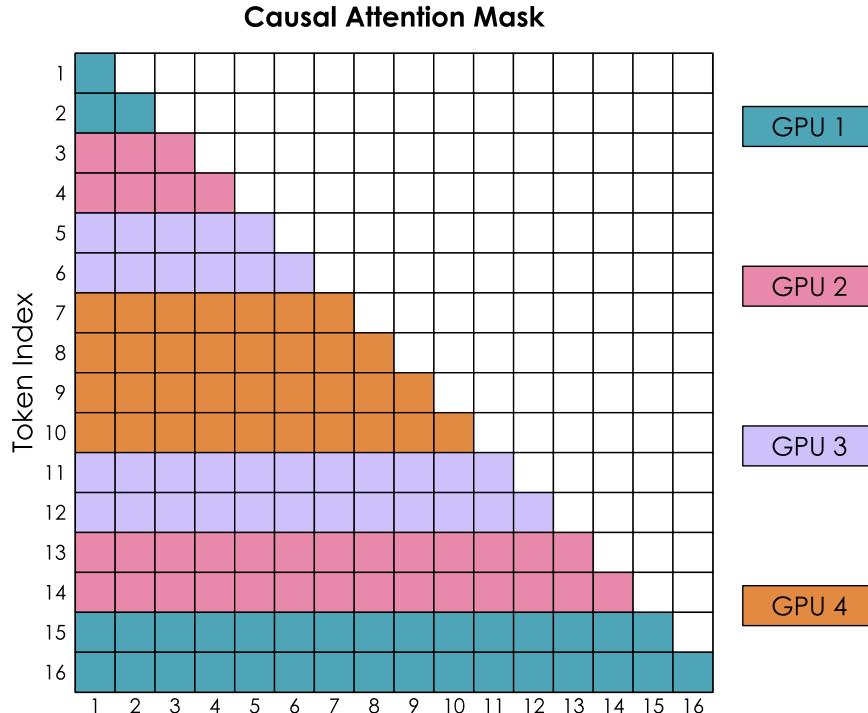
Source: Ring Attention with Blockwise Transformers For Near-Infinite Context (Liu et al., 2023)

# Context Parallelism (Ring Attention)



GPU 1 computes the pre-softmax-ed scores for Q<sub>1</sub>, Q<sub>2</sub>, Q<sub>3</sub>, Q<sub>4</sub>.. then becomes idle.

# Context Parallelism (Ring Attention)

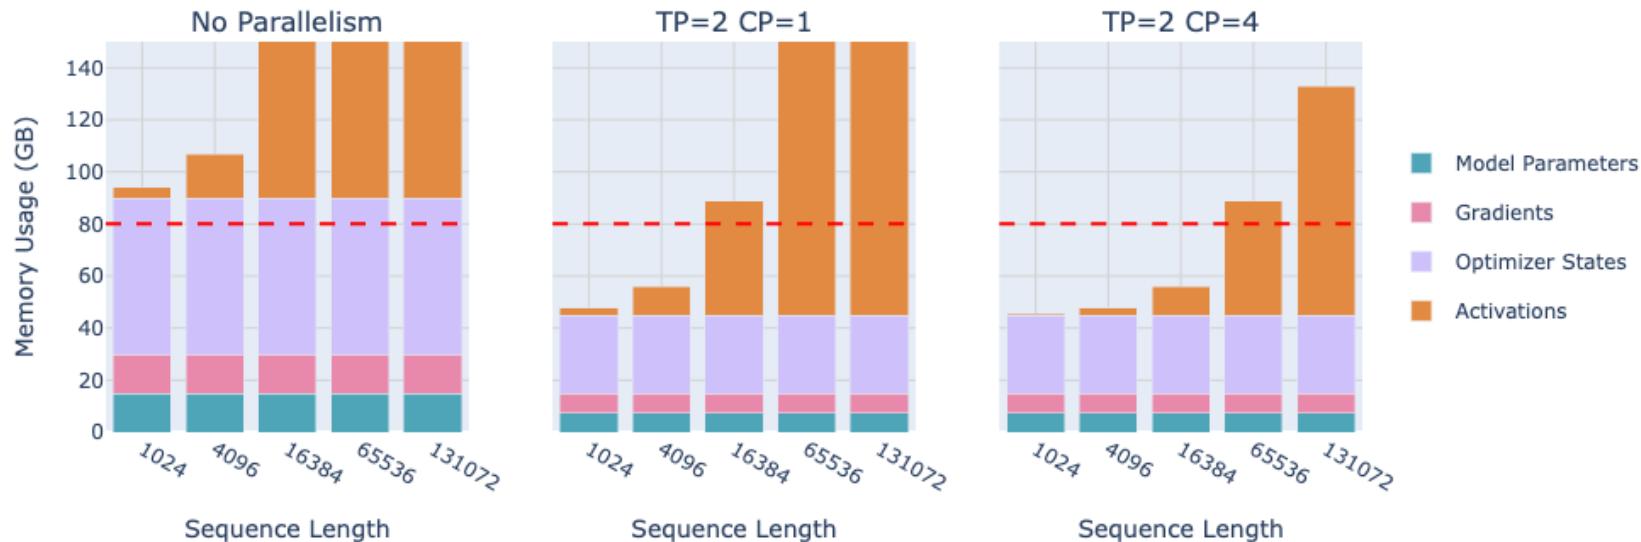


Balancing the workload  
for each individual GPUs.

Source: Striped Attention: Faster Ring Attention for Causal Transformers (Brandon et al., 2023)

# Context Parallelism (Ring Attention)

Memory Usage for 8B Model



Source: <https://nанотрон-ultrascale-playbook.static.hf.space/dist/index.html>

# Summarizing

	Sync overhead	Memory	Bandwidth	Batch size	Easy to use?
DDP/ZeRO1	Per-batch	No scaling	$2^* \# \text{ param}$	Linear	Very
FSDP (ZeRO3)	3x Per-FSDP block	Linear	$3 * \# \text{ param}$	Linear	Very
Pipeline	Per-pipeline	Linear	Activations	Linear	No
Tensor+seq	2x transformer block	Linear	8*activations per-layer all-reduce	No impact	No

Source: Tatsunori Hashimoto (Stanford)

# Solutions

- DeepSeek V3: DP=1, PP=16, EP (Expert Parallelism) = 8

## 3.2 Training Framework

The training of DeepSeek-V3 is supported by the HAI-LLM framework, an efficient and lightweight training framework crafted by our engineers from the ground up. On the whole, DeepSeek-V3 applies 16-way Pipeline Parallelism (PP) (Qi et al., [2023a](#)), 64-way Expert Parallelism (EP) (Lepikhin et al., [2021](#)) spanning 8 nodes, and ZeRO-1 Data Parallelism (DP) (Rajbhandari et al., [2020](#)).

- Llama 3: Staged Training

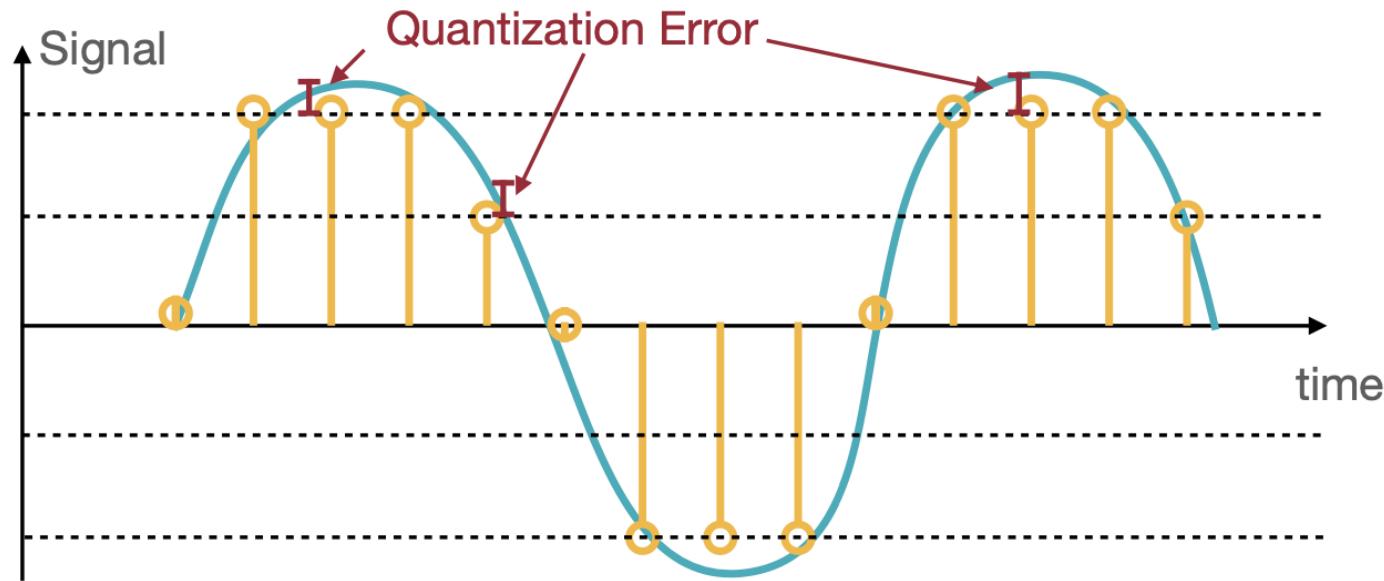
GPUs	TP	CP	PP	DP	Seq. Len.	Batch size/DP	Tokens/Batch	TFLOPs/GPU	BF16 MFU
8,192	8	1	16	64	8,192	32	16M	430	43%
16,384	8	1	16	128	8,192	16	16M	400	41%
16,384	8	16	16	8	131,072	16	16M	380	38%

**Table 4 Scaling configurations and MFU for each stage of Llama 3 405B pre-training.** See text and Figure 5 for descriptions of each type of parallelism.

# Quantization

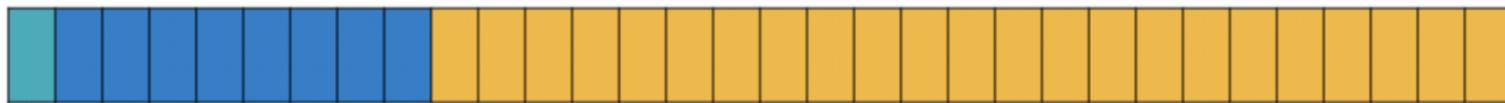
# Quantization: Mapping from high to low precision

— Continuous Signal     —○ Quantized Signal



# Numeric Data Types

- Example: 32-bit floating point number in IEEE 754 (FP32)



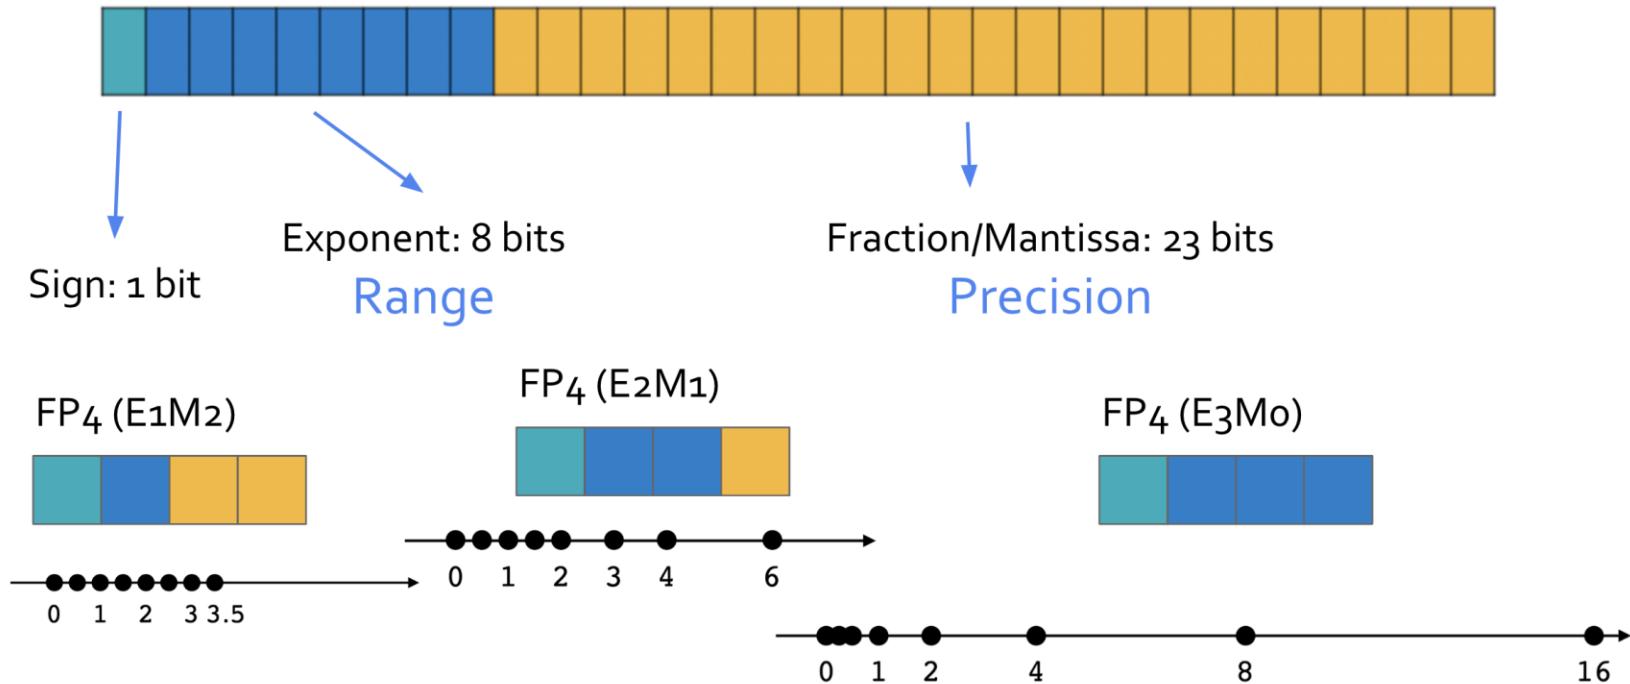
Sign: 1 bit

Exponent: 8 bits

Fraction/Mantissa: 23 bits

$$\text{Number} = (-1)^{\text{sign}} \times (1 + \text{Fraction}) \times 2^{\text{Exponent} - 127}$$

# Floating Point Numbers



# Floating Point Numbers

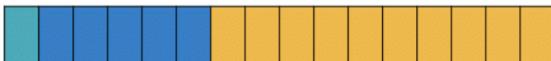
IEEE 754 Single Precision 32-bit Float (FP32)



Exponent      Fraction

8                23

IEEE 754 Half Precision 16-bit Float (FP16)



5                10

Google Brain Float (BF 16)



More range, less precision

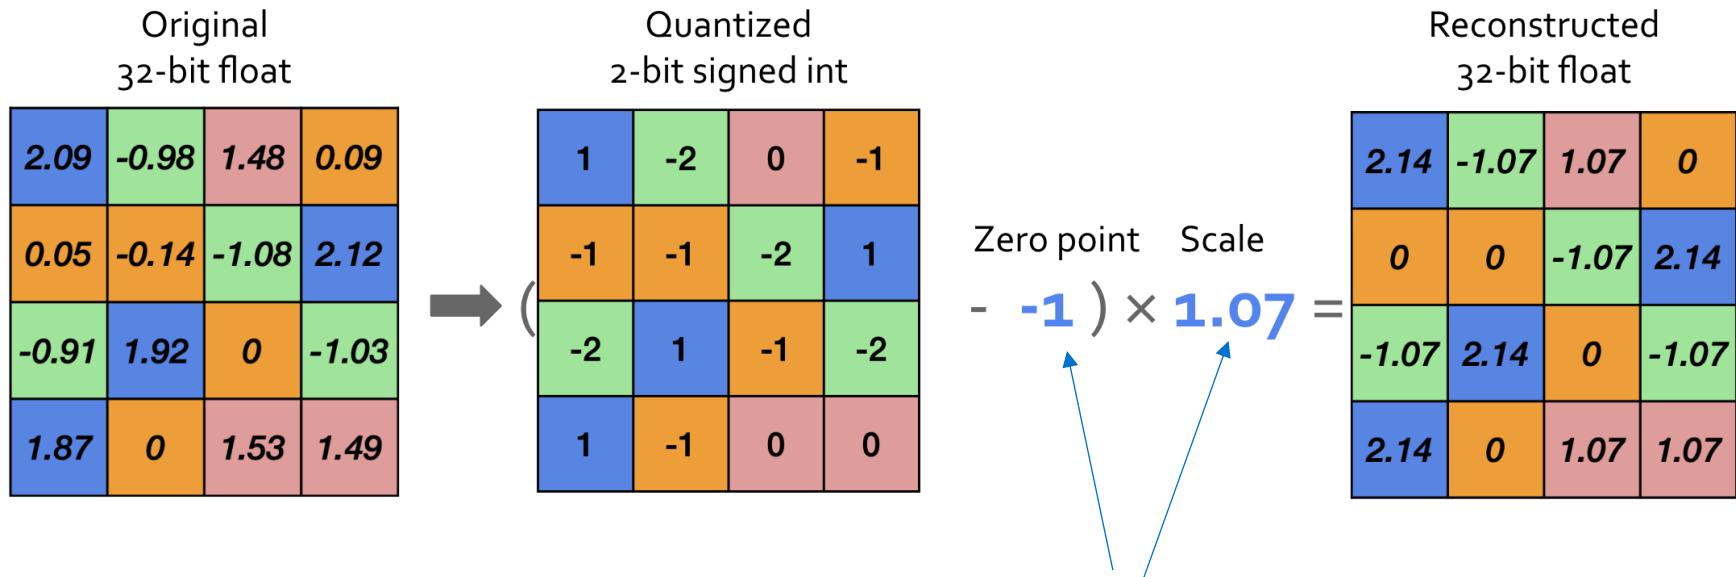
8                7

Nvidia FP8 (E4M3)



4                3

# Linear Quantization



How to find these numbers?

# Linear Quantization

Original 32-bit float			
2.09	-0.98	1.48	0.09
0.05	-0.14	-1.08	2.12
-0.91	1.92	0	-1.03
1.87	0	1.53	1.49

r

≈

Quantized 2-bit signed int			
1	-2	0	-1
-1	-1	-2	1
-2	1	-1	-2
1	-1	0	0

( q )

-

integer

$$\text{Zero point} \quad \text{Scale} \\ - (-1) \times 1.07 =$$

Z ) × S

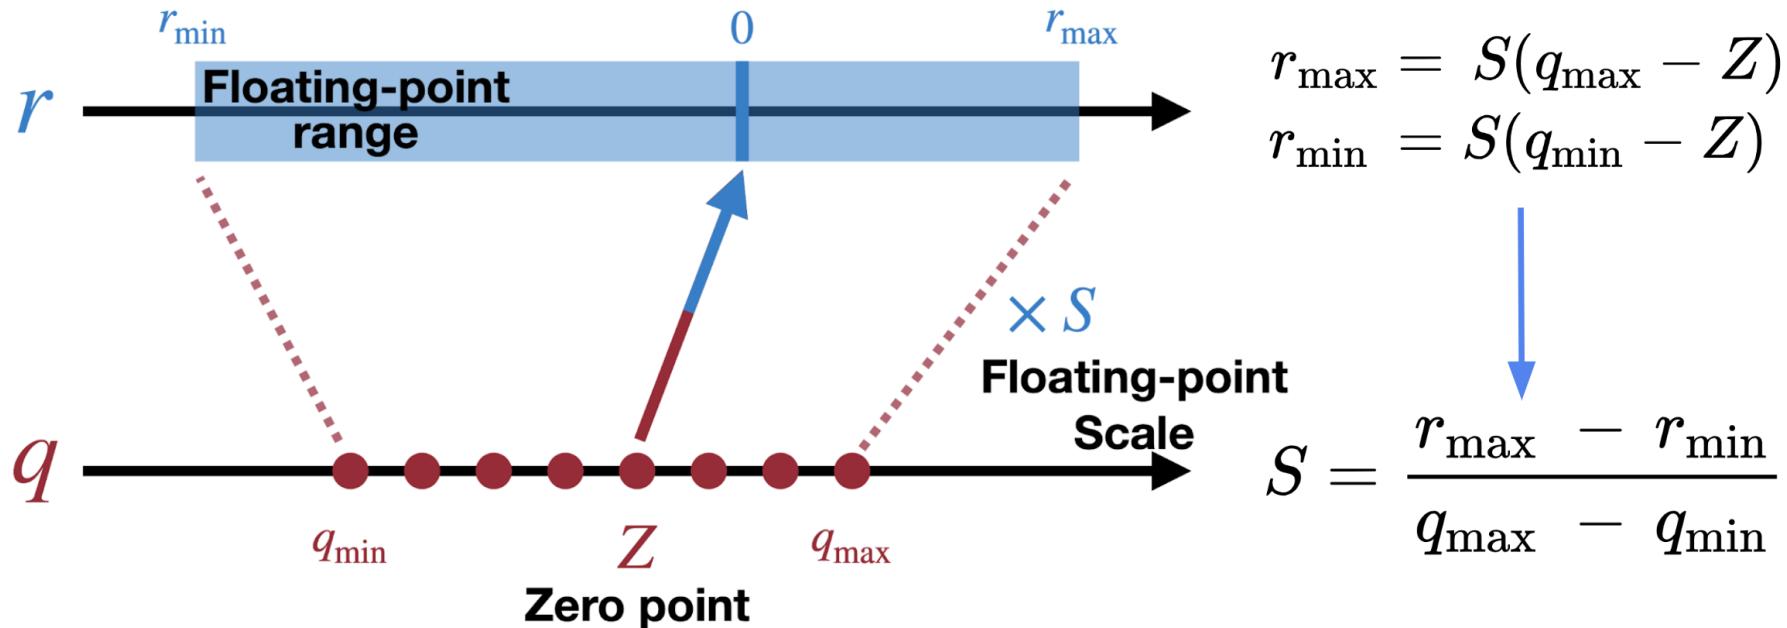
floating-point

Reconstructed  
32-bit float

2.14	-1.07	1.07	0
0	0	-1.07	2.14
-1.07	2.14	0	-1.07
2.14	0	1.07	1.07

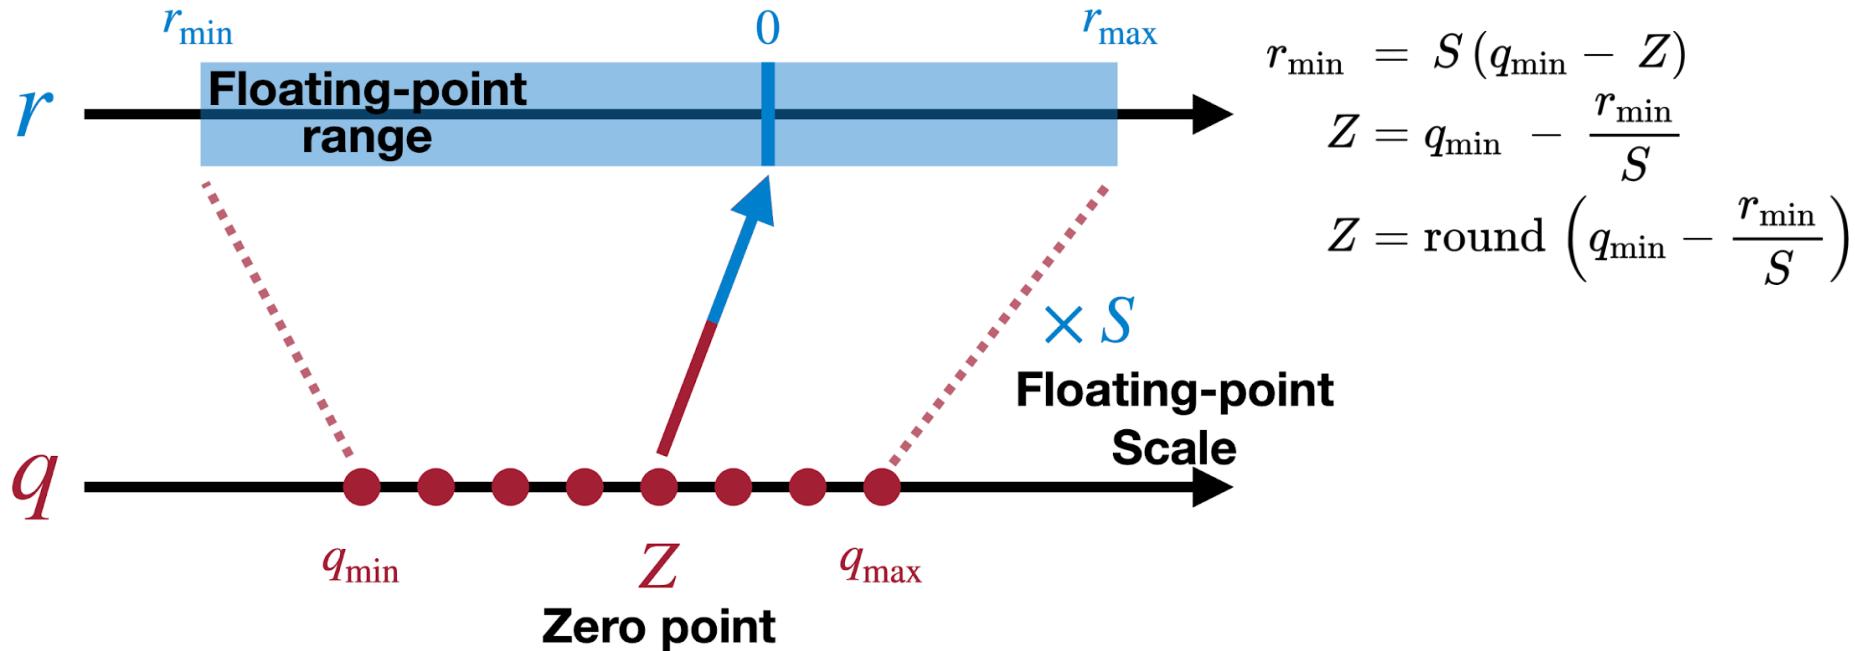
integer floating-point

# Linear Quantization: Scale



[Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference](#) (Jacob et al., CVPR 2018)

# Linear Quantization: Zero Point



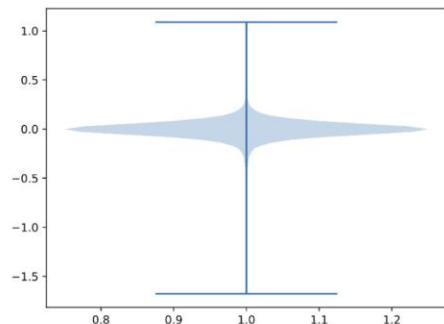
[Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference](#) (Jacob et al., CVPR 2018)

# Linear Quantization: Zero Point

“Absmax” Implementation

In practice, the weights are usually centered around zero ( $Z = 0$ ):

Therefore, we can find scale by using only the max.



Weight distribution of first conv  
layer of ResNet-50.

$$S = \frac{r_{\max} - r_{\min}}{q_{\max} - q_{\min}}$$

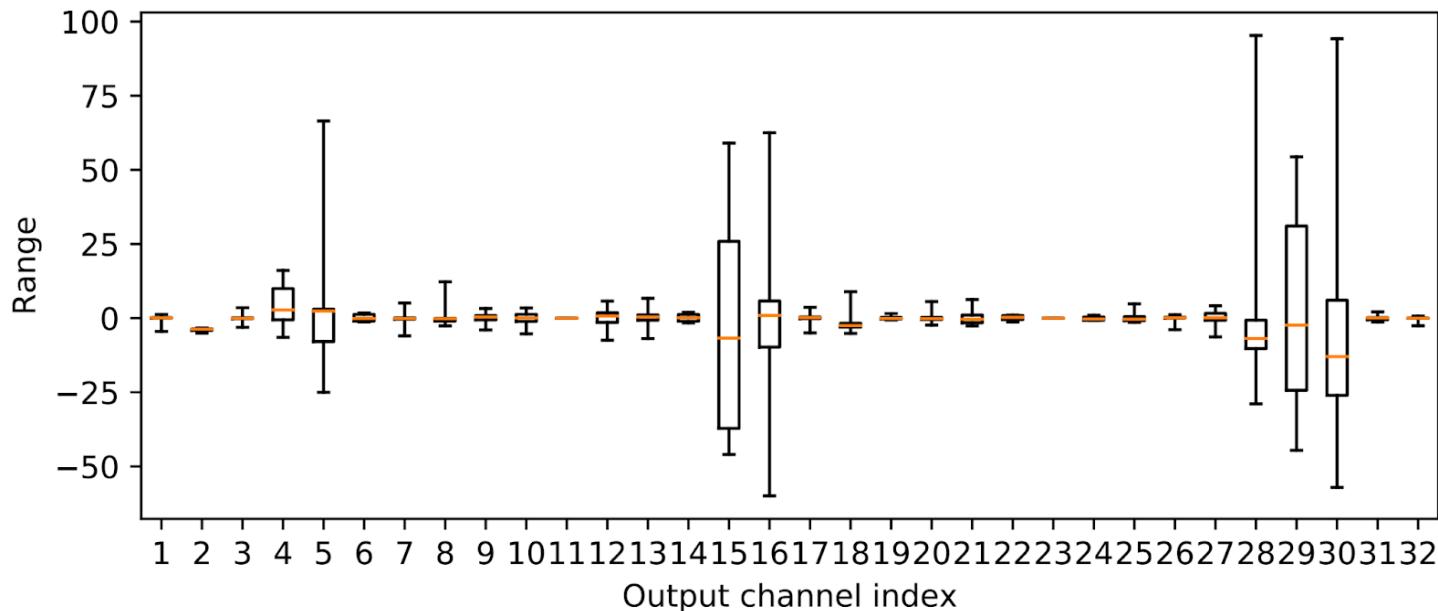


$$S = \frac{r_{\min}}{q_{\min} - Z} = \frac{-|r|_{\max}}{q_{\min}}$$

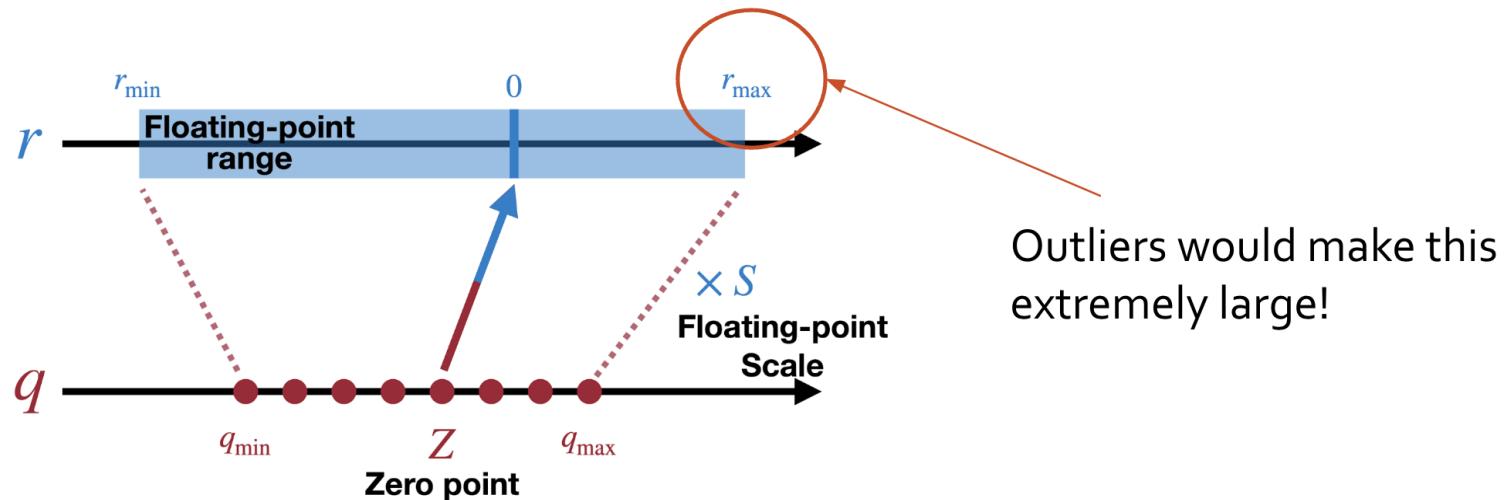
Used in Pytorch, ONNX

# Quantization of Language Models

There exists many outliers in activations (activations of the first layer MobileNetV2):

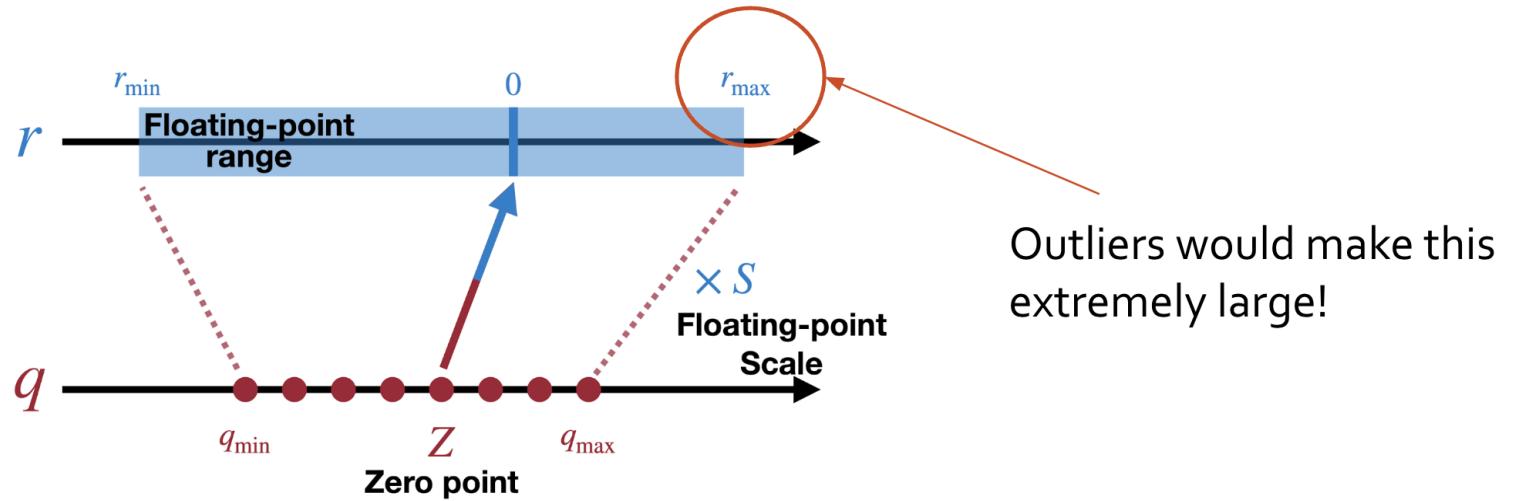


# Quantization of Language Models



Example: 15, 0.1, 0.02, 1.0, 0.01 -> 127, 1, 0, 8, 0  
(Everything under 0.05 gets mapped to 0)

# Quantization of Language Models

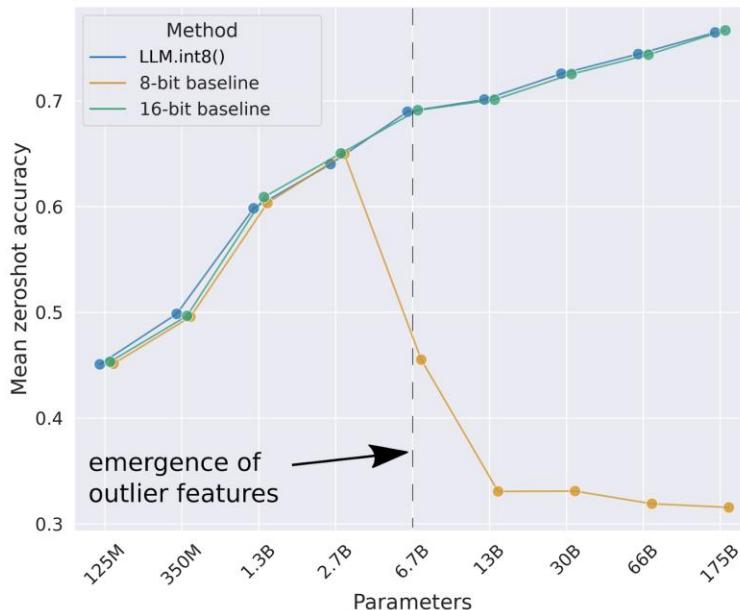


Example: 15, 0.1, 0.02, 1.0, 0.01 -> 127, 1, 0, 8, 0  
(Everything under 0.05 gets mapped to 0)

Quantize each channel individually, each channel gets its own scale and Zero-point!

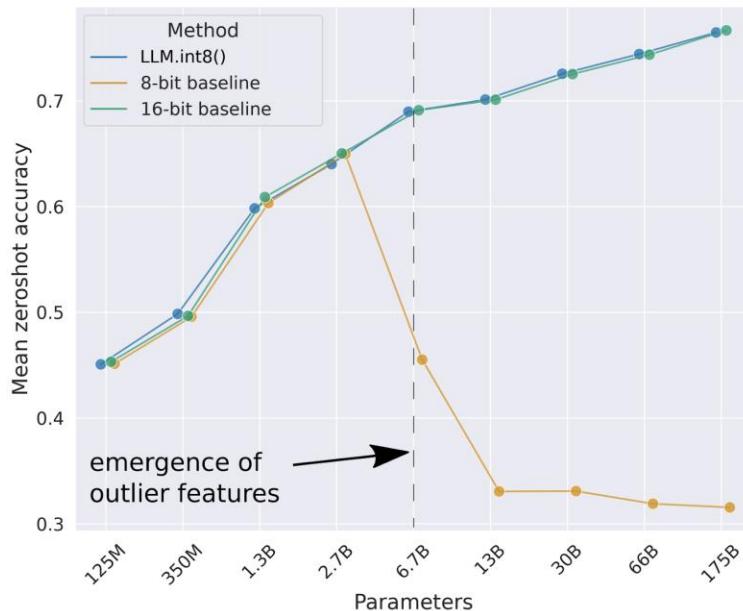
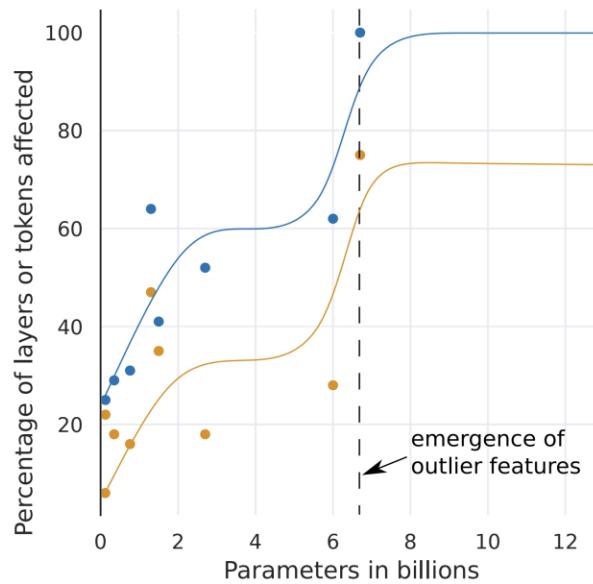
# Quantization of Language Models

Outlier features significantly harms performance after quantization in LMs.



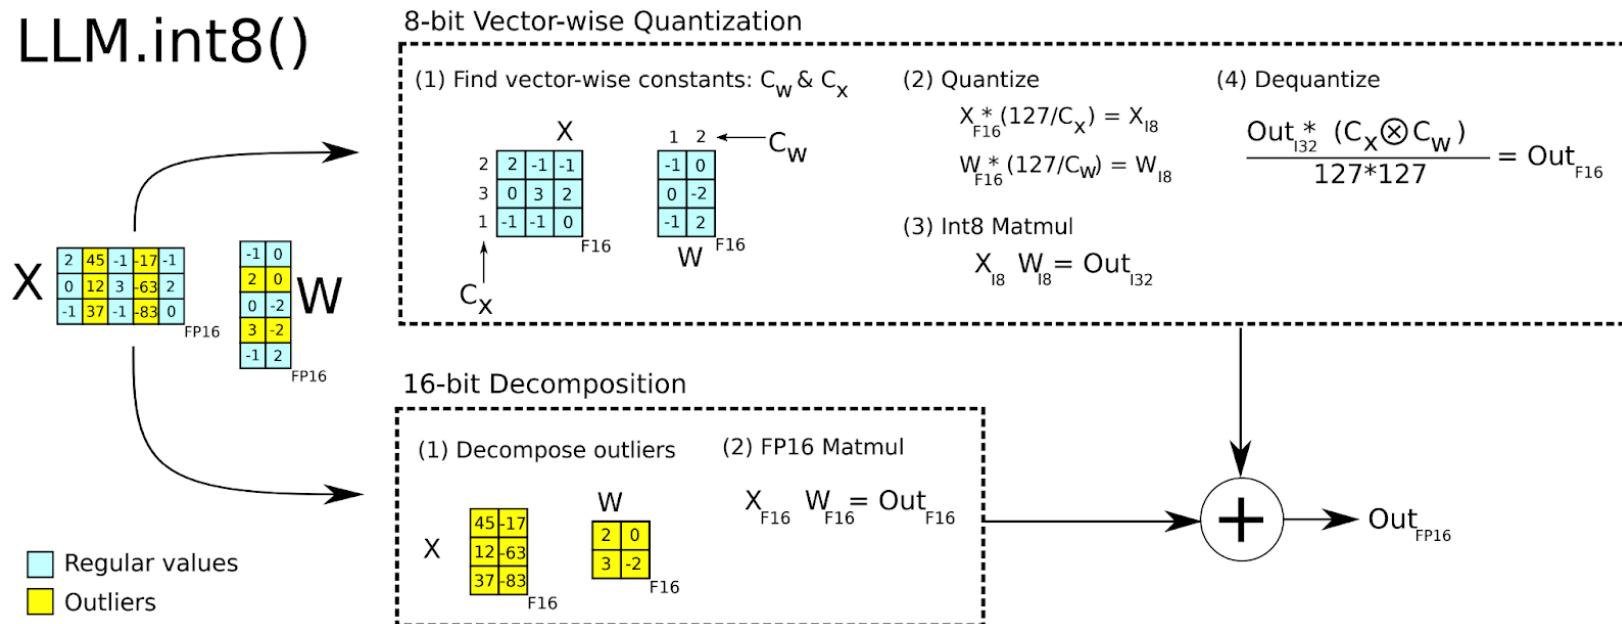
# Quantization of Language Models

Outlier features significantly harms performance after quantization in LMs.



# Quantization of Language Models

LLM.int8()



Keep outlier channels / features in 16-bit, quantize the rest.

# Quantization of Language Models

Parameters	125M	1.3B	2.7B	6.7B	13B
32-bit Float	25.65	15.91	14.43	13.30	12.45
Int8 absmax	87.76	16.55	15.11	14.59	19.08
Int8 zeropoint	56.66	16.24	14.76	13.49	13.94
Int8 absmax row-wise	30.93	17.08	15.24	14.13	16.49
Int8 absmax vector-wise	35.84	16.82	14.98	14.13	16.48
Int8 zeropoint vector-wise	25.72	15.94	14.36	13.38	13.47
Int8 absmax row-wise + decomposition	30.76	16.19	14.65	13.25	12.46
Absmax LLM.int8() (vector-wise + decomp)	25.83	15.93	14.44	<b>13.24</b>	<b>12.45</b>
Zeropoint LLM.int8() (vector-wise + decomp)	<b>25.69</b>	<b>15.92</b>	<b>14.43</b>	<b>13.24</b>	<b>12.45</b>

Zeropoint > absmax because outliers non-symmetric (either very large or very small, but not both)

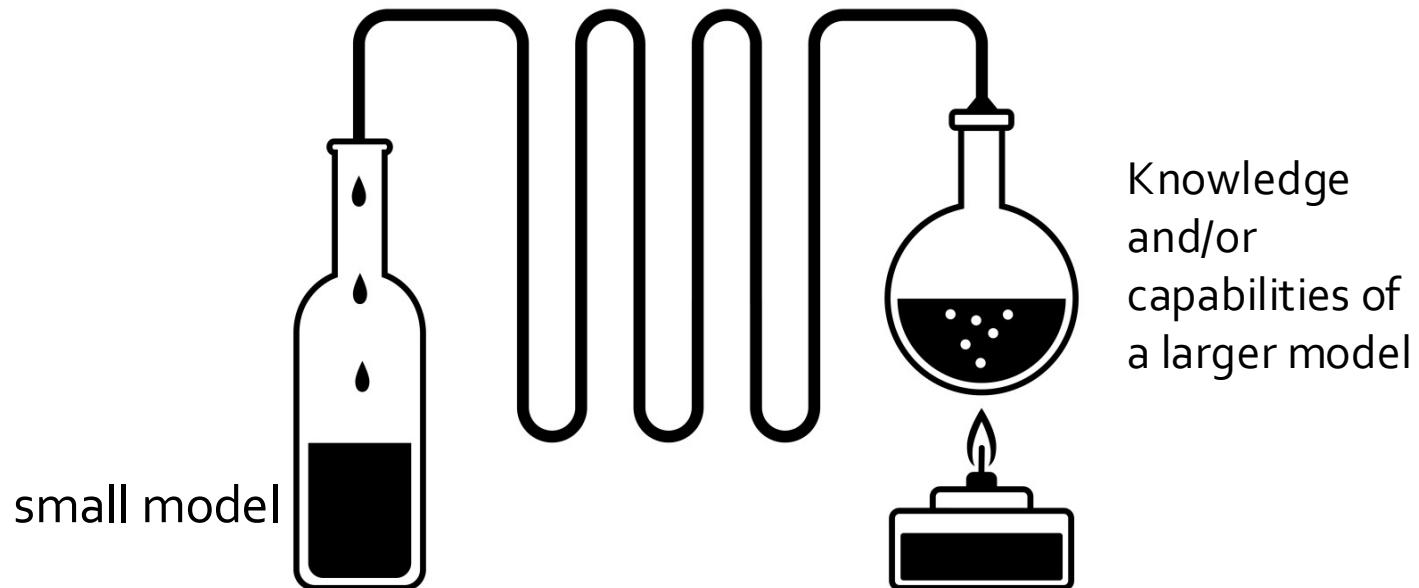
# Quantization of Language Models

- Maps floating point numbers (fp32, fp16, bf16) to low precision numbers (fp8, int8) to save memory.
- Is effective in reducing the memory required for both training / inference.
- 8-bit quantization loses minimal performance, while 4-bit quantization is hard, can be harmful to model performance.

# Distilling the knowledge of larger models

# Distillation

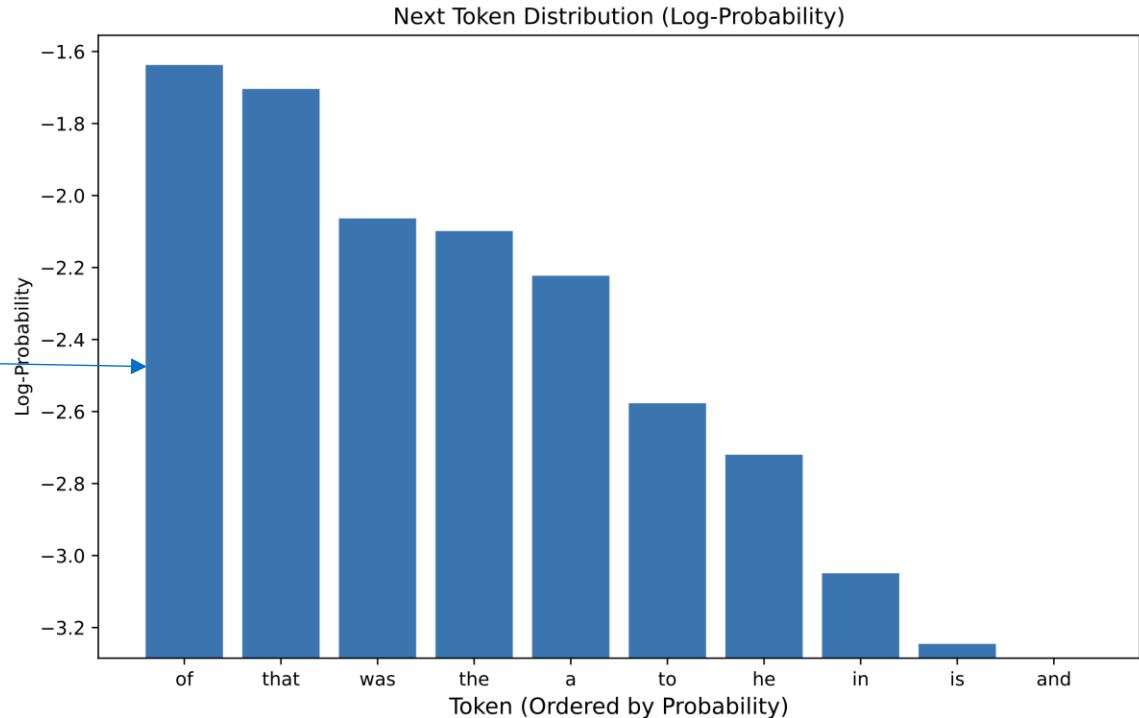
---



# Revisit: Standard Training (NLLoss)

prefix: The strange case \_\_\_\_  
groundtruth: of

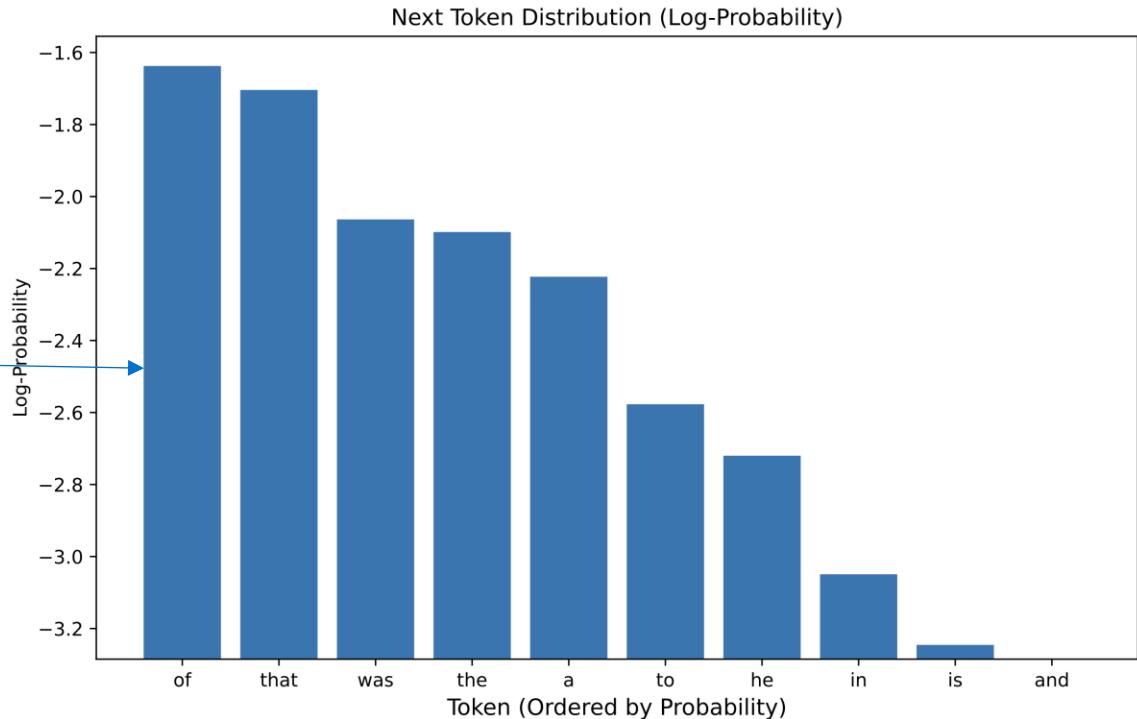
Loss =  $-\log p(\text{of})$   
= Cross Entropy( $y_{\text{pred}}$ ,  
groundtruth)



# Revisit: Standard Training (NLLloss)

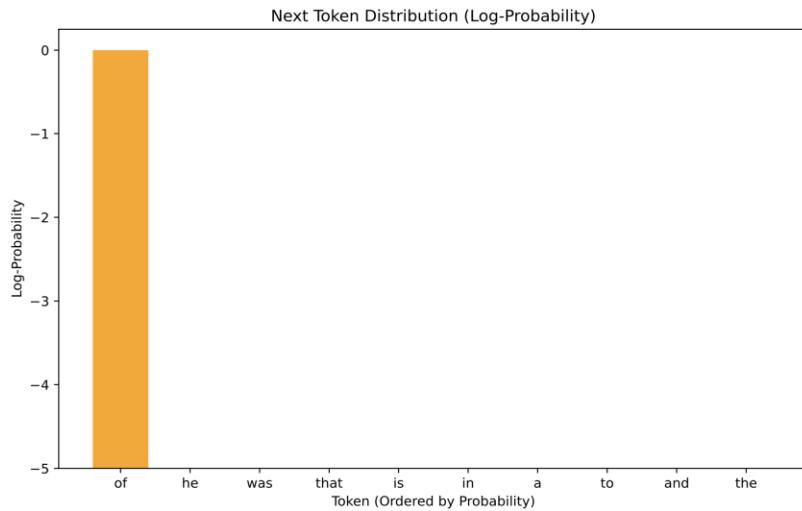
prefix: The strange case \_\_  
groundtruth: of

loss = -log p(of)

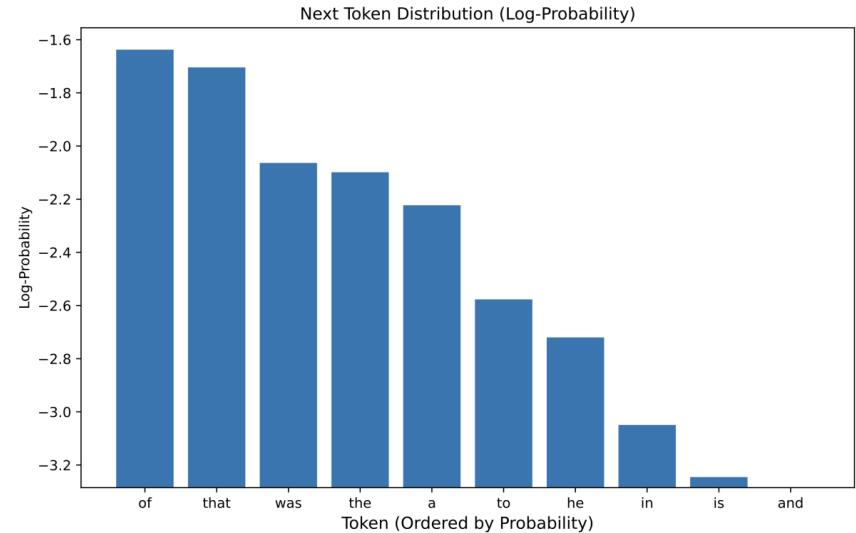


# Revisit: Standard Training (NLLloss)

$$\text{loss} = -\log p(\text{of}) = \text{Cross Entropy}(\text{groundtruth}, \text{y\_pred})$$



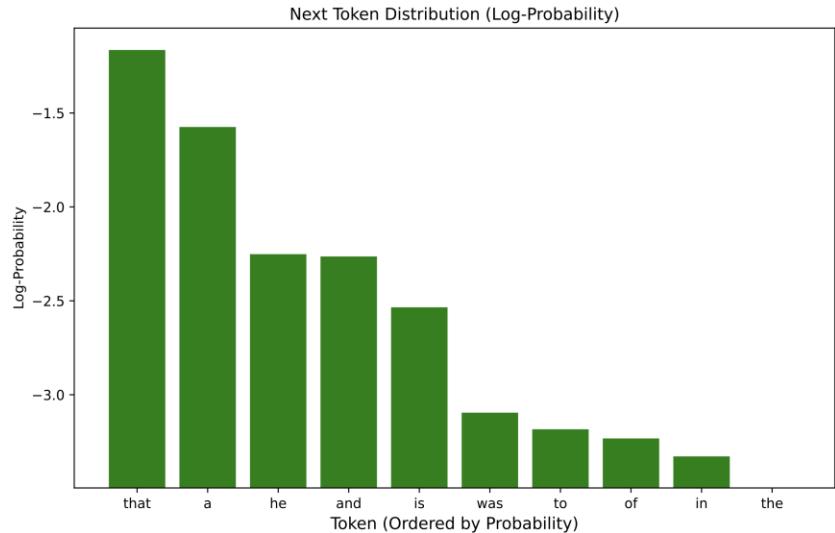
Groundtruth  
(one-hot)



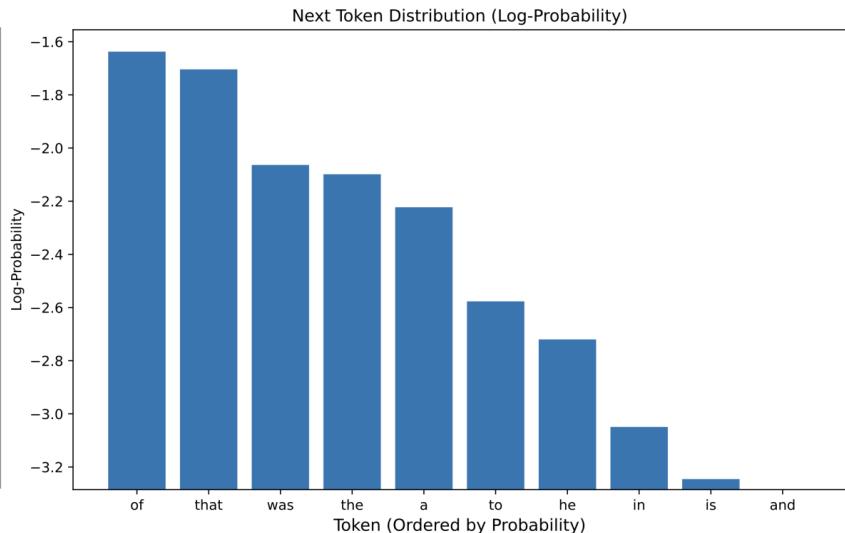
y\_pred

# Knowledge Distillation

$$\text{KD loss} = \text{Cross Entropy}(y_{\text{large}}, y_{\text{pred}})$$



Large model next token  
probs ( $y_{\text{large}}$ )

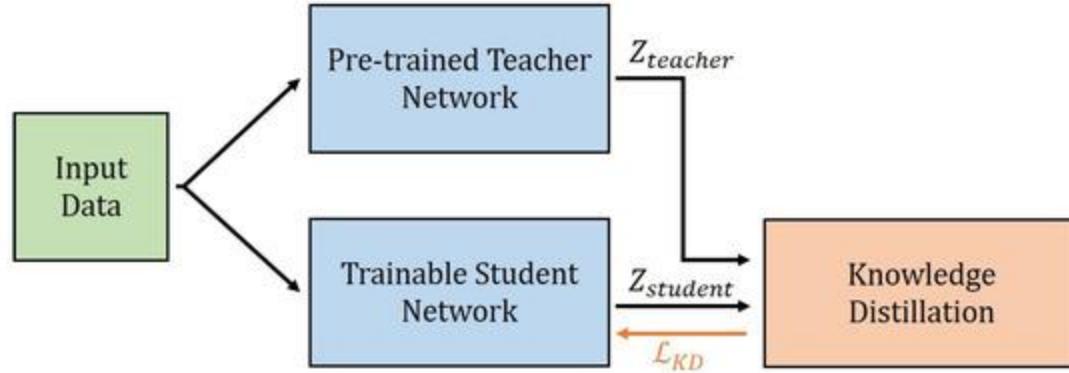


small model next token probs  
( $y_{\text{pred}}$ )

# Knowledge Distillation

Step 1: Initialize teacher model with a large and capable model

Step 2: Feed input data to both student and teacher (freezed)

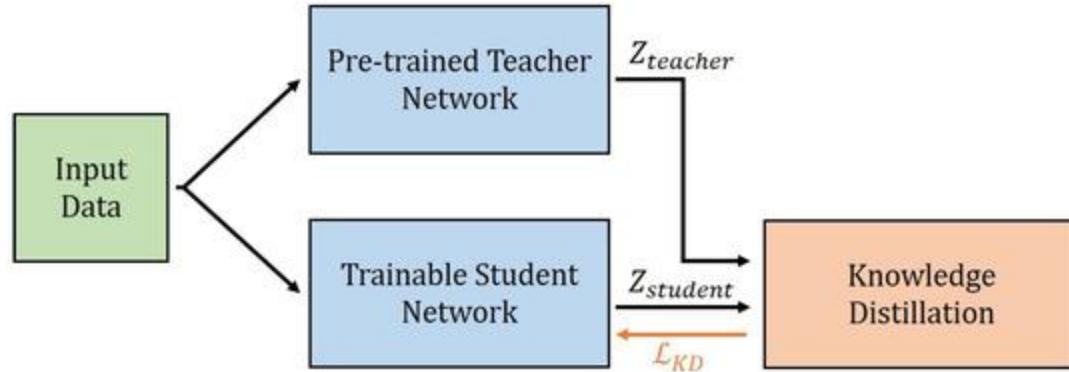


Step 3: Use teacher outputs to train student (Cross Entropy)

# What if the teacher is Proprietary (GPT)?

Step 1: Initialize teacher model with a large and capable model

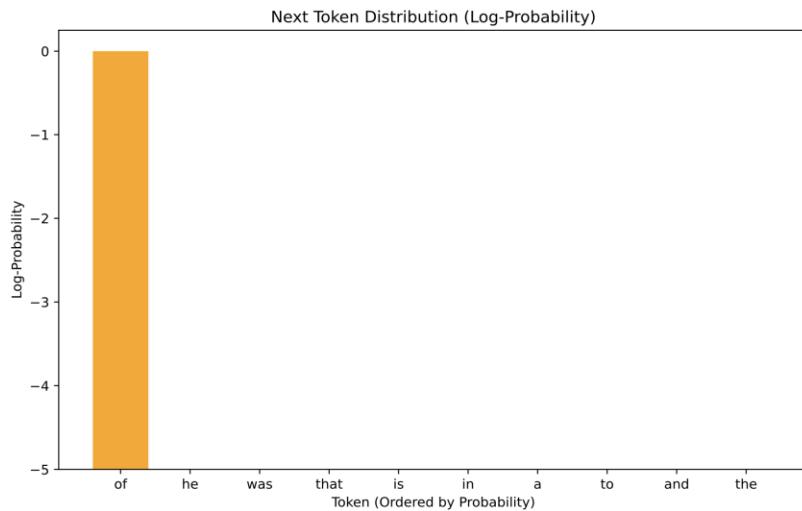
Step 2: Feed input data to both student and teacher (freezed)



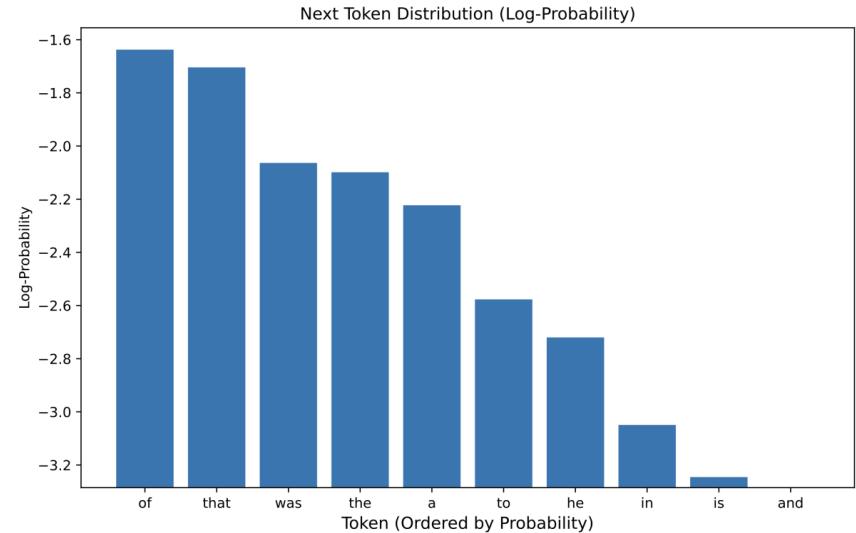
Step 3: Use teacher  
**generations** (instead of  
outputs) to train student!

# Revisit: Standard Training (NLLloss)

$$\text{loss} = -\log p(\text{of}) = \text{Cross Entropy}(\text{sampled text}, \text{y\_pred})$$

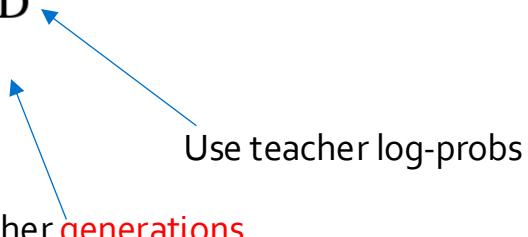


Sampled output  
(one-hot)



y\_pred

# What works better (a study in 2016)

Model	BLEU <sub>K=1</sub>	Δ <sub>K=1</sub>	BLEU <sub>K=5</sub>	Δ <sub>K=5</sub>
<i>English → German WMT 2014</i>				
Teacher Baseline $4 \times 1000$ (Params: 221m)	17.7	—	19.5	—
Baseline + Seq-Inter	19.6	+1.9	19.8	+0.3
<hr/>				
Student Baseline $2 \times 500$ (Params: 84m)	14.7	—	17.6	—
Word-KD	15.4	+0.7	17.7	+0.1
Seq-KD	18.9	<b>+4.2</b>	19.0	+1.4
				

[Sequence-Level Knowledge Distillation](#) (Kim & Rush, EMNLP 2016)

# Knowledge Distillation

---

- Train student (usually smaller model) on the output of a teacher (usually a larger model)
- The output can be log-probabilities or sampled outputs
- Effective in "distilling" the knowledge of large models to smaller ones.